

IS DATA VALUE PATH-ORDERED OR SET-ORDERED? A CARRIER-DEPENDENT SAME-SAMPLE TEST UNDER WSD SCHEDULES

Anonymous authors

Paper under double-blind review

ABSTRACT

Data attribution is usually indexed by a selected set even when optimization follows a schedule. We test whether marginal value changes with a target’s entry position by measuring the same target at multiple positions within each seed. On two frozen synthetic logistic carriers, the paired estimator gives different readings: the registered sine carrier has a negative departure ($g = -0.186$, 95% CI $[-0.258, -0.115]$), whereas the linear carrier is ambiguous ($g = +0.010$, CI $[-0.112, +0.132]$). The paired sign must not be conflated with the cross-sectional Kendall sign: negative g means that moving the same target earlier lowers its value, while negative cross-sectional τ ranks earlier-scheduled targets higher. A pure set-ordered ridge counter-arm reveals a non-negative estimator bias ($g = +0.102$, CI $[+0.028, +0.176]$ on the canonical allocation; positive point estimates on all four allocations). The canonical allocation therefore fires the sealed two-sided binary contract, but the positive bias cannot generate the negative sine departure; no magnitude correction is transported across carriers. Value-conditioned routing also reproduces the cross-sectional B2 signal in a set-ordered world. The study thus provides a controlled, carrier-dependent measurement, not a general decision between path-ordered and set-ordered value.

1 INTRODUCTION

Set-indexed attribution methods predict model behavior from training-set membership (Ilyas et al., 2022; Park et al., 2023), but they do not retain the path by which samples entered optimization. This abstraction leaves a prior measurement question: is per-sample marginal value invariant to entry order? If it is, an order-invariant attribution can summarize sample ranking. If it is not, the same sample set can induce different marginal-value relations under different entry schedules.

Distinguishing these cases requires more than a nonzero correlation. The entry construction must be explicit; the seed, rather than the item, must be the inference unit; a minimum effect must be separated from statistical significance; and the estimator must be measured on a known set-ordered case. We implement this contract in a deterministic logistic-regression bed. The same sample identities and held-out readout are evaluated under three entry rules, using both a path-sensitive leave-one-out value and an order-invariant comparator.

The result has three parts. First, same-sample pairing prevents a cross-sectional selection signal from being read as a path effect. Second, the sine carrier shows a negative paired departure while the linear carrier remains ambiguous. Third, the set-ordered ridge counter-arm exposes a non-negative paired-estimator bias, and a separate routing experiment shows that B2 is a reference-conditioned stress test. These results are limited to two synthetic carriers.

2 PROBLEM AND ESTIMANDS

For seed k , let D_k be the clean training set, G_k a fixed target set, and H_k a fixed held-out split. Entry construction c assigns each target $z \in G_k$ a schedule location $s_c(z)$. With $U(w; H_k)$ denoting




Figure 1: Qualitative measurement schematic extracted from the supplied frozen PDF. Each target is measured at early, middle, and late entry while identity and held-out readout remain fixed. The B2 lane is separate because reference-conditioned routing can create a cross-sectional signal in a set-ordered world. The schematic contains no experimental values.

held-out negative cross-entropy (higher is better), define empirical path value as

$$V_{c,k}(z) = U(w_{c,k}^{+G}; H_k) - U(w_{c,k}^{-z}; H_k), \quad (1)$$

where both fits use the same WSD budget and differ by leaving out target z . The seed-level path statistic is

$$\tau_{c,k} = \text{KendallTau}(\{s_c(z)\}_{z \in G_k}, \{V_{c,k}(z)\}_{z \in G_k}). \quad (2)$$

Negative τ means that earlier-scheduled targets rank higher in held-out marginal value. We average Equation 2 across seeds and form frozen 95% Student- t intervals over seed-level coefficients. A two-sided Wilcoxon signed-rank test corroborates paired constructions. The minimum meaningful magnitude is $\tau_a = 0.10$.

We report a cell as *null* when its interval includes zero, *sig-sub-threshold* when the interval excludes zero but remains inside the magnitude boundary, *sig-supra-threshold* when the entire interval lies beyond the boundary in the effect direction, and *reverse* for a significant reversal. The separate field `exceeds_threshold` records the magnitude decision.

On the same $D_k \cup G_k$ and H_k , the order-invariant comparator fits a regularized logistic model and evaluates the inverse-Hessian leave-one-out approximation

$$V_{\text{set},k}(z) = -\nabla_w U(w_k^*; H_k)^\top H_k^{-1} g_k(z), \quad (3)$$

where H_k is the regularized training Hessian and $g_k(z)$ is the target-sample loss gradient. This is the operational GLM reduction used here, not an implementation of the full TRAK or datamodel pipelines.

3 FROZEN EXPERIMENTAL DESIGN

The experiment uses two 24-dimensional synthetic logistic carriers. The linear carrier contains a clean class-separated pool and conflicting near-duplicates with flipped labels. The second carrier, retained under the frozen key *sine*, uses a shifted Gaussian pool with a fixed linear label boundary; targets negate the first feature and flip the label. The key therefore does not imply general nonlinear coverage.

Entry constructions. Construction A inserts the complete target block from one of the three schedule locations onward and computes leave-one-out value within that block. B1 assigns each

Table 1: Frozen experimental contract. Items within a seed are not independent inference units.

Component	Frozen setting
Carriers	linear; registered sine carrier
Dimension / pool	$d = 24$; 400 clean examples per seed
Split / targets	320 clean train; 80 held out; 40 conflicting targets
Optimizer	full-batch logistic GD; ℓ_2 coefficient 10^{-3}
WSD budget	$T = 300$; 10% warmup, 80% stable, 10% decay
Entry locations	$0.15T, 0.50T, 0.85T$ (rank-linear positions for B1)
Inference	16 seeds; per-seed Kendall τ ; 95% Student- t CI
Magnitude gate	$\tau_a = 0.10$
Negative control	20 seeded trials; Bonferroni/FWER-controlled detector




Figure 2: Frozen six-cell construction sensitivity (16 seeds per cell), extracted from the supplied PDF. Points are mean per-seed Kendall τ ; bars are 95% Student- t intervals; the shaded band is $|\tau| \leq \tau_a = 0.10$. Exact values appear in Table 2.

target a distinct injection step spread from $0.15T$ to $0.85T$. B2 ranks targets by the fixed set-valued reference, partitions the order into thirds, and injects high-, middle-, and low-ranked regions at the early, middle, and late locations. Thus B1 changes per-sample position, whereas B2 also couples reference rank to entry time. The comparison identifies the entry rule, not merely whether position changed.

For Construction A and $t \in \{0.15T, 0.50T, 0.85T\}$, targets outside the chosen block have $s_A(z) = 0$ and targets inside it have $s_A(z) = \lfloor tT \rfloor$. Both leave-one-out fits retain the same WSD budget, matching Equation 1.

Controls and frozen decision. The clean negative control is a strictly convex quadratic ridge substrate with additive, schedule-invariant influence. Independent random position labels use a per-seed critical value Bonferroni-adjusted over 20 trials; the frozen acceptance condition is no fires. A global construction-dependent verdict requires the threshold-crossing pattern to vary by construction or carrier and does not license a universal mechanism claim. The six cells, seeds, and analysis rules were frozen before the writing pass.

4 CONSTRUCTION SENSITIVITY

Cross-sectionally, B2 lies beyond $-\tau_a$ on both carriers, whereas A and B1 do not. Statistical significance alone would give a different headline: sine A and linear B1 exclude zero but remain below the minimum-effect threshold. Because B2 value-conditioned routing can reproduce this pattern on a set-ordered world, the paired analysis below is load-bearing.

Table 2: Frozen six-cell result. Only B2 crosses the magnitude boundary on both carriers. Intervals are per-cell and unadjusted; they are not treated as six independent tests.

Carrier	Construction	τ	95% CI	Four-state label	Exceeds
linear	A: block	+0.007	$[-0.017, +0.032]$	null	false
linear	B1: rank-linear	-0.065	$[-0.125, -0.005]$	sig-sub-threshold	false
linear	B2: region-lumped	-0.321	$[-0.384, -0.257]$	sig-supra-threshold	true
sine	A: block	-0.102	$[-0.130, -0.075]$	sig-sub-threshold	false
sine	B1: rank-linear	-0.011	$[-0.056, +0.033]$	null	false
sine	B2: region-lumped	-0.471	$[-0.502, -0.440]$	sig-supra-threshold	true

Table 3: Frozen construction-axis interval comparison.

Pair	linear		sine		Composite
B2 vs A	separated: $[-.384, -.257]$ [$-.017, .032$]	vs	separated: $[-.502, -.440]$ [$-.130, -.075$]	vs	both
B2 vs B1	separated: $[-.384, -.257]$ [$-.125, -.005$]	vs	separated: $[-.502, -.440]$ [$-.056, .033$]	vs	both
A vs B1	overlap 0.011		separated: $[-.130, -.075]$ [$-.056, .033$]	vs	mixed

4.1 CONSTRUCTION-AXIS SEPARATION

The frozen interval comparison separates B2 from A and B1 on both carriers. A and B1 overlap on linear by 0.011 but are separated on sine. These within-cell interval comparisons are descriptive; direct inference uses within-seed contrasts.

The six paired contrasts share the same 16 seeds and form one correlated family. Raw p values are retained; Bonferroni $\alpha/6 = 0.0083$ leaves the four B2-involving contrasts significant, while the two B1–A values straddle that threshold.

4.2 SAME-SAMPLE PAIRED IDENTIFICATION

Cross-sectional τ is not decisive because B2 routes high-reference-value samples early. We therefore move each target between $0.15T$ and $0.85T$ while holding every other target at its frozen B2 step. With $\Delta(z) = V(0.15T; z) - V(0.85T; z)$, define

$$g_k = \frac{\text{mean}_{z \in G_k} \Delta(z)}{\text{std}_{z \in G_k} \Delta(z)}, \quad (4)$$

where the denominator uses the 40 targets with $\text{ddf}=1$. Negative g means early entry lowers the same target’s value. This interpretation is opposite to negative cross-sectional τ , which ranks earlier-scheduled targets higher; the two signs are never equated.

The sine paired departure survives routing permutation and reverses, rather than reproduces, the cross-sectional sign interpretation. Linear remains unidentified. A per-seed sign-consistency fraction is omitted because, under this construction, it reduces to bookkeeping with no test power; an earlier selection-confound claim based on a sign-consistency code bug is revoked.

4.3 SET-ORDERED RIDGE COUNTER-ARM

The identical paired estimator is run on a strictly convex set-ordered ridge whose marginal value is order-invariant by construction. The sealed contract intended a 16-seed CI covering zero to count as silent and any CI excluding zero to count as a fire.

The canonical ridge allocation literally fires the sealed binary rule. Across all four allocations, however, point estimates are positive; 14 of 16 canonical seeds are positive, and dropping the most extreme seed of either sign preserves the positive direction. The contract’s sign-blind consequence clause is therefore an erratum: a positive set-ordered bias cannot produce the negative sine departure. This is a sign-only conservative argument. We do not subtract the ridge magnitude from the sine value because cross-carrier transport is untested.

Table 4: Same-seed construction contrasts. “Beyond” means the paired CI lies strictly outside $[-0.10, +0.10]$.

Carrier, pair	$\Delta\tau$	95% CI	p	Beyond
linear, B2–A	−0.328	[−0.395, −0.260]	3×10^{-8}	yes
linear, B2–B1	−0.256	[−0.325, −0.186]	1×10^{-6}	yes
linear, B1–A	−0.072	[−0.124, −0.021]	0.009	no
sine, B2–A	−0.369	[−0.421, −0.317]	2×10^{-10}	yes
sine, B2–B1	−0.460	[−0.509, −0.411]	3×10^{-12}	yes
sine, B1–A	+0.091	[+0.041, +0.141]	0.001	no

Table 5: Frozen same-sample paired identification (16 seeds).

Carrier	Paired slope	Four-state reading	Routing robustness
sine	−0.186, CI [−0.258, −0.115]; Wilcoxon $p = 9 \times 10^{-5}$	PATH-SUPPORTED-neg; early entry lowers value	placebo $g = -0.176$
linear	+0.010, CI [−0.112, +0.132]; Wilcoxon $p = 0.74$	AMBIGUOUS; not identified	—

4.4 B2 AS A REFERENCE-CONDITIONED STRESS TEST

On the set-ordered ridge, routing targets with the real B2 reference correlations reproduces the frozen cross-sectional magnitudes; at zero correlation, τ is near zero.

Thus B2 is a reference-conditioned stress test, not standalone path evidence. The paired reading remains subject to the ridge calibration above.

5 SET-VALUED BASELINE AND BED CERTIFICATE

On the same carriers, targets, splits, and constructions, we compute mean within-construction Kendall agreement τ_{AB} between empirical path value and Equation 3. Because frozen uncertainty intervals are unavailable, these values are descriptive.

The detector controls separate aliveness, power on a known path effect, and false alarms on a known set-ordered substrate.

Table 8: Detector calibration. Family sizes are explicit because the unadjusted and FWER-controlled rates use different trial counts.

Control	Family	Fires	Rate
Decoupled positive control (schedule dependence independent of B2 score)	16 seeds	14	87.5%
Order-invariant negative, unadjusted $\alpha = 0.05$	100 trials	3	3%
Order-invariant negative, frozen FWER control	20 trials	0	0% (SILENT)

Table 9: Descriptive agreement with the set-valued comparator.

Carrier	B1 rank-linear	B2 region-lumped
linear	+0.20	+0.30
sine	+0.57	+0.55

Table 6: Set-ordered ridge calibration and allocation sweep. The sine row is shown only as the headline reference.

Substrate / allocation	g	95% CI	Wilcoxon p	Covers 0
ridge, canonical $X_s[0:40]$	+0.102	[+0.028, +0.176]	0.011	no (fires)
ridge, $X_s[40:80]$	+0.192	[+0.130, +0.254]	0.00015	no (fires)
ridge, $X_s[80:120]$	+0.049	[−0.066, +0.165]	0.298	yes
ridge, random slice	+0.145	[+0.065, +0.225]	0.006	no (fires)
sine headline	−0.186	[−0.258, −0.115]	9×10^{-5}	no

Table 7: B2 routing stress test (16 seeds).

Carrier	operating ρ	routing-only τ	frozen B2 τ	zero-corr. τ
linear	0.30	−0.313	−0.321	−0.054
sine	0.55	−0.563	−0.471	−0.054

Table 10: Frozen bed-discriminability certificate. B2 fires as a routing-sensitive probe but is not the known-path positive control.

Reading	linear	sine
Known-path decoupled positive control	14/16 (87.5%)	14/16 (87.5%)
B2 routing-sensitive stress probe	$\tau = -0.321$	$\tau = -0.471$
Known-set negative-control fires	0/20 SILENT	0/20 SILENT
Set estimator alive (fraction distinct)	1.0	1.0
Live τ_{AB} aliveness read	+0.88	+0.72
Verdict	DISCRIMINATING	DISCRIMINATING

No single probe is certified on all three properties. The paired arm, not B2, supplies the evidence for entry-position sensitivity.

6 RELATED WORK

Datamodels and TRAK are indexed by selected sets rather than schedule position (Ilyas et al., 2022; Park et al., 2023). Our narrower question is whether empirical leave-one-out value has a schedule-dependent component when identities and held-out readout remain fixed. Machine-learning conclusions can be sensitive to power, run variance, and reporting choices (Card et al., 2020; Bouthillier et al., 2021; Dodge et al., 2019; Reimers & Gurevych, 2017); we therefore use seeds as inference units and separate confidence against zero from a minimum-effect decision. Training-trajectory suites motivate treating schedules as first-class experimental objects (Biderman et al., 2023), and the frozen analysis follows the rationale for preregistration in NLP (van Miltenburg et al., 2021).

7 LIMITATIONS AND SCOPE

The evidence is limited to deterministic CPU logistic-regression carriers with synthetic conflicts; the registered sine key does not represent general nonlinear models or language models. The paired test is primary, while B2 is a routing-reproducible stress test. Its set-ordered calibration has a non-negative bias, and the canonical allocation fires the sealed two-sided rule; only the sign-based conservative conclusion is supported. Linear is not identified. The comparator is an inverse-Hessian GLM reduction, not end-to-end TRAK or datamodels, and its agreement values lack frozen uncertainty. A single set-ordered negative control cannot exclude every artifact or calibrate natural data. Transfer requires a natural carrier, fresh controls, full same-bed baselines, and agreement uncertainty.

8 CONCLUSION

The same-sample paired test gives carrier-dependent readings: the sine carrier has a negative departure ($g = -0.186$, CI $[-0.258, -0.115]$) whose meaning reverses the cross-sectional sign, whereas the linear carrier is ambiguous ($g = +0.010$, CI $[-0.112, +0.132]$). The set-ordered ridge counter-arm has non-negative bias across four allocations, so it cannot explain the negative sine departure, but the canonical allocation fires the sealed sign-blind binary rule and no cross-carrier magnitude correction is justified. Because routing reproduces the B2 cross-sectional signal on a set-ordered world, B2 remains a stress test. These synthetic results do not decide whether data value is generally path-ordered or set-ordered.

REFERENCES

- Stella Biderman, Hailey Schoelkopf, Quentin Anthony, et al. Pythia: A suite for analyzing large language models across training and scaling. In *Proceedings of the 40th International Conference on Machine Learning*, 2023. arXiv:2304.01373; PMLR 202:2397–2430.
- Xavier Bouthillier, Pierre Delaunay, Mirko Bronzi, et al. Accounting for variance in machine learning benchmarks. In *Proceedings of MLSys*, 2021. arXiv:2103.03098.
- Dallas Card, Peter Henderson, Urvashi Khandelwal, Robin Jia, Kyle Mahowald, and Dan Jurafsky. With little power comes great responsibility. In *Proceedings of EMNLP*, 2020. arXiv:2010.06595.
- Jesse Dodge, Suchin Gururangan, Dallas Card, Roy Schwartz, and Noah A. Smith. Show your work: Improved reporting of experimental results. In *Proceedings of ACL*, 2019. arXiv:1909.03004.
- Andrew Ilyas, Sung Min Park, Logan Engstrom, Guillaume Leclerc, and Aleksander Madry. Data-models: Predicting predictions from training data. In *Proceedings of the 39th International Conference on Machine Learning*, 2022. arXiv:2202.00622.
- Sung Min Park, Kristian Georgiev, Andrew Ilyas, Guillaume Leclerc, and Aleksander Madry. TRAK: Attributing model behavior at scale. In *Proceedings of the 40th International Conference on Machine Learning*, 2023. arXiv:2303.14186.
- Nils Reimers and Iryna Gurevych. Reporting score distributions makes a difference: Performance study of LSTM-networks for sequence tagging. In *Proceedings of EMNLP*, 2017. arXiv:1707.09861.
- Emiel van Miltenburg, Chris van der Lee, and Emiel Krahmer. Preregistering NLP research. In *Proceedings of NAACL*, 2021. arXiv:2103.06944.

A REPRODUCIBILITY AND DISCLOSURE

All runs reported in the supplied manuscript are deterministic CPU experiments with the seed as the statistical unit. The quantitative plot is reproduced here by a faithful extraction from the supplied frozen PDF because the exact historical LaTeX and plotting source were unavailable. Generative tooling assisted language and typesetting; no generated asset created or modified quantitative evidence. The study uses synthetic data and no human participants or personal information.

B FROZEN PROVENANCE

The supplied PDF maps the following artifacts by SHA-256 prefix; this editorial pass transcribes but does not independently execute them.

61894856 R127 order run: candidateB_order_r127.json
2e08b7f9 R126 reconcile: candidateB_reconcile_r126.json
3a1c18da R128 four-state labels: candidateB_relabel_4state_R128.json
3178724b R128 bed certificate: candidateB_bed_discriminability_R128.json

378 **40b74cbc** R133 construction CI: candidateB_construction_axis_CI_R133.json
379 **d74eb619** R127 preregistration supplement.
380 **d0c898fb** R126 preregistration draft.
381 **547978c5** R139 same-sample paired and placebo artifact.
382 **b7028282** R140 rebuild (re-frozen R214).
383 **f6ba5097** R140 rebuild specification.
384 **502c13bf** R141 preregistration, revision 2.
385 **67c52c0b** R141 stress test and routing placebo.
386 **e3431b7f** R148 set-ordered ridge counter-arm preregistration.
387 **6b437ba5** R148 counter-arm output, canonical allocation.
388 **397dd466** R148 counter-arm code.
389 **b81168c4** R148 counter-arm finding record, including the allocation sweep.
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431