

NON-INTERCHANGEABLE SCORING PROTOCOLS AND A DRAW-COUNT-MATCHED REFERENCE: A NARROW MEASUREMENT STUDY ON LLaDA-8B AND GSM8K-MC

Anonymous authors

Paper under double-blind review

ABSTRACT

Multiple-choice accuracy for a diffusion language model is not determined by the model alone: the scorer chooses which option tokens are visible, when they are revealed, and whether revealed ground-truth tokens are fed back into later steps. We audit that measurement construction on one bounded bed: LLaDA-8B-Instruct with balanced four-choice GSM8K-MC.

On a confirmatory $n=400$ pool `s_conf` exceeds the non-injected `mc_uniform` reference by +13.0 pp at the per-candidate $n_{mc,i} = 2L_i$ total-NFE-fair equalization point ($b/c = 109/57$, exact $p = 6.6 \times 10^{-5}$; four tested budgets, conservative read +9.0 pp). This *refutes* the narrow account that the advantage is an artifact of the reference receiving less compute, but establishes no budget-independent law. On a disjoint fresh pool a fixed left-to-right order whose *schedule selection* reads no model output still beats `s_conf`; confidence ordering is therefore neither necessary nor uniquely optimal on this answer format. Rank-agreement and distractor-perturbation diagnostics support only a bed-local, sign-consistent non-interchangeability reading; we do not claim dose-response, cross-bed transfer, or isolated causal attribution to feedback.

1 INTRODUCTION

Masked diffusion language models assign a multiple-choice answer only after a scoring protocol turns option-level denoising losses into a ranking. That protocol is part of the measurement instrument. In particular, a scorer may reveal candidate tokens sequentially and feed the known candidate token back into the context, or it may score a randomly masked subset without such feedback. Even two feedback-based scorers can use different reveal schedules. Treating their accuracies as if they were readings from the same ruler hides a change in what was measured.

We audit this construction on one LLaDA-8B-Instruct $\times$ GSM8K-MC four-option bed: two historical balanced pilot pools, a confirmatory 400-question pool for the reference-arm contrast, and a disjoint fresh 819-question pool for reveal order. The study makes three pool-specific claims:

1. On the tested $n=400$ reference-budget curve, `s_conf`'s advantage over the non-injected reference is positive and significant at every budget and narrows as the reference receives more compute (Table 3, Figure 2). The registered matched- L point remains primary. At the per-candidate total-NFE-fair $2L$ endpoint, where both arms receive $2L_i$ forwards, the same bed retains the +13.0 pp headline gap (exact $p = 4.8 \times 10^{-5}$ and 6.6×10^{-5}); the highest-budget read remains the conservative one. This rules out a less-reference-compute-only account on the tested curve, but not budget dependence off the curve. Because the comparison is a complete protocol swap, it does not isolate feedback causally.
2. On the disjoint fresh pool, reveal order under the tested persisted schedule survives its $k=1$ multiplicity family. A fixed schedule whose selection reads no model output nevertheless beats `s_conf` on average (Table 6); the margin strengthens in the answer-pool construction stratum and reverses at answer position 0 (Appendix Table 8). Confidence ordering is




Figure 1: Conceptual scoring-protocol fork, and the measurement-level view of this study. The same item and model feed all five arms of Table 1: confidence, anti-confidence, fixed, seeded-random, and random-mask paths. Feedback arrows stay inside the four injected paths; the random-mask path has none. Each path ends in the same generic score-vector object, so no winner or item-level ranking is drawn, and the two rulers distinguish nominal draw count from total NFE. This generated artwork contains no experimental values; all quantitative results are data-derived.

therefore neither necessary nor uniquely optimal on this answer format, and the result does not establish an exclusively causal confidence mechanism.

3. The protocols are not interchangeable on this bed: changing the scorer changes per-question option rankings and choice-set sensitivity. Pairwise cross-protocol Kendall τ lies only just above the random arm’s own self- τ floor, so it bounds protocol-specific disagreement rather than proving distinctness. The deterministic `s_conf-fixed_inject` agreement provides a separate comparison (Table 17).

The contribution is a protocol audit rather than a new scoring algorithm: a paired, decision-record-level decomposition of comparisons that would otherwise collapse model and scorer into one accuracy number, with the negative boundary stated up front—neither the design nor its measurements support a dose-response curve, a concavity claim, or transfer to another model, dataset, or judge.

2 RELATED WORK AND LOCAL EVIDENCE CONTEXT

LLaDA introduced the 8B masked-diffusion model used here (Nie et al., 2025); iLLaDA later introduced confidence-based multiple-choice scoring (Nie et al., 2026). That scorer reveals the most confident candidate token and feeds the known token into later steps, so a comparison with a Monte Carlo-style scorer replaces protocol components as well as the accounting rule. We audit that replacement, not the published benchmark values. GSM8K (Cobbe et al., 2021) supplies only the controlled four-choice bed; our claims do not transfer beyond it.

Scorer dependence is established more broadly: benchmarks vary with task and prompt format (Brown et al., 2020; Hendrycks et al., 2021; Lin et al., 2022; Chen et al., 2021); frameworks distinguish raw from length-normalized scoring and expose prompt sensitivity (Liang et al., 2023; Biderman et al., 2024; Zhu et al., 2024); and audits cover judge design, contamination, and aggregation (Perez et al., 2022; Jain et al., 2024; Wang et al., 2023). Our contribution is the measured masked-diffusion protocol swing on a stated budget curve.

The closest mechanism literatures are mask-out objectives (Nie et al., 2025; 2026; Sahoo et al., 2024), multiple-choice normalization and position sensitivity (Holtzman et al., 2021; Zheng et al., 2024), and confidence-guided diffusion decoding (Nie et al., 2025). Our controls extend these con-

Table 1: The five scoring arms. L_i is candidate length. The scoring-forward count excludes the additional L_i adaptive schedule-selection forwards for confidence and anti-confidence order; nominal draw matching is therefore not total-NFE matching.

Arm	Reveal schedule	GT feedback	Scoring forwards
s_conf	min-CE first	yes	L_i (+ L_i select)
rand_inject	seeded random	yes	L_i
anti_inject	max-CE first	yes	L_i (+ L_i select)
fixed_inject	fixed order	yes	L_i
mc_uniform	random mask	no	1 ($n_{mc}=1$)

cerns to token injection: a fixed left-to-right schedule can beat confidence ordering on this answer format, but no cited work or result here establishes a transferable injected-feedback effect.

3 SETUP

Bed and pairing. Each item pairs one GSM8K question with four candidate answers scored by LLaDA-8B-Instruct (Nie et al., 2025) with frozen weights. Choice sets were frozen but differ by pool, and the fresh pool is *internally* heterogeneous by construction: the historical, $n = 400$, and first 336 fresh items use deterministic near-miss distractors; the remaining 483 fresh items use three distinct values drawn from a frozen GSM8K answer pool. Every arm reuses the recorded choice set. Historical and $n = 400$ pools are exactly answer-position balanced; the fresh pool is approximate. Fresh-pool contrasts therefore mix protocol effects with construction stratum, and post-hoc position standardization leaves the fresh-pilot difference, so we report them as strata-confounded. Inference uses BF16, no temperature, and no generation. Confidence and fixed schedules are deterministic; the seeded-random result is one persisted schedule. Because its runner used an unpinned Python string hash, replay is anchored to stored per-candidate orders, not the nominal seed alone. Appendix H lists identifiers, prompts, orders, and seeds. Historical pools use 12 or 16 items per answer position ($n = 48/64$); the $n = 400$ and fresh $n = 819$ pools are non-overlapping. Within each pool, arms score the same items, so contrasts are paired by question. Top-1 exact correctness is the sole endpoint; no partial credit is assigned.

Five arms. For candidate tokens $y_{1:L_i}$, an injected arm chooses a permutation of positions, reads cross-entropy before each reveal, inserts the known candidate token at that position, and repeats until the candidate is visible. It sums the per-step cross-entropies and selects the option with minimum score. s_conf chooses the next minimum-cross-entropy position; rand_inject uses a deterministically seeded random order; anti_inject chooses the maximum-cross-entropy position; and fixed_inject uses fixed left-to-right order. The fifth arm, mc_uniform, draws a masked subset of size $k \sim U\{1, \dots, L\}$ at each of n_{mc} draws (LLaDA upper-bound estimator), reads the partial-mask cross-entropy at the masked positions, and rescales per option by the standard factor L_i/k , so each draw is an unbiased length-corrected estimate of the full mask-out score and the arm returns their mean; ties break by first index. The $n_{mc}=1$ single draw is this paper’s minimum-budget arm; it differs operationally from the LLaDA/iLLaDA default confidence scorer, whose reveal targets the model’s own argmax token rather than the injected ground-truth candidate. Thus “protocol family” denotes the full tuple of feedback rule, reveal schedule, and scoring dose—not the model name alone.¹ Table 1 summarizes these protocol components.

Statistics and derived diagnostics. For paired binary correctness indicators, b and c count discordant directions, $m = b + c$, and $\hat{\Delta} = (b - c)/n$. We report the exact two-sided McNemar test (McNemar, 1947) and the locked 95% interval. Paired intervals default to the registered standard error $\widehat{SE} = \sqrt{m/n}$; each question-bootstrap or paired-bootstrap interval is named as such in its own caption. An alternative mean-corrected form moves the reported registered-SE bounds negligibly on the $n = 400$ pool and is convention-sensitive only on the small pilot pools, changing no verdict anywhere—Appendix Table 9 prints both conventions side by side. The pre-registered interpretation gate requires $m \geq 25$ before a paired contrast is promoted. For protocol agreement,

¹The anonymous artifact describes the exact sampler, prompt, mask construction, random-seed policy, and first-index tie rule.

Kendall τ (Kendall, 1938) is recomputed within each item from the two four-option score vectors after dropping tied pairs, then averaged across items. Choice-set sensitivity is an offline deterministic overlay: each of the three distractors is shrunk by 25%, and a flip is counted if the selected option changes (with perturbation counts stated per pool in the appendix). These are descriptive protocol diagnostics, not hypothesis tests.

Load-bearing estimand. The primary announcement is the paired mean accuracy gap $\Delta = \mathbb{E}_i[\mathbb{P}(A(q_i)=y_i) - \mathbb{P}(B(q_i)=y_i)] = \Pr(A(q)=y) - \Pr(B(q)=y)$ under nominal draw-count matching: q_i is a fixed item tuple (question, designated answer, deterministic four-option pool from the same scorer), the expectation is over the n items of a named pool, and the budget unit is the nominal per-candidate scoring draw (each arm spends L_i forwards per candidate; nominal, not total NFE, per Table 1). Estimation is the paired $\hat{\Delta} = (b - c)/n$ with the exact two-sided McNemar test; Δ is a protocol gap, not a mechanism attribution. Registration kept the nominal matched- L point ($n_{mc}=L$, draw-count only) primary; that primary point and the per-candidate total-NFE-fair equalization at $2L$ are both measured on the $n=400$ bed, and the reported curve still treats the +9.0 pp high-budget endpoint as its conservative read (Table 3).

Multiplicity. The confirmatory five-arm analysis uses all ten pairwise tests as one Holm family (Holm, 1979) at $\alpha = 0.05$; the smallest-step threshold is 0.005. The fresh fixed-order adjudication has one preregistered load-bearing pair (`s_conf-fixed_inject`, $k = 1$). As a conservative sensitivity, we also name a five-test reporting family: `s_conf` versus `rand_inject`, `anti_inject`, `fixed_inject`, `single-draw mc_uniform`, and `draw-count-matched mc_uniform`. Under Holm at $\alpha = 0.05$, the fixed-order result remains significant even under the worst-case factor-five adjustment ($p = 0.0194$). This named five-test set is a robustness disclosure, not a retroactive change to the registered $k = 1$ endpoint.

Registration and principal-contrast selection. The initial registration fixes the bed, five arms, $m \geq 25$ promotion gate, and exact paired testing. Draw-count matching was selected only after the small-pool scan, so every small-sample row is historical/exploratory. A separate confirmatory registration fixed the $n = 400$ draw-count-matched contrast, the ten-test Holm family, and the power stop before data. Its planning MDE and its magnitude gate are distinct, unchanged quantities: the planning MDE sized the $n = 400$ pool, while the same 4 pp promotion bar — a *point-estimate* gate, not an inferential lower-bound gate, so the high-budget read’s lower bound of 3.46 pp neither clears 4 pp nor needs to — applies to the matched- L , total-NFE-fair $2L$, and high-budget reads on that bed, all three of which clear it (Table 3). A separately pre-registered exploratory 25 pp tier applies only to the historical small rows, not as a stricter second bar for any $n = 400$ read; this unified rule *refutes* the earlier asymmetric thresholding. The corrected high-budget $n = 400$ run and the anonymous decision chronology are in Appendix Tables 15 and 14.

4 THE REFERENCE REMAINS LOWER ACROSS THE TESTED BUDGET CURVE ON THIS BED

The answer-level budget curve joins independently registered $n = 400$ runs, with $n_{mc}=128$ as the conservative high-budget main read. Its decisive row is the total-NFE-fair $2L$ point, where both arms spend $2L_i$ forwards per candidate (Table 3). The pre-registered nominal matched- L endpoint instead set $n_{mc,i} = L_i$, matching only the injected arm’s L_i scoring steps and not its schedule-selection passes, token-forwards, or total NFE; that draw-count-matched point is separately measured on this bed (paired $b/c = 106/54$), and the $2L/2L$ row returns the same +13.0 pp gap from a numerically distinct paired record ($b/c = 109/57$), so the two endpoints differ in discordance pattern but not in headline effect. The candidate-length distribution is in Appendix Table 18. At the headline operating point the paired result survives the ten-test Holm family; Table 3 and Figure 2 report the arm accuracies, discordances, estimate, interval, and exact test.

Same-bed likelihood readout. For each deterministic injected policy, summed candidate cross-entropy is the exact negative log-likelihood under that sequential unmasking policy. Since selection is its arg min, score and likelihood rankings coincide *by construction*; this identity is not an additional comparison. Table 2 shows both ground-truth signal and a policy shift on the same questions:




Figure 2: Data-derived paired forest of the tested $n=400$ reference-budget curve: the same four contrasts, frozen records, and intervals as Table 3, with no new decoding. Square area scales with the discordant count m ; filled squares mark question-bootstrap intervals and open squares registered paired-SE intervals. The matched- L primary endpoint and the decisive total-NFE-fair $2L$ equalization coincide at +13.0 pp, while the reference improves with budget, contracting the gap to the conservative +9.0 pp read. All four intervals exclude zero. The dotted 4 pp line is the registered *point-estimate* promotion bar, not a lower-bound requirement, so the high-budget interval reaching 3.46 pp is expected (§3). Regenerated by `paper/src/figures/decisive_paired_forest_refine10.py`.

Table 2: Same-bed per-token negative log-likelihood (nats; lower is better; $n=400$). ρ relates true-option nats to correctness. For `fixed_inject/s_conf`, paired true-option directions are 349/40 with 11 ties ($n_{\text{selected}}=389$, exact two-sided $p = 1.1 \times 10^{-62}$). `mc_uniform` is the separate mask-out family.

Arm	chosen nats/tok	true-opt nats/tok	ρ (true-opt, correct)	acc. %
<code>s_conf</code>	1.75	1.88	-0.12	42.75
<code>fixed_inject</code>	1.02	1.21	-0.28	45.50
<code>anti_inject</code>	1.20	1.43	-0.35	40.75
<code>rand_inject</code>	1.18	1.62	-0.41	39.00
<code>mc_uniform</code> (ELBO family)	0.19	0.77	-0.43	28.75

true-answer readout falls from 1.88 to 1.21 nats/token under `s_conf` versus fixed order. Thus the gap is localized to a protocol-dependent likelihood readout, without a cross-bed claim. These are teacher-forced option likelihoods, not generation; `mc_uniform` remains a separate mask-out family whose ELBO-style estimand is not point-comparable to sequential reveal.

These are operating points of one replacement, not independent triangulation. The monotone reference rise contracts but does not erase the gap, refuting a pure low-compute artifact on the tested curve, exact $2L$ included; any budget-independent law off the measured points remains untested, and the paired result isolates neither feedback, schedule, nor budget.

Table 4 shows every tested reference above its pool’s chance level at the nominal one-sided level and farther above it with budget; under Bonferroni-5 the two lowest-budget $n=400$ reads become *undetermined* against chance rather than at chance, which does not touch the paired contrasts in Table 3 and Appendix Table 15. The $n = 400$ high-budget read is significant and clears the same 4 pp bar as the primary, whereas the $n = 64$ pilot remains underpowered under its exploratory tier and is superseded, at the identical $n_{mc}=128$ operating point, by that measured $n=400$ row. Thus chance calibration and compute both *refute* a pure budget artifact; budget independence as a law off the tested curve stays *honestly unresolved*. The rest of the nested $n = 64$ scan is retained only as

³The $n=400$ single-draw row uses `results/c1_n400_20260819-104302/perq_r0.jsonl`; the matched- L row uses `results/nmc_sweep_grid128_20260823-023717/perq_r0.jsonl`; the $2L/2L$ (total-NFE-fair) and high-budget rows use `results/nmc_sweep_grid128_20260824-192820/perq_r0.jsonl` ($n_{\text{selected}}=400$ in each). “Same bed” is a checked property of the two records, not an assumption: `s_conf` is per-item identical across them (171/400) and so is $n_{mc}=128$ (135/400); only the two reference samplers differ.

Table 3: Tested reference-budget curve in ascending reference budget (reference/ s_{conf} total NFE in parentheses; row labels are the per-candidate budget labels of Appendix Table 19). The decisive row is total-NFE-fair equalization at $2L$, where the reference uses $n_{\text{mc},i} = 2L_i$ forwards and $s_{\text{conf}} L_i$ scoring plus L_i schedule-selection forwards. The pre-registered nominal matched- L point (reference $n_{\text{mc},i} = L_i$, nominal draw-count match only, primary endpoint per registration) is measured on the same $n=400$ items as that decisive row, with the same s_{conf} trajectories (171/400) and the same +13.0 pp headline, so the two endpoints are numerically distinct paired records with interchangeable headline readings.³ The two $n_{\text{mc},i}$ rows quote question-bootstrap intervals (20k, seed 20260824) and the rest the registered paired SE $\sqrt{m/n}$; tests are exact two-sided McNemar; Figure 2 plots the four $n=400$ rows. The $n=64$ scan is exploratory and superseded.

Reference point (ref/ s_{conf} NFE)	n	Acc. pair (s_{conf} /ref)	b/c	$\hat{\Delta}$	95% CI	Exact p	Role
Single draw (1/2L)	400	42.75/28.75%	113/57	+14.0 pp	[7.61, 20.39]	2.1×10^{-5}	corroboration
$n_{\text{mc},i} = L_i$ ($L/2L$)	400	42.75/29.75%	106/54	+13.0 pp	[7.00, 19.00]	4.8×10^{-5}	pre-registered primary
$n_{\text{mc},i} = 2L_i$ ($2L/2L$)	400	42.75/29.75%	109/57	+13.0 pp	[6.75, 19.25]	6.6×10^{-5}	NFE-fair equalization (decisive)
$n_{\text{mc}} = 128$ (128/2L)	400	42.75/33.75%	82/46	+9.0 pp	[3.46, 14.54]	0.0019	conservative read
$n_{\text{mc},i} = L_i$ ($L/2L$)	64	53.12/28.12%	24/8	+25.0 pp	[7.68, 42.32]	0.007	exploratory (superseded)

Table 4: Reference-arm calibration against the chance level of each declared four-choice pool (exact one-sided binomial tests on the arm’s own correct count). Rows 2–4 are the same $n = 400$ pool at three scoring budgets; the last row is a historical pilot on a different pool, not a large-pool substitute. These five tests are the paper’s one uncorrected family, so the last column scores each against the Bonferroni-5 bar $\alpha = 0.05/5 = 0.01$: three rows clear it and the two lowest-budget $n=400$ reads do not. The per-candidate draw-matched row sets $n_{\text{mc},i}=L_i$ per candidate. Those two remain corroborative chance checks either way; no headline claim rests on this family, because the headline claims are the paired McNemar contrasts of Table 3, which are unchanged by any correction applied here.

Reference arm	n	Correct	Accuracy	Chance	Exact p	Bonf.-5
Seeded-random reveal, fresh	819	398	48.60%	25%	6.56×10^{-48}	clears
mc_uniform, single draw	400	115	28.75%	25%	0.0486	no
mc_uniform, per-candidate draw-matched	400	119	29.75%	25%	0.0176	no
mc_uniform, $n_{\text{mc}} = 128$	400	135	33.75%	25%	5.6×10^{-5}	clears
mc_uniform, $n_{\text{mc}} = 128$ pilot	64	27	42.19%	25%	0.0019	clears

exploratory evidence (Appendix Table 15, Figure 5), and the identification projection in §7 is not answer accuracy.

Analysis map. Table 5 consolidates every headline estimand: pool, registration status, direction, multiplicity family, primary interval convention, and final verdict. It is the single index into the Tables 3, 4, 6, Appendix Table 8 and 17 evidence trail; per-estimand alternative conventions remain in Appendix Tables 15, 9.

5 SCORING PROTOCOLS ARE NOT INTERCHANGEABLE ON THIS BED

The score vectors produced by s_{conf} and fixed_inject —two deterministic protocols, so no schedule-seed noise enters—disagree clearly in rank ($\tau \approx 0.48$; Table 17), consistent with genuinely distinct protocols. Between s_{conf} and the seeded-random rand_inject , the item-level τ means sit only 0.01–0.10 above the random arm’s own self- τ floor (Figure 3); because that floor’s dispersion exceeds the gap, we read that value as an *upper bound* on protocol-specific rank disagreement rather than as evidence of distinctness, and the deterministic pair carries the load-bearing

Table 5: Analysis map: headline estimands of this paper. *Pool* = question bed; *Prereg. status* = whether the endpoint was fixed before observing data; *Direction* = pre-registered read direction (+ = first-listed arm better, − = second-listed arm better, 0 = symmetric/diagnostic); *Family* = multiplicity correction under which the displayed *p* is read; *Primary interval* = the registered or displayed uncertainty convention; *Verdict* = the registered or post-hoc reading (Holm survive / Bonf. clear / CI or LB excludes 0 or fails to). Discounted secondary diagnostics (chance calibration, strata) follow the same per-row convention.

Estimand		Pool	Status	Dir	Family	Verdict	Table
s_conf mc_uniform matched- <i>L</i>	vs	<i>n</i> =400	pre-reg. mary	pri- +	Holm-10	Holm (+13.0 pp, [7.00, 19.00])	survives 3
s_conf mc_uniform 2 <i>L</i> /2 <i>L</i>	vs	<i>n</i> =400	post-sel. (decisive)	sec. +	Holm-10	Holm (+13.0 pp, [6.75, 19.25])	survives 3
s_conf mc_uniform 128/2 <i>L</i>	vs	<i>n</i> =400	pre-registered	+ +	Holm-10	Holm vives (+9.0 pp, [3.46, 14.54])	sur- 3
s_conf rand_inject fresh	vs	<i>n</i> =819	post-sel. sec.	+ +	<i>k</i> =1 LB)	(1s LB >0 (+12.58 pp)	6
s_conf fixed_inject fresh	vs	<i>n</i> =819	pre-reg. single pair	−	<i>k</i> =1 (5-test Holm)	A: fixed (−5.62 pp)	wins 6
s_conf fixed_inject pilot	vs	<i>n</i> =400	pre-reg. con- trol	−	reg. MDE ₈₀	NOT DETECTED	6
five-option probe		<i>n</i> =64	pre-reg. Tier- 25	+ +	discordance gate	UNDETERMINED	15 [‡]
s_conf rand_inject pos 0	vs	fresh pos 0	post-hoc (Bonf. 4)	+ +	Bonferroni- 4	Bonf. (+12.56 pp)	survives App. 8
s_conf mc_uniform chance	vs	<i>n</i> =400 draws	post-hoc fam- 5	+ +	Bonferroni- 5	3/5 clear ($\alpha=0.01$)	4
protocol τ		<i>n</i> =64	descriptive	0 −		distinct by floor (0.48)	17

Primary interval conventions by row (all Table 3 and remaining paired rows): registered paired SE $\sqrt{m}/n$ for matched-*L* and 128/2*L*; question bootstrap (B=20k) for 2*L*/2*L*; paired bootstrap for fresh; registered one-sided LB for fresh $\hat{\Delta}$; exact McNemar *p* for strata; exact binomial for chance. [‡]Five-option probe full counts in Appendix Table 15.




Figure 3: Data-derived per-question Kendall τ distributions for s_conf versus rand_inject. Points are deterministic question-level values; boxes show quartiles and the center line shows the median. The mean printed below each box is the locked overlay value rounded to three decimals (Table 17 gives four). The comparison floor is the random arm’s own self- τ (mean 0.2361, Table 17); the cross-protocol means sit only 0.01–0.10 above it, i.e. within the floor’s dispersion, so this panel is descriptive, not inferential.

reading. Replacing the scoring protocol is accordingly not a harmless implementation detail, even among protocols that use ground-truth feedback.

Table 6: Reveal-order evidence and attribution control. Fresh is secondary, pilot undetermined, and pooled a heterogeneous-pool disclosure; their magnitudes differ ($z=2.41$, $p=0.016$; robust $z=2.32$). The fixed control bounds `s_conf`’s advantage at +2.32 pp (95% CI upper), and only 2.1% of order traces match. Cells mark one-sided (1s, constant 2.4865) or two-sided (2s, 2.8016) intervals/MDE₈₀, each as registered per endpoint (the fresh endpoint’s registered rule is one-sided); Appendix Figure 4 uses the paired-bootstrap two-sided CI for the fresh row as a display convention. Raw exact two-sided McNemar p per row: fresh 204/101, $p=3.7 \times 10^{-9}$; fresh-fixed 99/145, $p=0.0039$; fixed control $p=0.334$ (pilot/pooled as per their registered rules). All use paired SE $\sqrt{m/n}$. The fresh row collapses the two construction strata that Appendix Table 8 reports separately.

Reading (contrast)	n	$\hat{\Delta}$	m	Interval / LB	MDE ₈₀	Role
Fresh ^d	819	+12.58 pp	305	1s LB +9.07 pp 2s CI	5.30 pp (1s) 8.32 pp	post-selection secondary
Pilot ^d	400	+3.75 pp	141	[−2.07, +9.57] pp	(2s) 4.31 pp	undetermined secondary
Pooled ^d	1219	+9.68 pp	446	1s LB +6.83 pp 2s CI	(1s) 7.24 pp (2s)	disclosure ^a
Fresh fixed ^{b,e}	819	−5.62 pp	244	[−9.36, −1.88] pp 2s CI	5.34 pp (2s) 7.24 pp (2s)	A: fixed wins not detected
Fixed control ^e	400	−2.75 pp	107	[−7.82, +2.32] pp	CI upper +2.32 pp ^c	($p=0.334$)

^a A sample-size-weighted average of the two heterogeneous pools, *not* a cross-bed mechanism constant; retained for completeness after the registered heterogeneity check ($z=2.41$, $p=0.016$) triggered the disclosed post-data elevation of the fresh endpoint.

^b Fresh `s_conf` versus `fixed_inject` is the single pre-specified pair ($k=1$) and survives the worst-case five-test Holm family ($p=0.0194$). Tagged MDE₈₀ cells use the caption’s tail convention and claim no equivalence; Appendix Table 9 shows the alternative SE convention changes no verdict.

^c That row’s role is an exclusion rather than a detection, so it intentionally carries *both* its MDE₈₀ and its CI upper bound.

^d Contrast: `s_conf` versus `rand_inject` (seeded-random). ^e Contrast: `s_conf` versus `fixed_inject` (fixed order).

Choice-set perturbation provides a second, operational reading. Appendix Table 16 reports the exact-count rates for both protocols at $n = 64$; the historical $n = 48$ result is sign-consistent. The between-protocol gap persists in sign across the two samples, although it narrows in the larger sample. The 25% overlay is a design-chosen proportional shrink of each protocol’s own distractor scores, not a scale-invariant or scale-fair one, so we compare only persistence and sign, never cross-scorer magnitudes, and no significance or power claim attaches to these deterministic perturbations. Together with rank disagreement, they support only a sign-consistent, bed-local non-interchangeability reading.

6 STRUCTURED REVEAL BEATS THE TESTED SEEDED-RANDOM POLICY, BUT CONFIDENCE IS NOT UNIQUE

Sections 4–5 establish the reference gap and bed-local non-interchangeability, not reveal order. After the registered $n=400$ contrast was *undetermined* (Appendix A), a fresh, disjoint $n=819$ pool used the same pairing; its endpoint is one recorded seeded-random policy, not a schedule-seed average.

Persisted-schedule sensitivity. The single pinned seeded-random schedule is one realization, not an average over the seeded-random family. Two existing checks sharpen this boundary without new decoding. The registered pilot-versus-fresh heterogeneity test gives $\delta_{\text{diff}} = 8.83$ pp, $z = 2.41$, and two-sided $p = 0.016$, so pool composition alone produces shifts of roughly 8 pp with the schedule pinned. The persisted schedule matches the fixed left-to-right trace on only 2.1% of items (17/819; Table 6). Appendix Table 13 therefore records the fresh endpoint as post-selection secondary. Multi-seed replays with pinned schedule seeds remain the required follow-up (§7).

Fresh-pool strata audit. Appendix Table 8 re-tabulates frozen per-item contrasts by construction and answer position. The `s_conf`-over-random margin survives both constructions; the fixed-order advantage retains its sign but is significant only in the answer-pool stratum, and reverses at posi-

tion 0 (+12.56 pp, $p=0.0031$). The position audit is post-hoc (not pre-registered), but positions 0/1 ($p=0.0031/0.0032$) survive Bonferroni-4 ($\alpha/4=0.0125$), making the reversal multiplicity-controlled rather than selected. These are paired-record re-tabulations, not new decoding.

Endpoint decision tree. Appendix Table 13 separates the pre-data rules from the post-data endpoint elevation. Preregistration did *not* assign fresh-only priority: step 4 is an explicitly post-data choice, motivated by the failure of a common-effect interpretation.

Pool heterogeneity is disclosed, not merged. Appendix Figure 4 and Table 6 retain every pool. The fresh comparison is significant under its registered one-sided $\alpha = .05$ rule but was elevated only after heterogeneity was observed, so it is post-selection secondary evidence; pooled remains a weighted disclosure. Charging two-pool selection still gives a fresh lower bound of +8.40 pp at $\alpha = 0.025$, and the pre-specified fresh `s_conf-fixed_inject` adjudication survives both multiplicity accounts (Table 6, note b). This supports a structured-order boundary, not equivalence or a uniquely confidence-driven mechanism.

Boundary: what “reveal order” attributes here, and what it does not. The `s_conf-rand_inject` ablation covers one persisted random schedule. A fixed order that selects its schedule without model output beats `s_conf` on the fresh pool on average, but the margin is construction-dependent and reverses at answer position 0 (Appendix Table 8); the pooled-fresh sign repeats at $n=400$ and reverses on historical $n=64$ (Figure 5). No fixed-order cross-pool test supports merging. All six $n=400$ order contrasts remain Holm-nonsignificant, including the cleanest same-bed zero-effect row, `s_conf` versus `anti_inject` (+2.00 pp, exact $p=0.573$; Appendix Table 9). Thus confidence is neither necessary, uniquely optimal, nor exclusively causal on this bed.

Boundary effect and winner-course description. Appendix A gives the pilot/extension winner-course analysis, the $n=400$ undetermined reading, both registered MDE forms, and power rules.

7 LIMITATIONS

Bed, mechanism, and reliability. One checkpoint and one four-choice GSM8K bed support no cross-model, task, or judge prediction. The `s_conf-mc_uniform` swap changes feedback, path, and passes jointly, so the tested compute result does not factorize the mechanism. The frozen re-run and leave-one-out audit pass their stated gates, but neither adds multi-seed evidence; multi-seed accuracy and τ require new decoding, and the 48-item self- τ remains a noise floor (Appendix Table 21).

External validity and compute. The design supports no dose response, concavity, or cross-bed prediction band; the five-choice probe remains *UNDETERMINED* (Appendix Table 11), and length-normalized likelihood was not measured. The compute ledger covers scoring rather than generation. Exact $2L$ rules out the narrow artifact only on the tested curve, while the historical high-budget $n = 64$ row remains underpowered (Table 3; Appendix Table 20).

8 CONCLUSION

On this single bed, the scoring protocol materially changes the measurement. Confidence injection beats the tested non-injected reference across the measured budget curve, including exact $2L$, which rules out only a less-reference-compute account. A fixed schedule wins on average on the fresh pool but reverses at position 0 (Appendix Table 8), so confidence ordering is neither necessary nor uniquely optimal on this answer format. These findings establish bed-local protocol non-interchangeability; off-curve, multi-seed, causal, and generalization claims remain unresolved (Appendix H).

REFERENCES

Stella Biderman, Hailey Schoelkopf, Lintang Sutawika, Leo Gao, Jonathan Tow, Baber Abbasi, Alham Fikri Aji, Farrukh Akhtar, Taufiq Alam, Alon Albalak, Wissam Ammar, et al. Lessons

- from the trenches on reproducible evaluation of language models, 2024. Venue: arXiv preprint; volume and pages not applicable.
- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners, 2020. Venue: NeurIPS 2020.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, Felipe Petroski Such, Dave Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William Hebgen Guss, Alex Nichol, Alex Paino, Nikolas Tezak, Jie Tang, Igor Babuschkin, Suchir Balaji, Shantanu Jain, William Saunders, Christopher Hesse, Andrew N. Carr, Jan Leike, Josh Achiam, Vedant Misra, Evan Morikawa, Alec Radford, Matthew Knight, Miles Brundage, Mira Murati, Katie Mayer, Peter Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. Evaluating large language models trained on code, 2021. Venue: arXiv preprint; volume and pages not applicable.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems, 2021. Venue: arXiv preprint; volume and pages not applicable.
- Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Chapman & Hall/CRC, 1994.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding, 2021. Venue: ICLR 2021.
- Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.
- Ari Holtzman, Peter West, Vered Shwartz, Yejin Choi, and Luke Zettlemoyer. Surface form competition: Why the highest probability answer isn’t always right. In *Proceedings of EMNLP*, 2021.
- Naman Jain, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar-Lezama, Koushik Sen, and Ion Stoica. Livecodebench: Holistic and contamination free evaluation of large language models for code, 2024. Venue: ICLR 2025.
- Maurice G. Kendall. A new measure of rank correlation. *Biometrika*, 30(1/2):81–93, 1938. doi: 10.2307/2332226.
- Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, Benjamin Newman, Binhang Yuan, Bobby Yan, Ce Zhang, Christian Cosgrove, Christopher D. Manning, Christopher Ré, Diana Acosta-Navas, Drew A. Hudson, Eric Zelikman, Esin Durmus, Faisal Ladhak, Frieda Rong, Hongyu Ren, Huaxiu Yao, Jue Wang, Keshav Santhanam, Laurel Orr, Lucia Zheng, Mert Yuksekgonul, Mirac Suzgun, Nathan Kim, Neel Guha, Niladri Chatterjee, Omar Khattab, Peter Henderson, Qian Huang, Ryan Chi, Sang Michael Xie, Shibani Santurkar, Surya Ganguli, Tatsunori Hashimoto, Thomas Icard, Tianyi Zhang, Vishrav Chaudhary, William Wang, Xuechen Li, Yifan Mai, Yuhui Zhang, and Yuta Koreeda. Holistic evaluation of language models, 2023. Venue: Transactions on Machine Learning Research (TMLR), 2023; volume and pages not applicable.
- Stephanie Lin, Jacob Hilton, and Owain Evans. Truthfulqa: Measuring how models mimic human falsehoods, 2022. Venue: ACL 2022.

- Quinn McNemar. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2):153–157, 1947. doi: 10.1007/BF02295996.
- Shen Nie, Fengqi Zhu, Zebin You, Xiaolu Zhang, Jingyang Ou, Jun Hu, Jun Zhou, Yankai Lin, Ji-Rong Wen, and Chongxuan Li. Large language diffusion models, 2025. Venue: ICML 2025.
- Shen Nie, Qiyang Min, Shaoxuan Xu, Zihao Huang, Yuxuan Song, Yong Shan, Yankai Lin, Wayne Xin Zhao, Chongxuan Li, and Ji-Rong Wen. Improved large language diffusion models, 2026. Venue: arXiv preprint; volume and pages not applicable.
- Ethan Perez, Sam Ringer, Kamilė Lukošiušė, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, Andy Jones, Anna Chen, Benjamin Mann, Brian Israel, Bryan Seethor, Christopher McKinnon, Christopher Olah, Da Yan, Daniela Amodei, Dario Amodei, Dawn Drain, Dustin Li, Eli Tran-Johnson, Guro Khundadze, Jackson Kernion, James Landis, Jamie Kerr, Jared Mueller, et al. Discovering language model behaviors with model-written evaluations, 2022. Venue: ACL 2022.
- Subham Sekhar Sahoo, Marianna Arriola, Yair Schiff, Aaron Gokaslan, Edgar Marroquin, Justin Chiu, Alexander Rush, and Volodymyr Kuleshov. Simple and effective masked diffusion language models. In *Proceedings of NeurIPS*, 2024.
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models, 2023. Venue: ICLR 2023.
- Edwin B. Wilson. Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927. doi: 10.1080/01621459.1927.10502953.
- Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. Large language models are not robust multiple choice selectors. In *Proceedings of ICLR*, 2024.
- Kaijie Zhu, Jindong Wang, Jiaheng Zhou, Zichen Wang, Hao Chen, Yichong Xiong, Ziqian Ma, Wei Liu, Lingjuan Sun, Xinyu Li, Neil Zhenqiang Gong, and Xing Xie. Promptrobust: Towards evaluating the robustness of large language models on adversarial prompts, 2024. Venue: arXiv technical report, v5 July 2024; volume and pages not applicable.

A BOUNDARY ANALYSIS AND REGISTERED VERDICTS

Fresh-pool recipe. The $n = 819$ study uses four-option GSM8K multiple-choice items drawn without question overlap with the $n = 400$ pool; unlike the exactly balanced $n = 400$ pool, its answer positions are only approximately balanced. All arms share the same model checkpoint, prompts, item order, answer extractor, and recorded choice sets. The seeded-random arm is bound to its persisted per-candidate orders because the original runner’s nominal seed derivation included an unpinned Python string hash. The paired record contains one row per question and arm with the prediction, reference answer, correctness, candidate-score vector, and order trace. The anonymous artifact bundle additionally contains the preregistration, execution environment, verdict, and the heterogeneity recomputation recorded in Table 13.

Table 7: Boundary ledger for `s_conf` versus `rand_inject`, except the directional defence. The original $n = 400$ reveal-order endpoint remains UNDETERMINED; no post-data power calculation changes that label.

Reading	n	$\hat{\Delta}$	m	Exact p	95% CI	Registered reading
First bed	48	+16.67 pp	22	–	[−2.49, +35.82]	Below discordance gate
Extension	64	+7.81 pp	23	–	[−6.87, +22.50]	Below discordance gate
Pilot	400	+3.75 pp	141	0.238	[−2.07, +9.57]	UNDETERMINED
Directional defence	400	+1.75 pp	159	0.634	[−4.43, +7.93]	Wrong sign

As Table 7 shows, the small-bed rows do not distinguish regression to the mean from sampling variation. We use boundary effect only as an analysis-risk description, not as an effect-halving claim.

Table 8: Fresh-pool strata split on the enlarged $n=819$ bed; b/c and p are paired discordants and exact two-sided McNemar. The three positive estimates are distinct estimands rather than replications: fresh s_conf -vs- $rand_inject$ +12.58 pp ($n=819$), paired-bed s_conf -vs- $mc_uniform$ +13.0 pp ($n=400$ at both $L/2L$ and $2L/2L$), and heterogeneous pooled s_conf -vs- $rand_inject$ +9.68 pp ($n=1219$). The position audit is post-hoc and uses the Bonferroni-4 control stated in the main text.

Split	vs	n	Acc. pair (s_conf /other)	$b/c; \hat{\Delta}$	exact p
Constr.: near-miss (336)	$rand_inject$	336	47.02/33.33%	90/44; +13.69 pp	8.7×10^{-5}
Constr.: answer-pool (483)	$rand_inject$	483	71.01/59.21%	114/57; +11.80 pp	1.6×10^{-5}
Constr.: near-miss (336)	$fixed_inject$	336	47.02/50.00%	51/61; -2.98 pp	0.395
Constr.: answer-pool (483)	$fixed_inject$	483	71.01/78.47%	48/84; -7.45 pp	0.0022
Answer pos. 0	$fixed_inject$	199	65.83/53.27%	46/21; +12.56 pp	0.0031
Answer pos. 1	$fixed_inject$	207	61.84/72.46%	15/37; -10.63 pp	0.0032
Answer pos. 2	$fixed_inject$	226	62.83/74.34%	18/44; -11.50 pp	0.0013
Answer pos. 3	$fixed_inject$	187	53.48/65.78%	20/43; -12.30 pp	0.0052

For the $n = 400$ pilot, the interval includes zero and relevant positive effects; the pre-registered terminal reading is thus UNDETERMINED, not equivalence or no effect.

Provenance. The fresh estimate draws on $c1_decisive_20260819-212944$ and $c1_fixed_inject_fresh_20260821-062533$ ($n=819$); the paired-bed estimate uses Trial 1895183 ($n=400$); and the pooled estimate is their sample-size-weighted average ($n=1219$). The beds, comparators, and estimands differ, so no single mechanism is attributed to their differences.

Reveal-order effect by pool (Delta = s_conf - $rand_inject$)




Figure 4: Three pool contrasts from Table 6. Fresh uses a paired-bootstrap 95% CI, pilot the registered-SE two-sided CI, and pooled a one-sided lower bound (open upper cap is display-only). Square area scales with m ; significant heterogeneity makes pooled secondary.

Table 9: All six injected-order contrasts on the $n = 400$ pool. Exact tests are two-sided and Holm-adjusted within this registered family of six (the six pairwise order contrasts are the declared secondary family; the confirmatory $n=400$ budget-curve family of §4 is the separate ten-test Holm family). Each interval is $\hat{\Delta} \pm h$. The last two columns give the two paired-SE conventions: the manuscript quotes the registered $h = 1.96\sqrt{m}/n$ throughout, and the mean-corrected $h = 1.96\sqrt{m - (b - c)^2/n}/n$ differs from it by at most 0.04 pp here (and by at most 1.2 pp on the small $n = 48/64$ pools, where no verdict changes either). Holm adjustment is step-down with monotonicity enforced, so adjusted p is non-decreasing in raw p ; in particular the largest raw p in the family inherits Holm $p = 1.000$.

Pair	m	$\hat{\Delta}$	Exact p	Holm p	Mean-corr. h	Registered h
s_conf-rand_inject	141	+3.75 pp	0.238	0.953	5.807 pp	5.818 pp
s_conf-anti_inject	154	+2.00 pp	0.573	1.000	6.078 pp	6.081 pp
s_conf-fixed_inject	107	-2.75 pp	0.334	1.000	5.061 pp	5.069 pp
rand_inject-anti_inject	159	-1.75 pp	0.634	1.000	6.176 pp	6.179 pp
rand_inject-fixed_inject	138	-6.50 pp	0.033	0.198	5.721 pp	5.756 pp
anti_inject-fixed_inject	153	-4.75 pp	0.145	0.727	6.043 pp	6.061 pp

B LENGTH-IDENTITY CHECK (NO ACCURACY CLAIM)

This arithmetic projection consumes only the designated-candidate length row of Table 18; it contains no accuracy outcome and does not replace the measured $n=400$ total-NFE-fair result.

Table 10: Length-identity sensitivity (arithmetic projection, no accuracy content). The observed contrast is zero because the probe flag is arm-independent; nonzero cells are closed-form functions of the disclosed designated-candidate lengths. This does not substitute for the measured total-NFE-fair $2L/2L$ row of Table 3. Intervals resample lengths, not accuracy effects. The last row is a lower audit counterfactual based on the nominal all-candidate distribution, not the estimator input.

Restart fraction	0.4	0.6	0.8
Observed-window gap (pp)	0.00	0.00	0.00
Projected gap (pp)	+3.161	+0.692	+0.061
95% CI (pp)	[3.063, 3.252]	[0.657, 0.726]	[0.056, 0.066]
All-candidate substitution (pp)	2.61	0.51	0.04

For restart fraction m , designated-candidate indicator ι_i , and $n_c(q)$ candidates in question q , the closed form is

$$\text{NFE}_i(m) = \min(6/L_i, 2) (1 - m_L^*)^{-2} m, \quad m_L^* = (1 - 2^{-1/L_i})^{-1},$$

$$\hat{\Delta}_q(m) = n_c(q)^{-1} \sum_{i \in q} \iota_i (1 - m)^{L_i}.$$

The reported mean spans $Q=819$ questions and its interval uses $R=2000$ question bootstraps (seed 20260821) (Efron & Tibshirani, 1994); the zero row is an arm-independent identity, and the nonzero rows are length-model projections rather than evidence about an accuracy effect.

C FIVE-OPTION EXTERNAL-VALIDITY PROBE

Table 11: Five-option probe. The terminal label follows the pre-registered rule and is not reinterpreted after observing power. The significant point estimate below therefore carries no evidence-of-absence interpretation and no external-validity claim; it is not cited as a measured conclusion anywhere outside this table and the matching limitation in §7.

Metric	Value
Items decoded	400/400
s_conf accuracy	62.25%
rand_inject accuracy	53.75%
Δ	+8.5 pp
Discordant b/c	84 / 50
One-sided exact p	0.0021
One-sided 95% lower bound	+3.74 pp
Planning point	7.59 pp at $m = 149$
Realized discordance	$m = 134$
Registered verdict	UNDETERMINED

The realized power calculation is descriptive. It neither replaces the registered discordance gate nor upgrades this result to supporting evidence.

The Tier-25 gate ($m \geq 25 \wedge |\hat{\Delta}| \geq 25$ pp) is a historical-pool promotion threshold, frozen at commit 911d60d3 on 2026-08-20, that predates this five-option probe. Its 25-pp magnitude was fixed from earlier historical-pool uplift records and was not tuned post hoc against the probe’s outcomes. The gate applies only to promotion; it does not referee any §1 headline.

D REGISTRATION AND ANALYSIS CHRONOLOGY

Table 12: Complete interval and decision conventions behind Table 5. The question bootstrap is the printed primary interval for matched- L and $2L$; the registered paired-SE convention is the printed primary interval for single draw and $n_{mc}=128$. Thus the two conventions are explicit rather than silently mixed.

Rows	Primary uncertainty	Test / family	Verdict rule
Single draw; $n_{mc}=128$	Registered $\hat{\Delta} \pm 1.96\sqrt{m}/n \times 100$ pp	Two-sided exact McNemar; ten-test Holm family	Adjusted $p < .05$ and $ \hat{\Delta} \geq 4$ pp
Matched- L ; total-NFE $2L$	Question bootstrap (intervals in Table 3); registered-SE values are printed alongside in the main analysis map	Same ten-test Holm family	Same adjusted- p and 4 pp promotion rule
Fresh s_conf- rand_inject	Registered one-sided $\hat{\Delta} - 1.645\sqrt{m}/n \times 100$ pp; paired-bootstrap two-sided CI only in Figure 4	Single registered directional contrast	Positive lower bound; pooled/fresh sign conflict forces UNDETERMINED
Fresh / pilot s_conf- fixed_inject	Paired-bootstrap two-sided CI (fresh); registered-SE two-sided CI (pilot)	Single fresh pair; fixed five-test sensitivity uses $\alpha = .01$	Negative CI for fresh; pilot retains its registered no-detection reading
Five-option probe	Registered one-sided lower bound	Exact directional McNemar, separate Tier-25 gate	$m \geq 25$, positive lower bound, and $ \hat{\Delta} \geq 25$ pp jointly; otherwise UNDETERMINED
Reference chance calibration	Exact two-sided binomial interval/test versus $p_0 = .25$	Four rows on the $n = 400$ bed; locked per-row $\alpha = .01$	Only $p < .01$ clears; no row is relabelled from its paired McNemar verdict

Table 13: Pooled-to-fresh endpoint decision tree. The final elevation is post-data and is disclosed as such; it is not presented as the original primary endpoint.

Step	Status	Rule or observation	Consequence
1	Pre-data	Pooled $n = 1219$ primary; fresh $n = 819$ sensitivity. Direction fixed as <code>s_conf > rand_inject</code> ; one-sided 95% lower bound must exceed zero. Opposite pooled/fresh signs force UNDETERMINED.	No favorable sign selection
2	Registered data check	Between-pool magnitude heterogeneity: $\delta_{\text{diff}} = 8.8263$ pp, $z = 2.4148$, two-sided $p = 0.0157$; robust $z = 2.3229$.	Reject a common-effect estimate
3	Pre-data branch	Both pool effects are positive (fresh +12.58 pp, pilot +3.75 pp).	Sign-consistency branch passes
4	Post-data	Elevate the disjoint, no-reuse fresh pool; retain pooled +9.68 pp only as a weighted disclosure.	Fresh becomes the secondary answer-level result

Table 14: Anonymous chronology of the load-bearing decisions. Pre-data and post-data steps are kept distinct.

Status	Decision	Consequence
Pre-data	Pooled $n = 1219$ primary; fresh $n = 819$ sensitivity; opposite signs force UNDETERMINED.	Prevents choosing a favorable sign.
Pre-data	One-sided direction and heterogeneity check fixed before the fresh outcome.	Separates direction from pool compatibility.
Registered data check	Fresh and reused-pilot magnitudes differ ($z = 2.4148$, $p = 0.0157$).	Common-effect pooling is rejected.
Post-data, disclosed	Fresh disjoint pool elevated; pooled estimate retained as secondary disclosure.	The fresh priority is not presented as original.
Pre-data extension	Fixed identity order compared with <code>s_conf</code> on the fresh pool as a single load-bearing pair.	Tests whether confidence ordering is necessary or uniquely optimal.
Post-data reading	Fixed order beats <code>s_conf</code> .	Supports only a structured order boundary, with format/structure caveats.

E HISTORICAL DIAGNOSTICS AND BUDGET LEDGER

Table 15: Historical $n = 64$ estimator-budget scan (upper rows, exploratory) and the corrected $n = 400$ confirmatory scan (lower block, ascending budget). The historical $n_{\text{mc}} = 128$ row remains underpowered under the separate 25 pp exploratory tier. On the $n = 400$ bed, the pre-registered draw-matched $n_{\text{mc},i} = L_i$, total-NFE-fair $n_{\text{mc},i} = 2L_i$, and high-budget $n_{\text{mc}} = 128$ contrasts share the same 4 pp promotion bar and all three pass; the last is the conservative main read. Intervals follow Table 3: question bootstrap for the two $n_{\text{mc},i}$ rows, registered paired SE elsewhere.

Reference	Δ	m	Exact p	95% CI
Single draw ($n=64$)	+25.00 pp	30	0.0052	[8.23, 41.77]
$n_{\text{mc}} = 128$ ($n=64$)	+10.94 pp	29	0.2649	[-5.56, +27.43]
Per-candidate draw-matched ($n=64$)	+25.00 pp	32	0.0070	[7.68, 42.32]
<i>Corrected $n = 400$ confirmatory run</i>				
Single draw ($n=400$)	+14.00 pp	170	2.1×10^{-5}	[7.61, 20.39]
Draw-matched $n_{\text{mc},i} = L_i$ ($n=400$)	+13.00 pp	160	4.8×10^{-5}	[7.00, 19.00]
Total-NFE-fair $n_{\text{mc},i} = 2L_i$ ($n=400$)	+13.00 pp	166	6.6×10^{-5}	[6.75, 19.25]
$n_{\text{mc}} = 128$ ($n=400$)	+9.00 pp	128	0.0019	[3.46, 14.54]




Figure 5: Five-arm top-1 accuracy on the $n = 400$ pool, regenerated from frozen item-level records; exact counts are printed above each bar. Whiskers are descriptive Wilson 95% intervals (Wilson, 1927) on each arm’s *marginal* accuracy. Registered comparisons are paired and are reported in Table 15, not read off these unpaired intervals.

Table 16: Descriptive exact-count flip rates under the frozen 25% distractor-score shrink. The shrink rescales each protocol’s *own* score scale; no cross-scorer magnitude or inferential claim is attached.

Pool	s_conf	rand_inject
$n = 48$	63/144 (43.75%)	87/144 (60.42%)
$n = 64$	90/192 (46.88%)	105/192 (54.69%)




Figure 6: Tie-dropping Kendall τ distributions on the historical $n = 48$ and $n = 64$ pools. Each panel explicitly overlays the random-order self- τ floor (0.2361 ± 0.451) and the deterministic self-agreement anchor (1.000); both are descriptive calibration aids, not inferential intervals.

Table 17: Rank-agreement diagnostics. Deterministic rows are reproducibility anchors; random-order self-agreement is the stochastic comparator floor.

Diagnostic	Scope	Value	Interpretation
s_conf-rand_inject tie-drop	$n = 48$	0.2514	Full grid 0.2500; 8 tied items
s_conf-rand_inject tie-drop	$n = 64$	0.3323	Full grid 0.3255; 9 tied items
Random-order self- τ	Shared 48	0.2361 ± 0.451	Stochastic floor
Deterministic self- τ	Shared 48	1.000	Reproducibility anchor
s_conf-fixed_inject	$n = 400/\text{fresh } 819$	0.48 / 0.47	Cross-protocol anchor

Table 18: Candidate-length distributions used by draw matching and strict NFE sensitivity. The first three rows count all four nominal candidates per question. The final row counts only the unique designated candidate in each of the 819 fresh questions; this is the distribution used by the closed form in Table 10.

Pool	Items	$L = 3$	$L = 4$	$L = 5$	$L = 6$	$L = 7$	$L = 8$	$L = 9$	Mean L
Historical	64	39	120	72	17	8	0	0	4.355
Primary	400	262	751	424	118	30	15	0	4.343
Fresh sensitivity	819	364	1237	1010	390	167	101	7	4.722
Fresh designated	819	163	400	173	48	20	13	2	4.278

Table 19: Per-candidate budget ledger, in ascending reference budget. Scoring draws and total forward calls are deliberately separated; schedule selection adds another L forwards to the confidence and anti-confidence arms. The reference-arm rows carry the same budget labels as Table 3, so every row of that budget curve closes back to a ledger line: single draw, $n_{\text{mc},i} = L_i$ (nominal draw match only), $n_{\text{mc},i} = 2L_i$ (total-NFE match), and $n_{\text{mc}} = 128$.

Protocol	Scoring draws	Selection forwards	Total forwards	Reading
s_conf / anti_inject	L	L	$2L$	Model-derived ordering plus injection
rand_inject / fixed_inject	L	0	L	Seeded-random or fixed ordering plus injection
mc_uniform single	1	0	1	One masked-subset estimator draw
mc_uniform draw-matched	L	0	L	Matches nominal injection scoring draws only
$n_{\text{mc},i} = L_i$				
mc_uniform total-NFE-fair	$2L$	0	$2L$	Matched to s_conf’s total NFE (scoring plus selection); the decisive row of Table 3
$n_{\text{mc},i} = 2L_i$				
mc_uniform high budget	128	0	128	High-budget point on the historical pilot and on the corrected $n=400$ run

F CONTRIBUTION AND POWER DETAILS

Table 20: Answer-accuracy contribution ledger. The $n = 400$ total-NFE-fair $2L/2L$ row is the decisive equalization and the draw-matched $n_{mc,i} = L_i$ row is the pre-registered primary endpoint; the four $n = 400$ rows form the reported budget curve and the high-budget endpoint is the conservative main read. “Per-candidate draw-matched” sets the reference estimator budget at $n_{mc,i} = L_i$ per candidate. The $n = 64$ row alone retains the separate historical 25 pp exploratory tier.

Reference point	n	$\hat{\Delta}$	m	Exact p	95% CI	Role
Single draw	400	+14.0 pp	170	2.1×10^{-5}	[7.61, 20.39]	Curve point
Draw-matched $n_{mc,i} = L_i$	400	+13.0 pp	160	4.8×10^{-5}	[7.00, 19.00]	Pre-registered primary
Total-NFE-fair $n_{mc,i} = 2L_i$	400	+13.0 pp	166	6.6×10^{-5}	[6.75, 19.25]	NFE-fair decisive
$n_{mc} = 128$	400	+9.0 pp	128	0.0019	[3.46, 14.54]	Conservative main read
Per-candidate draw-matched	64	+25.0 pp	32	0.007	[7.68, 42.32]	Exploratory (25pp tier)

G POWER AND SAMPLE-SIZE LEDGER

The decisive total-NFE-fair $2L/2L$ contrast clears the registered Holm threshold and shared promotion bar; the historical small-pool draw-matched result is detected with only a slim margin over its detectable magnitude, while the historical small-pool high-budget arm is underpowered and remains UNDETERMINED (not zero) under its exploratory tier. The corrected $n = 400$ high-budget read is determined and clears the same promotion bar as the decisive $2L/2L$ row (Table 20). Rank agreement and fixed perturbation overlays are descriptive diagnostics rather than hypothesis tests.

Table 21: Headline robustness-gate ledger. The frozen rerun is within its abort tolerance, and the leave-one-out check is a nondegenerate PASS; the hostile re-adjudication counts state the boundary at which significance or direction would fail.

Diagnostic	Recorded quantity	Reading
Frozen rerun	$n=400$; s_conf accuracy 42.75%; deviation 0.0 pp; abort tolerance 1 pp	Within abort tolerance
Hostile $b \rightarrow c$ re-adjudication	166 discordants; one change moves the discordant proportion by 0.60 pp and $\hat{\Delta}$ by $2/n=0.50$ pp; 12 changes lose significance and 26 reverse the sign	Change-count boundary
Three-bucket rule	$0 < infl \leq margin$: PASS; $infl > 2/n$: FAIL; zero influence: degenerate	Prespecified categorization
Leave-one-out	Influence 0.283 pp on the 400-item bed; $1/n=0.25$ pp; margin/influence ≈ 21	Nondegenerate PASS

H ANONYMOUS REPRODUCTION CHECKLIST

Table 22: Stable evidence-ID manifest for the quantitative statements exposed in the manuscript. Each ID names an immutable object in the anonymous artifact bundle without exposing a local machine path, execution identifier, or revision token.

ID	Bundle object	Scope	Independent check
E10	Reference-bed single-draw paired record	$n = 400$; sconf 171, mc_uniform 115	Recompute accuracy, discordants, exact McNemar, and registered SE
E11	Reference-bed matched- L paired record	$n = 400$; sconf 171, mc_uniform 119; $b/c = 106/54$	Recompute question bootstrap and Holm entry
E12	Reference-bed $2L$ and 128-draw paired records	$n = 400$; reference 119/135; $b/c = 109/57$ and $82/46$	Recompute both intervals, exact tests, and total-forward labels
E13	Fresh confidence/seeded-random paired record	$n = 819$ persisted schedule	Replay schedules and recompute fresh, pilot, and pooled disclosures
E14	Fresh fixed-order paired record	$n = 819$ plus construction and position strata	Recompute the primary pair and all disclosed strata
E15	Reference chance-calibration ledger	Four $n = 400$ reference rows; exact two-sided test versus .25	Recompute counts, exact p values, and locked $\alpha = .01$ verdicts

Table 23: Components included in the anonymous reproduction bundle. Each row has an immutable manifest entry and a one-command validation check.

Component	Contents	Independent check
Sampling frames	Question identifiers, split membership, option count, and reference-option index for every pool	Verify sample sizes, disjointness, and four-versus five-option chance levels
Distractor construction	Deterministic rules for generating, deduplicating, and shuffling numerical distractors	Rebuild each choice set and compare all candidate strings
Prompt template	Exact instruction, question, choices, answer suffix, and mask placement	Byte-compare rendered prompts before decoding
Protocol implementation	Confidence, anti-confidence, seeded-random, fixed, and random-mask schedules; tie policy; seed policy and persisted order traces	Run sampler self-checks, replay recorded schedules, and detect the documented cross-process seed limitation
Judge and extraction	Top-1 argmin rule, tie handling, answer parser, and invalid-output policy; no partial credit	Recompute predictions and paired correctness from score vectors
Statistics	Exact McNemar, paired intervals, Holm families, binomial chance tests, heterogeneity check, and SESOI formula	Recompute every table cell from the item-level records
Figures	Deterministic regeneration scripts plus the plotted data payloads they emit	Regenerate quantitative PDFs and byte-compare their numerical payloads
Manifest	Content hashes, environment specification, and command sequence	Validate dependencies before building the canonical manuscript

Table 23 is the public traceability surface. Internal workflow labels, machine paths, task identifiers, and revision tokens are intentionally absent from the double-blind manuscript while remaining available to the artifact auditor.