

Q-CoMem: QUANTIZED SPLIT REPLAY FOR MEMORY-CONSTRAINED HYBRID LLMs

Anonymous authors

Paper under double-blind review

ABSTRACT

Repeated queries over stable documents create a capacity problem for memory-constrained inference: exact prefix caching minimizes warm latency but retains state across all layers, whereas dense recomputation retains nothing but repeatedly processes the document. We present Q-CoMem, a capacity-first extension of CoMem split replay for hybrid LLMs. An offline Write runs the document through a split depth, then group-quantizes the boundary residual together with the lower-layer KV, convolution, and recurrent state required for later queries. Online Read dequantizes and forks that document state, keeps mutable request state private, replays the query through the lower layers, and reconstructs the suffix. A fail-closed FORKAUDIT trace validates immutability, copy-on-write, recurrent rebinding, and call provenance because output equality alone can miss state aliasing.

On an archival 60-item Qasper/2WikiMQA validation cohort with Qwen3.5-35B-A3B and H20-3e, a frozen Q4/Q4/Q8 policy reduces mean retained tensor payload from 136.235 to 9.661 MiB/document ($14.10\times$ smaller). Its mean F1 (0–100) is 54.24 versus 54.62 for split Q16, a paired difference of -0.39 points with a 95% bootstrap interval of $[-2.04, 1.08]$; this is a selected trade-off point among the evaluated split-replay policies, not a statistical-equivalence claim. A separate systems execution on the eight-item indices-6–9 subset records TTFT: full-prefix Q16 is faster (0.163 versus about 0.673 s) and has higher generated-token throughput (13.38 versus 7.15–7.19 tokens/s), while Q-CoMem TPOT is within 2.5% of its 54.11 ms value. Thus the supported contribution is capacity reduction with an explicit online cost, not TTFT acceleration. Separately, shared-document KV lowers the post-priming allocated delta from 4.901 to 2.229 GiB at $N = 32$, and a fail-closed ownership audit exposes a historical alias invisible to output equality. We do not claim lower total process memory, measured recall, edge-device performance, or broad quality preservation.

1 INTRODUCTION

Repeated queries over stable documents are common in document QA, agents, tree search, and best-of- N decoding. Reusing a document prefix avoids repeated prefill, but exact caching retains per-document state at every model layer. On a memory-constrained deployment, that state competes with model weights, request workspace, and other resident documents. The resulting question is therefore capacity-first: how much reusable document state must remain resident for later queries, and what latency and quality costs buy a smaller footprint?

Dense recomputation and exact prefix caching form two endpoints. The former stores no document state but repeatedly processes the prefix; the latter gives low warm TTFT but retains the complete prefix cache. CoMem instead writes a document to an intermediate depth and reuses the resulting per-token residual (Liu et al., 2026a;b). Hybrid LLMs make this trade-off harder: lower full-attention layers also need KV, lower linear layers need convolution and recurrent state, and those mutable states cannot be shared blindly across requests (Yang et al., 2025; Lenz et al., 2025).

We introduce Q-CoMem, a quantized hybrid split-replay design. Its offline Write persists the split residual and all lower-layer state needed to replay a future query, using state-type-aware low-bit packing. Its online Read dequantizes one document entry, shares only immutable document tensors, creates request-local mutable state, and reconstructs the upper suffix. This saves retained bytes by




Figure 1: Q-COMEM Write/Read pipeline. Offline Write runs the document to split depth j and packs the complete retained state: boundary residual, lower full-attention KV, and lower convolution/recurrent state. Online Read dequantizes the entry, shares only immutable document tensors, creates a request-local mutable fork through COW/rebinding, replays the query through the lower layers, and reconstructs the suffix. The resulting design trades online work for a smaller retained document entry; FORKAUDIT checks the ownership boundary independently of output equality.

spending suffix computation; it is not intended to dominate exact caching on TTFT. Because equal outputs can still hide a mutable-state alias, we pair the design with FORKAUDIT, a phase-aware ownership trace that checks immutability, copy-on-write (COW), recurrent rebinding, and selected call provenance.

Our contributions follow this memory–latency–quality logic:

- **Complete quantized hybrid split state.** Q-COMEM stores the depth- j residual together with the necessary lower full-attention KV and lower convolution/recurrent state. Real groupwise Q4/Q8 packing and a state-type/per-layer bit assignment reduce the whole retained entry rather than counting the residual alone.
- **Capacity–latency–quality measurement.** On 60 archived Qasper/2WikiMQA validation items, the frozen Q4/Q4/Q8 point is $14.10\times$ smaller than full-prefix Q16 and changes mean F1 by -0.39 points relative to split Q16. A separate registered systems execution on the eight-item indices-6–9 subset reports TTFT, TPOT, and generated-token throughput with the adverse TTFT trade-off visible. Both tables keep retained payload separate from allocator and process memory, and leave unmeasured Recall cells blank.
- **Ownership-safe reuse.** The seven-target FORKAUDIT contract couples semantic relations to phase-indexed storage evidence across 96 ownership configurations. A historical alias corrupts the persistent base while output and terminal request state remain exact; the repaired path is storage-clean in 8/8 cells. At $N = 32$, shared-document KV separately reduces the post-priming allocator delta by 54.5% relative to full-copy KV.

The quality and latency evidence shares the reported Qwen3.5/PyTorch/ Transformers/H20 stack. The timing execution uses indices 6–9 from each dataset, hence an eight-item subset of the 60-item quality cohort, under a distinct protocol; the larger ownership study uses fixed sequential schedules. “Store” is retained tensor payload, not total VRAM, process memory, or admitted serving capacity. Section 4.8 states the full boundary.

2 RELATED WORK

Prior work spans depth reuse, exact prefix caching, cache compression, hybrid state, and relation-based testing. Q-CoMEM combines these threads around one specific objective: reduce the complete retained state of a hybrid split-replay entry while making request-local ownership auditable.

Split-depth reuse. CoMem persists one depth- j residual per document token and executes the upper suffix for later queries (Liu et al., 2026a;b). We build on that interface rather than claiming a new split-replay primitive. The added system problem is to quantize the residual *and* the lower KV/recurrent state required by a hybrid model, then expose the capacity–latency–quality trade-off under an explicit storage denominator.

Paged and prefix-aware inference. PagedAttention, Prompt Cache, RadixAttention, ChunkAttention, Hydragen, and Preble establish block sharing, prefix reuse, shared-prefix computation, and scheduling (Kwon et al., 2023; Gim et al., 2024; Zheng et al., 2024; Ye et al., 2024; Juravsky et al., 2024; Srivatsa et al., 2025). These systems are the low-TTFT reference point; Q-CoMEM deliberately trades additional suffix work for a smaller retained document entry.

Hybrid state and cache management. Gated Delta Networks maintain mutable compact state (Yang et al., 2025); Jamba, Marconi, and HYPIC address hybrid architectures or reusable hybrid prefixes (Lenz et al., 2025; Pan et al., 2025; Liu et al., 2026c), while Palu, SqueezeAttention, and HeadKV compress or budget KV (Chang et al., 2025; Wang et al., 2025; Fu et al., 2025). Q-CoMEM differs by jointly accounting for the split residual and the lower hybrid state. FORKAUDIT validates the immutable-document/mutable-request boundary rather than replacing these cache policies; Appendix Table 16 separates published numbers from our controls.

Evaluation and metamorphic testing. SCBench emphasizes cache lifecycles (Li et al., 2025); metamorphic testing, MT-DLComp, and CRADLE motivate preserved relations and independent checks (Chen et al., 2018; Xiao et al., 2022; Pham et al., 2019); DNNMem separates memory denominators (Gao et al., 2020). Appendix Table 11 gives our component-wise boundary.

3 METHODOLOGY

3.1 CAPACITY-FIRST OBJECTIVE AND STATE DECOMPOSITION

Let a stable document D be written once and queried repeatedly. A deployment must fit model weights, shared adapters, active workspace, retained document entries, and runtime state:

$$M_{\text{device}} \geq M_{\text{weights}} + M_{\text{shared}} + M_{\text{active}} + \sum_{D \in \mathcal{H}} S(D) + M_{\text{runtime}}, \quad (1)$$

where $\mathcal{H}$ is the hot document set. Q-CoMEM targets $S(D)$; it does not compress model weights or promise a smaller active workspace. We therefore optimize retained bytes subject to measured latency and quality, and report allocator measurements under a separate denominator while leaving process/NVML memory unmeasured.

For split depth j , the complete reusable hybrid state comprises the document boundary residual h_j^D , lower-layer full-attention KV, and lower-layer convolution/recurrent state. With quantizer Q_b , its retained payload is

$$S_{\text{QCoMem}}(D; j, \mathbf{b}) = |Q_{b_h}(h_j^D)| + \sum_{\ell < j, \ell \in \mathcal{A}} |Q_{b_\ell}(K_\ell^D, V_\ell^D)| + \sum_{\ell < j, \ell \in \mathcal{G}} |Q_{b_\ell}(C_\ell^D, G_\ell^D)|. \quad (2)$$

Counting only h_j^D would understate the deployed state, especially for a hybrid architecture.

3.2 QUANTIZED HYBRID SPLIT REPLAY

Offline Write. The document runs through layers $[0, j)$. We retain h_j^D and the lower attention and recurrent cache leaves needed to process a later query. Each floating tensor is partitioned into groups

of 64 values. For a b -bit group with minimum m and maximum u , we store

$$s = (u - m)/(2^b - 1), \quad q = \text{clip}(\text{round}((x - m)/s), 0, 2^b - 1), \quad (3)$$

with packed unsigned values and BF16 scale/bias metadata. Q16 is the BF16 unpacked reference. State types may use different widths; the evaluated Q4/Q4/Q8 point assigns Q4 to the residual and lower full-attention state and Q8 to lower linear state. The reported per-layer policy uses a Q4 residual, Q16 projections, and active-layer widths [8, 8, 4, 4, 8, 8, 8] selected on a disjoint four-item calibration set. We treat it as one frozen operating point, not as a general claim that layer-wise selection is optimal.

Online Read. For each query, the packed entry is dequantized into a logical document view. Immutable document tensors may be shared, whereas cache leaves that kernels can mutate are cloned or rebound into request-local storage. Query tokens traverse the lower layers using this state; document and query residual chunks then seed the suffix, which is executed through layers $[j, L]$ before incremental greedy decoding. Keeping the document and query chunks distinct is necessary for the Qwen3.5 hybrid recurrence. The design thus exchanges suffix reconstruction and dequantization for a smaller per-document retained entry.

3.3 OWNERSHIP-SAFE FORKS AND FAIL-CLOSED VALIDATION

For resident requests $r \in \{1, \dots, N\}$, let $\mathcal{A}$ and $\mathcal{G}$ denote full-attention and recurrent layers. A logical fork is

$$\mathcal{F}_r(D) = (\{K_\ell^D, K_\ell^r\}_{\ell \in \mathcal{A}}, \{G_\ell^D, G_\ell^r\}_{\ell \in \mathcal{G}}), \quad (4)$$

where document state K_ℓ^D, G_ℓ^D is immutable and may be shared, whereas append reservations K_ℓ^r and mutable recurrent state G_ℓ^r are request-local. A borrowed recurrent base is read-only at setup and must rebind to private storage when the registered transition first mutates it; a partial KV tail is copied before append.

FORKAUDIT validates this lifecycle with seven targets: frozen identity, prefix immutability, private ownership, tail-safe COW, bounded host-side call provenance, cross-arm semantic equivalence, and cross- N prefix consistency. For case specification Σ , execution trace τ , mandatory slots $M_i(\Sigma)$, and replay predicates Φ_i ,

$$\text{Pass}_i(\tau) \iff [\text{Coverage}_i(\tau) = \text{complete}] \wedge \text{Bind}_\Sigma(\tau) \wedge \bigwedge_{\phi \in \Phi_i} \phi(\tau). \quad (5)$$

Coverage is therefore separate from verdict: a missing, duplicated, modified, or unbound mandatory receipt cannot silently pass. Targets 2–5 are pre-output storage/call obligations; targets 6–7 compare registered semantic observables and cannot substitute for ownership evidence.

This is an offline, non-adversarial regression contract, not security attestation. Capture/replay, the mandatory-slot schema, correct slot-to-live-tensor binding, and framework storage semantics form the trusted computing base (TCB). A pass establishes only that the registered predicates held at captured points under that TCB; it does not cover coherent omission/fabrication, transient writes between observations, compiled GDN, underlying ATen/CUDA, or device/driver execution.

3.4 CAPTURE, REPLAY, AND SENSITIVITY CHECKS

The workflow freezes code, model, data, schedule, and runtime roles; captures setup, transition, and final digests, ranges, calls, semantics, and allocator snapshots; then replays every applicable predicate with complete/partial/open coverage reported separately from pass/fail/not-evaluated. The selected-cell hook additionally binds the actual builder, direct clone lineage, and serializer rows at three phases, while the separate dispatch run binds the selected host-side attention artifact/configuration immediately before invocation and seals it after normal return. Appendix D and Table 21 give the full target-to-record matrix.

Sensitivity uses matched live mutations, designer–executor-separated controls, a dense CPU-FP32 attention sidecar, and a candidate-import-free NumPy GDN oracle at frozen operator boundaries. A historical pre-fix request preserved output and terminal request state while retaining an invalid mutable base alias; the trace failed at transition binding and base immutability, and the repaired path

passes in 8/8 cells. This motivates the audit component without turning the fixed controls into a population-rate claim.

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Quality cohort. The broadest Store-F1 panel is an archival, raw-first validation run on Qwen3.5-35B-A3B, PyTorch 2.11.0+cu129, Transformers 5.14.1, and eight H20-3e GPUs. It contains 60 paired items: Qasper and 2WikiMQA source indices 6–35, or 30 examples per dataset, under official LongBench-v1 prompts and a 4,096-token input cap. Qasper permits at most 128 generated tokens and 2WikiMQA at most 32. We compare dense recomputation, full-prefix Q16/BF16, split-replay Q16/BF16, frozen state-type Q4/Q4/Q8, same-memory per-layer mixed precision, and an aggressive mixed policy. Four calibration prompts (indices 4–5 from each dataset) precede and are disjoint from validation. The archived raw shards provide retained Store, predictions, references, and LongBench answer F1, but not timing or Recall.

Latency execution. A separate registered systems execution uses the same model family in BF16, PyTorch 2.11.0+cu129, and Transformers 5.14.1. Eight validation items comprise four Qasper and four 2WikiMQA examples at source indices 6–9. These items are a subset of the 60-item quality cohort, not an independent replication; the execution and aggregation protocol is separate and the two tables are not pooled. Each item uses one H20-3e, a 4,096-token input cap, greedy decoding, at most 32 generated tokens, and three timing repeats per configuration. It compares no-cache dense recomputation, full-prefix Q16/BF16 reuse, uniform Q8 and Q4 Q-CoMEM states, state-type Q4/Q4/Q8, and the frozen per-layer point. Metrics are retained Store per document, TTFT, time per output token (TPOT), generated-token throughput, and answer F1. Recall was not recorded.

Ownership-audit cohort. Primary feasibility uses one fixed Qwen3.5-35B-A3B configuration (Qwen Team, 2026a;b): ten vLLM-0.26 Triton full-attention layers with BF16 KV (the protocol’s “Q16” storage label, not quantization) and 128-token pages, plus thirty Transformers-torch GDN layers. Eight H20-3e ranks run one book each with batch-one sequential calls on one CUDA stream/rank. In the fresh formal rerun, each registered attention call binds the selected hashed Triton kernel-cache artifact and compile configuration before invocation and seals the receipt only after the original launcher returns normally on the same assigned GPU and stream. GDN calls bind the mutually exclusive frozen eager routes; each GDN state digest covers 60 tensors.

Each distinct PG-19 book (Rae et al., 2020) supplies a 4,095-token document and 32-token registered transition. Four cells evaluate $N \in \{1, 8, 32\}$ and eight argmax outputs; the query plus seven fed-back outputs give 39 state-appended tokens and exercise COW on the 127-token final page. Appendix Table 5 separates all supporting cohorts. A distinct `tiiuae/Falcon-H1-0.5B-Base` track (Zuo et al., 2025) freezes Transformers 5.14.1 on H20 and the official `DynamicCache` reference, $N \in \{1, 2\}$, and persistent lossless-Q16 versus deep-materialized candidate arms. It tests exact relational transfer for that naive-path configuration, not performance, memory, or quality.

4.2 OWNERSHIP-AUDIT PROCEDURE

Each primary-factorial cell binds identity, rebuilds allocator and witness executions, records setup-transition-final ownership, and compares tokens, logits, logical KV, and GDN state. Replay adds fixed-axis projections, one dense-FP32 sidcar per rank, and matched live faults. A separate PDF-only designer freeze supplies five executor-separated pairs. Aggregation requires every byte binding. Appendix Table 6 states target strength.

The stronger hook wraps only the selected $N = 8$ shared-document/materialized-GDN ownership witness. Per rank, $60 = 30 \times 2$ denotes persistent layer/family coordinates, $480 = 8 \times 60$ direct clone destinations, and each phase contains the full $540 = (1 + 8) \times 60$ serializer-row universe with six owner-row anchors. The same schedule separately enumerates $12 = 3 \times 4$ clean-memory calls per rank (96 total). Zero selected-context overlaps means only that this context never entered an allocator cell, not that the wrapper was absent or perturbation-free.

Table 1: Retained-state–quality trade-off on the 60-item archival validation cohort. The cohort contains 30 Qasper and 30 2WikiMQA examples. TTFT and Recall are retained as target columns but are blank because this run did not record either metric.

Dataset: Qasper (30 items) + 2WikiMQA (30 items)						
Retained state	Store (MiB/doc)	vs. prefix	TTFT (s)	F1 (0–100)	Recall	Δ F1 vs. split Q16 [95% CI]
No-cache dense recompute	–	–		54.29		–0.33 [–4.61, 2.67]
Full-prefix Q16/BF16	136.235	1.00×		54.68		+0.06 [0.00, 0.18]
Q-CoMEM split Q16/BF16	34.683	3.93×		54.62		0.00 [0.00, 0.00]
Q-CoMEM frozen Q4/Q4/Q8	9.661	14.10×		54.24		–0.39 [–2.04, 1.08]
Q-CoMEM same-memory mixed	9.395	14.50×		53.36		–1.26 [–4.28, 0.98]
Q-CoMEM aggressive mixed	7.536	18.08×		49.19		–5.44 [–11.56, –0.17]

Store and F1 are means over the 60 items. Store is retained-document tensor payload, not process/NVML memory or total VRAM. Compression is relative to full-prefix Q16/BF16. F1 and its differences use the 0–100 scale; intervals are paired 10,000-resample bootstrap intervals against split Q16. The frozen policy uses Q4 residuals, Q4 lower attention state, and Q8 lower linear state. Its catastrophic-regression count is 0/60 under the registered per-item Δ F1-versus-dense ≤ -50 point rule (equivalently -0.5 on the $[0, 1]$ scale). An interval crossing zero is not evidence of statistical equivalence. Blank TTFT and Recall cells denote unmeasured metrics, not zeros.

4.3 CONTROLS, SCHEDULES, AND ORACLES

The primary factorial uses sequential round-major steps on one CUDA stream per rank. Full logits use shape/dtype-bound CPU-FP32 bytes; KV, GDN, bases, calls, storage, and allocator records use typed digests. Detached replay cannot reconstruct unarchived tensors or recapture the producer. A preregistered NumPy oracle checks four GDN transitions at global model-layer indices 0, 10, 20, 38 after native q/k normalization against four wrong recurrences.

Supporting protocols move GDN reconstruction to a spawn-created CUDA-IPC observer and record CUDA-event-bracketed model-call intervals around cancel, scrub, reclaim, and replacement. Appendix Table 5 fixes their geometry and authorization.

4.4 MEMORY DENOMINATORS

We separate logical state, pages, synchronized allocator bytes, and unmeasured service capacity. In both deployment tables, “Store” denotes retained tensor payload, but the accounting implementations differ: the 60-item panel sums the stored persistent components (packed residual data, scales, biases, and unique lower-cache tensor storage), whereas the eight-item execution measures the physical byte-range union owned by one retained entry. Neither denominator includes non-tensor/Python metadata, allocator pools, process/NVML memory, active workspace, or admission capacity, and their differently aggregated values are not pooled.

4.5 RETAINED STATE, LATENCY, AND ANSWER QUALITY

Table 1 reports the broadest current retained-state–quality cohort. Moving from full-prefix Q16 to split Q16 first reduces mean Store from 136.235 to 34.683 MiB/document (3.93× smaller). Quantizing the complete split state with the frozen Q4/Q4/Q8 policy reduces it further to 9.661 MiB/document: 14.10× smaller than the full-prefix reference, or a 92.91% reduction. Mean F1 is 54.24 versus 54.62 for split Q16, with a paired difference of -0.39 points and 95% bootstrap interval $[-2.04, 1.08]$. No item crosses the registered catastrophic threshold. This is a selected trade-off point among the evaluated split-replay policies; the interval crossing zero does not prove equivalence or noninferiority.

The same-memory per-layer policy saves only 2.75% Store relative to the frozen state-type policy, while its mean-F1 difference is more negative and one of 60 items crosses the catastrophic threshold. The aggressive policy reaches 18.08× compression but its interval against split Q16 lies below zero and five of 60 items cross that threshold. We therefore use frozen Q4/Q4/Q8 as the headline operating point and treat both mixed policies as ablations, not improvements.

Table 2 provides the separate timing characterization on the eight-item indices-6–9 subset. On these eight items, Q8 retains 15.89 rather than 140.34 MiB/document (88.68% less) at the same measured mean F1 of 39.137. The calibrated per-layer point retains 9.74 MiB/document (93.06% less) with a mean-F1 change of -0.022 points. Uniform Q4 reaches 8.41 MiB but loses 3.047 F1 points, making it an aggressive-bit failure endpoint rather than the headline operating point.

Table 2: Capacity–latency–quality trade-off from a separate execution on the registered eight-item Qasper/2WikiMQA validation slice (source indices 6–9 from each dataset), which is a subset of the 60-item panel rather than an independent replication. All rows use Qwen3.5-35B-A3B, Transformers 5.14.1, one H20-3e per item, a 4,096-token cap, greedy decoding, at most 32 generated tokens, and three timing repeats.

<i>Dataset: Qasper (4 items) + 2WikiMQA (4 items)</i>							
Retained state	Store (MiB/doc)	Reduction	TTFT (s)	TPOT (ms)	tok/s	F1 (0–100)	Recall
No-cache dense recompute	0.00	–	0.649	648.75	1.36	39.42	
Full-prefix Q16/BF16	140.34	reference	0.163	54.11	13.38	39.14	
Q-CoMEM Q8	15.89	88.68%	0.673	55.44	7.15	39.14	
Q-CoMEM Q4/Q4/Q8	10.01	92.87%	0.674	54.97	7.19	39.01	
Q-CoMEM per-layer mixed	9.74	93.06%	0.673	55.07	7.19	39.11	
Q-CoMEM Q4	8.41	94.01%	0.674	55.02	7.19	36.09	

Store is median retained-document tensor payload: the physical byte-range union owned by one retained entry. It excludes non-tensor/Python metadata, allocator pools, process/NVML memory, active workspace, and admission capacity. Store and timing/throughput columns are medians; F1 is the mean over the eight workloads on the 0–100 LongBench scale. This execution is not pooled with the 60-item panel. Reductions are relative to full-prefix Q16. The per-layer policy uses a Q4 residual, Q16 projections, and layer bits [8, 8, 4, 4, 8, 8, 8] selected on a disjoint four-item calibration set. Mean-F1 deltas for Q8 and per-layer mixed are 0.000 and -0.022 points, respectively. The all-Q4 row is the aggressive failure endpoint, not a quality-preserving operating point. Recall was not measured in this cohort; its cells are deliberately left blank and no recall claim is made.

The latency result is not a speedup over exact caching. Full-prefix Q16 records 0.163 s TTFT, whereas the Q-CoMEM rows are approximately 0.673–0.674 s because they dequantize and rebuild the suffix. Their TPOT is within 2.5% of full-prefix (54.97–55.44 versus 54.11 ms), but median generated-token throughput falls from 13.38 to 7.15–7.19 tokens/s; dense recomputation has 648.75 ms TPOT. Thus the supported result is a retained-capacity trade-off with explicit online cost, not universal latency improvement.

4.6 AUDIT CONTROLS AND GOVERNANCE

Pre-output ledgers bind code, model, data, queries, and protocol. Nine matched live mutations, a pre-registered target-gate-suppression matrix, and five designer–executor-separated constructed controls use frozen loci, predicates, and precedence; they establish fixed-case sensitivity and localization, not a rate. A separate no-injection historical protocol compares fresh pre-fix, repaired-borrowed, and materialized $N = 2$ lanes for one previously observed alias mechanism. Appendix E records the detailed rows.

4.7 OWNERSHIP AND CORRECTNESS RESULTS

4.7.1 PRIMARY TRACE VALIDATION

All seven targets have complete mandatory-trace coverage and satisfy the registered pass criterion for the fresh frozen primary rerun. For target 5, host-side receipts for 209,920 attention calls record the selected hashed Triton artifact/configuration and autotune/no-autotuner decision immediately before the original Python launcher, then seal after normal return. Another 635,520 receipts record the mutually exclusive GDN eager route and functional cache rebind. The totals factor as $20,992 \times 10$ attention calls and $20,992 \times 30$ request-GDN calls plus 192×30 one-time prefill calls. They establish host-side selected launch/configuration provenance only, not device-side operator identity, execution, or completion, compiled GDN, underlying ATen/CUDA, or driver/device binaries.

4.7.2 OWNERSHIP EVIDENCE AND ALLOCATOR RESULT

Across $8 \times 3 \times 4 = 96$ configurations (192 rebuilt allocator/witness cells), tokens, full logits, final logical KV, final GDN, both fixed-axis projections, and all 288 adjacent-fan-out comparisons match exactly. These semantic relations are paired with physical ownership evidence: borrowed GDN starts as 60 immutable base aliases/request, materialized setup is disjoint, and completed requests in both setup arms rebind at the registered 32-token transition boundary to 60 request-private tensors disjoint from bases and completed peers. The evaluated ownership and immutability predicates pass at setup, transition, and final capture; the bounded host-side receipts also pass (Section 4.7.1). The $N = 1$ cells directly exercise request–base disjointness.

Selected-cell builder-to-receipt live binding. The V29 rerun is terminally complete and scientifically valid on all eight ranks. Its same-run selected witnesses yield $8 \times 3 \times 6 = 144$ selected-owner-row checks, $8 \times 3 \times 540 = 12,960$ full serializer/live-object rows, $8 \times 480 = 3,840$ direct clone edges, and $8 \times 3 = 24$ stable artifacts. All 96 separately enumerated clean-memory calls occur outside the active selected context. Other cells and the honest-process assumption remain outside this stronger binding.

Allocator reduction with canonical equality. With materialized GDN setup at $N = 32$, switching from full-copy to shared-document KV reduces final allocated delta from 4.901 to 2.229 GiB, a 54.5% reduction above the frozen post-priming baseline, while the registered observables are canonically equal across the compared ownership arms. Borrowed GDN changes allocation timing rather than the shared-KV end state: its corresponding full-copy and shared-document endpoints are 4.890 and 2.229 GiB, and both shared-document cells finish at 2.229 GiB. Appendix Table 15 reports all four endpoints under the distinct allocator denominator.

Falsification beyond output equality. All 44 clean captured-boundary numerical rows pass and all 44 seeded wrong-operator controls are rejected. Nine matched controls and five designer-executor-separated controls also pass, while every paired mutant is rejected at its frozen gate; token equality catches 0/5 of the semantic-complete mutants and exact logits catch 1/5. In the historical case, output and terminal request state remain exact in 8/8 cells despite persistent-base corruption, whereas the repaired path is storage-clean in 8/8. Appendix Tables 7–14 give the fixed-case rows and captured-boundary qualifications.

Supporting cohorts extend the lifecycle through cancellation/reclamation, derive 1,080 GDN observations and 96,660 relations from a source-distinct expected-slot census, and reproduce them in a second execution. A separate Falcon-H1 configuration matches its official Transformers reference on all registered token, logit, state, and ownership checks. These bounded checks do not supply additional Q-COMEM performance or quality evidence; protocol and transfer details remain in the appendix.

4.8 LIMITATIONS AND SCOPE

The quality panel contains 60 Qasper/2WikiMQA validation items on one Qwen3.5/H20 stack; the latency panel is a separate execution on its eight-item indices-6–9 subset, not an independent replication. The tables are not pooled. They support the stated bounded point estimates, not broad benchmark quality preservation. The 60-item run is archival raw-first evidence: hashes, raw predictions, F1, bootstrap intervals, and Store accounting replay, but the original launcher used a mutable code directory and retained neither a pre-run full-source ledger nor a full model-weight ledger. We therefore do not describe it as fresh source-frozen regeneration. The 60-item run has no timing measurements; neither cohort measured Recall, whose table cells remain blank.

H20 is not an edge device; memory-constrained local or edge use is the motivation, not an evaluated hardware claim. The calibrated per-layer row is one operating point selected from four disjoint items, and the current implementation has substantially higher TTFT than full-prefix reuse. We do not evaluate continuous batching, QPS, admission, eviction, or scheduler effects.

Store is retained tensor payload: component accounting in the 60-item panel and an entry-owned physical byte-range union in the eight-item execution. It excludes non-tensor/Python metadata, allocator pools, active workspace, process/NVML memory, and model weights; reducing Store does not by itself prove lower peak VRAM or serving capacity. The separate 54.5% result is a post-priming allocator delta from the ownership factorial, so the two reductions are not additive or interchangeable.

Within the TCB in Section 3, FORKAUDIT upgrades semantic equality to phase-indexed ownership evidence. Its selected-cell binding remains an honest same-process check: other cells and coordinates, framework/IPC semantics, compiled GDN, ATen/CUDA, and device/driver execution remain unattested. Constructed controls and one historical mechanism demonstrate fixed-case sensitivity, not population error rates. Appendix H gives the exhaustive boundary.

5 CONCLUSION

Q-COMEM reduces the complete retained hybrid split-replay entry by jointly packing the document residual and required lower KV/recurrent state. On the 60-item archival validation cohort, frozen Q4/Q4/Q8 is $14.10\times$ smaller than full-prefix Q16; its observed mean-F1 difference from split Q16 is -0.39 points with 95% interval $[-2.04, 1.08]$. The separate timing execution on the eight-item subset makes the cost explicit: exact prefix reuse remains much faster in TTFT and higher in generated-token throughput, while TPOT is within 2.5% on this eight-item slice. FORKAUDIT complements this capacity result by validating request-local ownership and exposing a historical silent alias. The claim is deliberately capacity-first: smaller retained per-document state at bounded measured quality, not measured Recall, lower total VRAM, or universally faster inference.

REPRODUCIBILITY STATEMENT

Claims map to registered evidence identifiers. The 60-item table maps to a raw-first archival package that rechecks 66 mirrored files, 48 shards, 360 item-level F1 values, 24 paired-bootstrap intervals, and Store accounting. The latency table is hash-bound to its eight workload shards and aggregate; manifest-first replay covers the primary ownership factorial and named supporting cohorts. Dispatch and live-binding results bind frozen archives to post-run ledgers, terminal manifests/trees, and audit mirrors. Appendix A reports replay counts, costs, and non-reexecuted boundaries.

REFERENCES

- Chi-Chih Chang, Wei-Cheng Lin, Chien-Yu Lin, Chong-Yan Chen, Yu-Fang Hu, Pei-Shuo Wang, Ning-Chi Huang, Luis Ceze, Mohamed Abdelfattah, and Kai-Chiang Wu. Palu: KV-cache compression with low-rank projection. In *International Conference on Learning Representations*, volume 2025, pp. 50222–50249, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/7da6e0e00702c60607a6ae05c802ef85-Abstract-Conference.html.
- Tsong Yueh Chen, Fei-Ching Kuo, Huai Liu, Pak-Lok Poon, Dave Towey, T. H. Tse, and Zhi Quan Zhou. Metamorphic testing: A review of challenges and opportunities. *ACM Computing Surveys*, 51(1):4:1–4:27, 2018. doi: 10.1145/3143561. URL <https://doi.org/10.1145/3143561>.
- Yu Fu, Zefan Cai, Abedelkadir Asi, Wayne Xiong, Yue Dong, and Wen Xiao. Not all heads matter: A head-level KV cache compression method with integrated retrieval and reasoning. In *International Conference on Learning Representations*, volume 2025, pp. 99269–99290, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/f649556471416b35e60ae0de7c1e3619-Abstract-Conference.html.
- Yanjie Gao, Yu Liu, Hongyu Zhang, Zhengxian Li, Yonghao Zhu, Haoxiang Lin, and Mao Yang. Estimating GPU memory consumption of deep learning models. In *Proceedings of the 28th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, pp. 1342–1352. Association for Computing Machinery, 2020. doi: 10.1145/3368089.3417050. URL <https://doi.org/10.1145/3368089.3417050>.
- In Gim, Guojun Chen, Seung-seob Lee, Nikhil Sarda, Anurag Khandelwal, and Lin Zhong. Prompt cache: Modular attention reuse for low-latency inference. In *Proceedings of Machine Learning and Systems*, volume 6, pp. 325–338, 2024. URL https://proceedings.mlsys.org/paper_files/paper/2024/hash/a66caa1703fe34705a4368c3014c1966-Abstract-Conference.html.
- Jordan Juravsky, Bradley Brown, Ryan Ehrlich, Daniel Y. Fu, Christopher Ré, and Azalia Mirhoseini. Hydragen: High-throughput LLM inference with shared prefixes. *arXiv preprint arXiv:2402.05099*, 2024. URL <https://arxiv.org/abs/2402.05099>.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems*

- Principles*, pp. 611–626. ACM, 2023. doi: 10.1145/3600006.3613165. URL <https://doi.org/10.1145/3600006.3613165>.
- Barak Lenz, Opher Lieber, Alan Arazi, Amir Bergman, Avshalom Manevich, Barak Peleg, Ben Aviram, Chen Almagor, Clara Fridman, Dan Padnos, Daniel Gissin, Daniel Jannai, Dor Muhl-gay, Dor Zimberg, Edden Gerber, Elad Dolev, Eran Krakovsky, Erez Sa, Erez Schwartz, Gal Cohen, Gal Shachaf, Haim Rozenblum, Hofit Bata, Ido Blass, Inbal Magar, Itay Dalmedigos, Jhonathan Osin, Julie Fadlon, Maria Rozman, Matan Danos, Michael Gokhman, Mor Zusman, Naama Gidron, Nir Ratner, Noam Gat, Noam Rozen, Oded Fried, Ohad Leshno, Omer Antverg, Omri Abend, Or Dagan, Orit Cohavi, Raz Alon, Ro’i Belson, Roi Cohen, Rom Gilad, Roman Glozman, Shahar Lev, Shai Shalev-Shwartz, Shaked Meïrom, Tal Delbari, Tal Ness, Tomer Asida, Tom Ben Gal, Tom Braude, Uriya Pumerantz, Joshua Cohen, Yonatan Belinkov, Yuval Globerson, Yuval Levy, and Yoav Shoham. Jamba: Hybrid transformer–mamba language mod-els. In *International Conference on Learning Representations*, volume 2025, pp. 67959–67984, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/a9ed43fa31dc8b4a7d7a673d713dcb5f-Abstract-Conference.html.
- Yucheng Li, Huiqiang Jiang, Qianhui Wu, Xufang Luo, Surin Ahn, Chengruidong Zhang, Amir Abdi, Dongsheng Li, Jianfeng Gao, Yuqing Yang, and Lili Qiu. SCBench: A KV cache-centric analysis of long-context methods. In *International Conference on Learning Representations*, volume 2025, pp. 66063–66093, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/a540b17fb2295c736d5afd6c507acf66-Abstract-Conference.html.
- Hanzuo Liu, Xuan Qi, Chunyu Liu, Haotian Zhong, Yulong Wang, Key, Rayying, Alex Lamb, and Mingyu Gao. CoMem: Reusing transformer depth across queries with persistent intermediate residuals. In *COLM 2026 Workshop on Efficient Reasoning*, 2026a.
- Hanzuo Liu, Xuan Qi, Chunyu Liu, Haotian Zhong, Yulong Wang, Rayying, Key, Alex Lamb, and Mingyu Gao. Understanding is done early: A depth division of labor in large language models and its use for unbounded-context memory. *arXiv preprint arXiv:2607.28263*, 2026b. doi: 10.48550/arXiv.2607.28263. URL <https://arxiv.org/abs/2607.28263>.
- Yifei Liu, Juntong Wu, Yang Liu, Junhao Hu, Minghao Li, Xiaoxu Chen, and Weihang Chen. HYPIC: Accelerating hybrid-attention LLM serving with position-independent caching. *arXiv preprint arXiv:2607.01299*, 2026c. doi: 10.48550/arXiv.2607.01299. URL <https://arxiv.org/abs/2607.01299>.
- Rui Pan, Zhuang Wang, Zhen Jia, Can Karakus, Luca Zancato, Tri Dao, Yida Wang, and Ravi Netravali. Marconi: Prefix caching for the era of hybrid LLMs. In *Proceedings of Machine Learning and Systems*, volume 7. MLSys, 2025. URL https://proceedings.mlsys.org/paper_files/paper/2025/hash/7c180af017258d239bac6248d1eb26ac-Abstract-Conference.html.
- Hung Viet Pham, Thibaud Lutellier, Weizhen Qi, and Lin Tan. CRADLE: Cross-backend validation to detect and localize bugs in deep learning libraries. In *2019 IEEE/ACM 41st International Conference on Software Engineering (ICSE)*, pp. 1027–1038. IEEE, 2019. doi: 10.1109/ICSE.2019.00107. URL <https://doi.org/10.1109/ICSE.2019.00107>.
- Qwen Team. Qwen3.5: Towards native multimodal agents. Official Qwen technical blog, February 2026a. URL <https://qwen.ai/blog?id=qwen3.5>.
- Qwen Team. Qwen3.5-35B-A3B-Base model card. Hugging Face model repository, 2026b. URL <https://huggingface.co/Qwen/Qwen3.5-35B-A3B-Base/blob/main/README.md>. Accessed 2026-08-16.
- Jack W. Rae, Anna Potapenko, Siddhant M. Jayakumar, Chloe Hillier, and Timothy P. Lillicrap. Compressive transformers for long-range sequence modelling. In *International Conference on Learning Representations*, 2020. URL https://iclr.cc/virtual_2020/poster_SylKikSYDH.html.

- Vikranth Srivatsa, Zijian He, Reyna Abhyankar, Dongming Li, and Yiyang Zhang. Preble: Efficient distributed prompt scheduling for LLM serving. In *International Conference on Learning Representations*, volume 2025, pp. 37057–37082, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/5bc342f48de8264779952fac378f96dc-Abstract-Conference.html.
- Zihao Wang, Bin Cui, and Shaoduo Gan. SqueezeAttention: 2d management of KV-cache in LLM inference via layer-wise optimal budget. In *International Conference on Learning Representations*, volume 2025, pp. 23593–23606, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/3b0a8df568ec496a717566a7f8158aaa-Abstract-Conference.html.
- Dongwei Xiao, Zhibo Liu, Yuanyuan Yuan, Qi Pang, and Shuai Wang. Metamorphic testing of deep learning compilers. *Proceedings of the ACM on Measurement and Analysis of Computing Systems*, 6(1):1–28, 2022. doi: 10.1145/3508035. URL <https://doi.org/10.1145/3508035>.
- Songlin Yang, Jan Kautz, and Ali Hatamizadeh. Gated delta networks: Improving Mamba2 with delta rule. In *International Conference on Learning Representations*, volume 2025, pp. 29687–29707, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/4904fad153f6434a7bcf04465d4be2cc-Abstract-Conference.html.
- Lu Ye, Ze Tao, Yong Huang, and Yang Li. ChunkAttention: Efficient self-attention with prefix-aware KV cache and two-phase partition. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 11608–11620, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.623. URL <https://aclanthology.org/2024.acl-long.623/>.
- Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying Sheng. SGLang: Efficient execution of structured language model programs. In *Advances in Neural Information Processing Systems*, volume 37, pp. 62557–62583. Curran Associates, Inc., 2024. doi: 10.52202/079017-2000. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/724be4472168f31balc9ac630f15dec8-Abstract-Conference.html.
- Jingwei Zuo, Maksim Velikanov, Ilyas Chahed, Younes Belkada, Dhia Eddine Rhayem, Guillaume Kunsch, Hakim Hacid, Hamza Yous, Brahim Farhat, Ibrahim Khadraoui, et al. Falcon-H1: A family of hybrid-head language models redefining efficiency and performance. *arXiv preprint arXiv:2507.22448*, 2025. doi: 10.48550/arXiv.2507.22448. URL <https://arxiv.org/abs/2507.22448>.

A DETAILED REPRODUCIBILITY AND ARTIFACT SCOPE

The 60-item quality package preserves 66 authoritative mirrored files from QS Job 234340 / Trial 1830116, including 48 configuration-by-rank shards. Its raw-first verifier recomputes all 360 item-level F1 values, 24 paired-bootstrap intervals, Store component accounting, and the archived aggregate. The data hash and policy receipt are bound and calibration precedes validation, but the original execution did not freeze a complete source or model-weight ledger; offline outcome replay therefore does not imply source-frozen GPU regeneration.

Manifest-first CPU replay covers the primary-factorial package and named replayable supporting cohorts, not every contextual table. The primary package rechecks 536 artifacts, 96 cells, 288 cross- N comparisons, eight attention rows, and nine fault pairs. The target-gate-suppression package rechecks 18 cases and 24 FP32 sidecars; its disclosed wrapper changes only an inherited execution-identifier source and preserves every candidate byte. The independent-census package rechecks six IPC captures, 1,080 receiver rows, 96,660 relations, and omission/duplication/relabel controls against a preproducer 180-slot expectation. The dual-producer package binds two fresh serial executions of the same frozen producer implementation and four distinct observers, and rechecks exact zero-tolerance equality of the 1,080 observations and 96,660 relation labels per execution. The bounded two-stream replay recomputes four full-logit comparisons and both interval overlaps.

The fresh compiled-dispatch rerun binds a frozen source archive to a passing eight-rank formal aggregate, eight rank receipts, 209,920 attention calls, 635,520 GDN calls, and 224 rejected post-run negative-control decisions. The local mirror preserves the primary and formal aggregates, runtime preflight, raw/scientific ledgers, and both terminal ledgers; it does not duplicate the model view or every raw shard. The receipt closes target 5 only at the declared honest-process fixed-stack boundary: normal launcher return is not a device-side completion attestation, and compiled GDN, underlying ATen/CUDA identity, and driver/device binaries remain outside scope.

The separate live-binding rerun binds a different frozen archive to eight scientifically valid rank shards and a terminally complete formal tree. Its selected-cell hook validates 144 selected rows, 12,960 full live-storage rows, 3,840 direct clone edges, and 24 stable phase artifacts; a separately enumerated schedule counter sees all 96 clean-memory calls with zero selected-context overlaps. The post-run audit also rejects 184 formal controls and verifies 31 bytecode files, 13 descendant directories, and the 1,367-node final tree. This supports only the selected $N = 8$ shared-document/materialized-GDN cell and six preregistered owner-addressed rows per phase under an honest same-process runtime.

Within each selected witness call, the verifier (1) enumerates 60 persistent sources before the builder, (2) records the actual builder’s persistent-rooted Torch-dispatch clone edges, (3) exhaustively matches 480 destinations and passes a one-use capability bound to the exact group and verifier, (4) rereads all 540 live tensors against the production serializer at three phases, and (5) publishes hash-bound phase paths for terminal reread. Any count, object, descriptor, lineage, phase-order, or path mismatch aborts before acceptance. Trusted components are hook placement, Python/Torch-dispatch and tensor/storage semantics, and honest process scheduling; group identity, object continuity, lineage cardinality, serializer fields, and publication paths are checked.

The designer–executor-separated package binds the PDF-only author freeze, a lifecycle-amended executor, five clean/mutant pairs, detached clean and pair replays, and 208 terminal files. All five clean gates and expected-first classifications pass; the pre-injection amendment was frozen before all five ranks were rerun in a new output root.

The historical-regression package binds the pre-discovery chronology, frozen pre-fix/repair sources, outcome-blind protocol and resource amendment, eight three-lane rank artifacts, 24 FP32-logit sidecars, and 24 fresh nonces. A fresh candidate-import-free replay of every rank and the aggregate is byte-identical to the archived output.

The expanded numerical-sweep package binds 272 sidecars totaling 140,199,936 bytes and rechecks 20 attention and 24 GDN clean rows plus 44 controls. The Falcon package compares candidate arms against a separately executed official Transformers reference on 8 ranks and detached-replays the aggregate. HYPIC binds 4,404 timing artifacts and hash-replays all 16 Store cells. No package enumerates OS/driver allocations; every conclusion remains bounded to its declared configuration. Appendix I maps exact artifacts and boundaries.

On one Apple Mac16,8 CPU host, three consecutive fixed-order replays without a cache flush took a median 100.146 s (99.248–110.526 s) for the core package and 105.717 s (104.846–119.566 s) for the serial profile that additionally ran five locally complete supporting verifiers; the paired within-profile increment was 5.598 s (5.571–9.039 s). The canonical core distribution has 630 files and 892,284,156 logical bytes (850.95 MiB); 536 raw traces occupy 888,785,811 bytes, or 99.62% of manifest-payload bytes. These are local CPU replay and logical-storage measurements, not H2O capture cost, GPU perturbation, cold-start latency, online inference cost, or a matched uninstrumented delta.

B ETHICS STATEMENT

This work introduces no human-subject data collection. It uses public PG-19 text for the trace-validation protocols and a fixed 60-item Qasper/2WikiMQA validation cohort; the latency measurements use its eight-item indices-6–9 subset in a separate execution. It does not assess model safety, bias, privacy, or downstream deployment quality. Memory efficiency can lower deployment cost but may also increase the scale of model use; those broader effects are outside this controlled systems study.

C PROTOCOL GEOMETRY




Figure 2: FORKAUDIT ownership factorial. Four KV-by-GDN setup cells share immutable document state, keep append tails private, and rebind borrowed recurrent state at the registered transition. The audit rail records lifecycle, storage, call/semantic, and fault-to-gate receipts. This schematic details the validation component; it is not experimental evidence.

Table 3: Frozen primary-factorial execution geometry. The eight ranks provide distinct books and query banks, not repeated stochastic measurements.

Component	Frozen value
Model	Qwen3.5-35B-A3B; 10 attention + 30 GDN layers
Hardware	8 × H20-3e; one book/rank
Software	Python 3.11.13; PyTorch 2.11.0+cu129; Transformers 5.14.1; vLLM 0.26.0+cu129; Triton 3.6.0; CUDA 12.9
KV backend	vLLM fused attention; BF16/Q16; page 128
Document / query	4,095 / 32 tokens; PG-19 train only
Generation	8 output tokens; 39 state-appended input tokens (32 query + 7 generated-token feedback)
Resident counts	1, 8, 32; four KV-by-GDN ownership cells
Primary schedule	all requests resident; batch-one sequential round-major calls on one CUDA stream/rank
Validation	96 factor configurations; separate memory/witness cells
Fault campaigns	$N = 2$; the primary all-gates-on campaign has nine live mutants with rebuilt controls; the target-gate-suppression campaign has nine fresh clean plus nine suppressed-gate cases on 8 H20s; the designer-executor-separated campaign has five new clean/mutant pairs on 5 H20s
Historical regression	one previously encountered one-token alias defect; 8 ranks × 3 fresh $N = 2$ lanes; 3 archived-coordinate reproductions + 5 additional frozen inputs; no mutation

Table 4: Meaning of the V29 selected-cell evidence units. Counts are per rank unless marked total; all belong to the same formal run.

Unit	Count	Check and authorization
Persistent sources	60	independently enumerated before the actual builder; exact object, storage, descriptor, and content remain fixed
Selected anchors	6/phase	preregistered persistent/request semantic rows; request-owned subset consumes the group-bound lineage capability
Full live/serializer rows	540/phase	(1 persistent + 8 requests) $\times$ 60 rows; live objects are re-enumerated and every serialized field, role, interval, content digest, and phase relation is checked
Builder clone edges	480	8×60 materialized destinations exhaust the captured persistent-rooted ledger and are direct <code>aten.clone</code> outputs
Phase artifacts	3	setup, post-transition, and post-generation products are hash-bound, published to stable rank paths, and terminally reread
Clean-memory context separation	12 (96 total)	separate allocator-cell calls are all observed; zero selected-context overlaps mean the selected context did not enter them, not that the wrapper was absent or perturbation-free

With field order (owner, request, layer, family, state), the six anchors are (Pers, $\emptyset$, 0, rec, 0), (Req, 0, 0, conv, 0), (Req, 0, 1, conv, 0), (Req, 0, 2, conv, 0), (Req, 0, 2, rec, 0), and (Req, 1, 0, rec, 0). The capability covers the five request rows; the persistent row is their source anchor.

Table 5: Cohort-to-claim authorization. Cohorts are never pooled; each row supports only the listed inference.

Cohort	Frozen scope	Authorized claim / explicit exclusion
Primary ownership factorial	8 books; $N = 1, 8, 32$; four KV $\times$ GDN cells; 8 decode steps	ownership lifecycle, cross-cell equivalence, selected attention oracle, live-mutant sensitivity, and four-cell allocator endpoints; not transfer or scheduler evidence
Primary target-5 rerun	same frozen 8-rank, 96-configuration Qwen/H20 factorial; 209,920 attention and 635,520 GDN calls	per-call selected Triton kernel-cache artifact/configuration and eager-route binding at the declared fixed-stack scope; not device-side completion, driver/device, compiled-GDN, ATen/CUDA, performance, or cross-stack evidence
Selected-cell live binding	same frozen 8-rank factorial; stronger hook active only for the $N = 8$ shared-document/materialized-GDN cell; actual builder, production serializer, and direct clone lineage	144 selected rows, 12,960 live-storage rows, 3,840 clone edges, and 24 phase artifacts pass; all 96 clean-memory calls have zero selected-context overlaps; not evidence for other cells or coordinates, malicious runtimes, framework independence, device binaries, or performance
Fresh-process control	one independent $N = 2$ process; one clean path and M1	input reconstruction, cleanup, and one intended-gate reproduction; not pooled with the factorial
Lifecycle transfer	8 books; aligned 4,096-token prefix; $N = 4$; two cancelled slots	exact outputs/state, zero-scrub, exact reuse, and stale-lease rejection on the same adapter; not scheduler, concurrency, or performance evidence
Scheduler interleave	8 books; $N = 4$; page/prefix/tail geometries 64/3,985/17 and 128/4,033/65; cancel slot 3 after round 2	deterministic request-step interleaving, zero-scrub, epoch change, exact reuse, uninterrupted-control equality, and three frozen scheduler-fault gates; not concurrent kernels, continuous-batching performance, detection rate, or production end-to-end correctness

Continued on next page

Table 5 continued.

Cohort	Frozen scope	Authorized claim / explicit exclusion
Independent preproducer slot census	frozen Qwen/H20 geometry and schedule; 180 slots/capture; two $N = 2$ GDN-policy cells and three phases/cell	all 6 captures, 1,080 receiver rows, and 96,660 relations close; omission, duplication, and semantic-relabel controls fail closed; producer manifests/rows are not expected-set authority, but slot-ID-to-live-tensor binding, IPC, and paused capture remain trusted
Dual-producer repeat	same frozen input/protocol and producer implementation; two fresh serial executions and four distinct out-of-process observers	both executions reproduce 1,080 semantic coordinates, byte digests, stable descriptors, and 96,660 relation labels at zero tolerance; not malicious-producer resistance, independent receiver model execution, or cross-stack evidence
Falcon-H1 bounded transfer	tiuuae/Falcon-H1-0.5B-Base; Transformers 5.14.1/H20 naive path; $N = 1, 2$; separately executed official Transformers reference; persistent-Q16 and deep-materialized arms	8/8 ranks, 96 generated-token and 96 full-vocabulary FP32-logit comparisons, 13,824 state rows across KV key/value, convolution, and Mamba2 recurrent families, 192 ownership checks, cross-arm/cross- N relations, and registered controls pass; no performance, memory, quality, or other-stack evidence
Bounded two-stream lifecycle	one H20; one 4,033-token document; two 16-token query requests at $N = 2$; two host threads and distinct CUDA streams; pre-cancel and post-reclaim phases; cancellation after synchronization	positive call-interval overlap plus exact four lifecycle outputs, 20 logical-KV layer pairs, two GDN slot aggregates, ownership, scrub, and reuse; not simultaneous kernels, native continuous batching, in-flight cancellation, production scheduling, or performance
GDN transition oracle	1 frozen book/window; global model-layer indices 0, 10, 20, 38; 4 query tokens	clean recurrence outputs/states and four seeded wrong-transition rejections conditional on post-normalization inputs; not all GDN layers or end-to-end semantics
Expanded captured-boundary numerical sweep	2 distinct PG-19 inputs; all 10 attention layers and 12 of 30 GDN layers; 20 attention and 24 GDN clean rows plus one seeded control/row	captured-boundary numerical agreement and rejection of all 44 seeded wrong operators; not upstream-activation, independent-capture, full-model, or end-to-end validation
Target-gate suppression matrix	9 fresh clean + 9 target-suppressed cases; same fixed model/data/ $N = 2$; 8 distinct H20s	per-fault comparison with redundant audit gates, exact allowlisted production assertions, and token/FP32-logit equality; M8's injected sentinel is not a detector; not a blind-fault study or rate
Designer-executor-separated faults	separate designer reads only the prior PDF; 5 new frozen $N = 2$ clean/mutant pairs; 5 H20s; all gates enabled	all 5 clean gates pass and all 5 mutants fail first at their frozen primary predicate; per-fault sensitivity only, not a detection rate, natural-defect study, or cross-stack evidence
Historical alias regression	one pre-discovery integration defect; 3 archived coordinates + 5 additional frozen inputs; fresh pre-fix, repaired, and materialized $N = 2$ one-token lanes; 8 H20s	pre-fix binding gate and base violation with output/terminal-state equality, plus storage-clean exact repair; one retrospective defect only, not prevalence, recall, statistical replication, or cross-stack evidence
Mac common-mode control	Apple M4 Pro; six documents/queries; five fresh MLX processes; vanilla dense plus CoMem BF16/Q8/Q4 at depth 7	separates split-family agreement from agreement with vanilla dense; motivation only and separate from the H20 case

Continued on next page

Table 5 continued.

Cohort	Frozen scope	Authorized claim / explicit exclusion
H20 archival Store-F1 validation	60 validation items from Qasper/2WikiMQA (30 each); dense, full-prefix, split Q16, frozen Q4/Q4/Q8, and two mixed policies	raw-first replay of Store, predictions, mean F1, paired bootstrap intervals, and fixed-case failure counts; not source-frozen regeneration, TTFT/TPOT/throughput, Recall, total memory, edge, or broad-benchmark evidence
H20 deployment timing	source indices 6–9 from each of Qasper/2WikiMQA (8 items total), a subset of the 60-item validation cohort; no-cache dense recompute, full-prefix KV reuse, and CoMem Q8/Q4/mixed variants (six configurations total) with 3 timing repeats	same-stack retained-Store, TTFT, TPOT, tokens/s, and mean F1 comparison from a separate execution; not an independent replication, not pooled with the 60-item panel or primary factorial, and not scheduler/continuous-batching evidence
Matched serving controls	8 validation items from Qasper/2WikiMQA; vLLM 0.26 and SGLang 0.5.17 cache off/on plus 24 fixed-configuration HYPIC timing/quality cells; one H20-3e per item	within-block cache-hit, quality, and client-wall timing context; not pooled with in-process CoMem timings, not a cross-framework leaderboard, and not continuous-batching evidence
HYPIC retained-state receipts	8 Prefix Cache + 8 HYPIC cells on the same fixed TP=1 model/slice; receipt-instrumented SGLang lifecycle	median target-entry-owned physical tensor payload only; not timing, NVML/process memory, capacity, continuous batching, or ForkAudit ownership evidence
Related-work operator transfer	Hydragen at $N = 8/32$ on one captured layer-3 attention family; Palu on one layer with 8 calibration and 1 held-out PG19 row	same-checkpoint numerical and logical-state trade-offs only; not an all-layer implementation, LongBench/quality result, speedup claim, or jointly optimized system
Local replay/storage accounting	one Apple Mac16,8 CPU host; three consecutive fixed-order runs without cache flush; manifest-only core package plus a serial profile with five locally complete supporting verifiers	local CPU replay wall time and logical file/byte counts only; not H20 capture, GPU perturbation, cold-start latency, online inference, compressed-upload size, or an uninstrumented-baseline delta

Table 6: Trace coverage and replay verdict for the seven contract targets in fresh frozen primary case. Equality denotes canonical byte-digest identity as recorded. Overall trace coverage is complete and all seven target verdicts pass at their declared fixed-stack scope. A passing primary target verdict remains trace-relative and conditional on capture honesty, slot-ID-to-live-tensor binding, and the target boundary in Section 3. The census removes producer rows as expected-set authority only for its bounded GDN cohort; supporting cohorts do not upgrade these targets to comprehensive system correctness.

Target	Trace coverage	Replay verdict	Decisive witness	Boundary
Frozen identity	complete	pass	identical preflight/terminal source ledgers, terminal model rehash, and stable data/query/protocol/execution bindings	one primary-factorial formal run; the dual repeat recaptures only the bounded $N = 2$ GDN cohort
Prefix immutability	complete	pass	physical document-KV digests and GDN base guards across all three phases	one frozen configuration/geometry
Private ownership	complete	pass	all-pairs request-base and request-peer range replay at the primary registered transition; actual-builder, serializer, and direct-clone receipts for the selected cell	pointer-free conservative intervals; stronger live binding covers only six selected rows per phase in one cell
Tail-safe append	complete	pass	one-time 127-token private-tail copy before append	one partial-tail geometry
Host-side launch/config.	complete	pass	per-call ID, selected kernel-cache artifact/configuration, autotuner outcome, GPU/stream, and post-return seal; source-bound GDN eager route	normal launcher return, not device completion; no driver/device, compiled-GDN, or ATen/CUDA attestation
Cross-arm equivalence	complete	pass	tokens, shape/dtype-bound CPU-FP32 logit digests, logical KV, and GDN	one primary sequential schedule
Cross- N consistency	complete	pass	288 all-arm adjacent-fan-out comparisons over 72 rank-query identities	relational, not an independent oracle

Selected numerical oracles, constructed controls, lifecycle/scheduler checks, the independent census and dual-producer repeat, bounded two-stream and Falcon cohorts, and the historical reproduction are supporting checks rather than additional contract targets. Their captured-input, same-adaptor, process-separated-but-shared-IPC, second-configuration, and schedule boundaries are reported separately.

D POINTER-FREE STORAGE WITNESS SCHEMA

Target-to-record dependencies. The following matrix is normative. I, L, E, and F denote the Identity, Lifecycle, Execution, and Target/fault classes in Table 21; parenthesized entries name mandatory fields. Missing any marked field makes that target open. A preregistered geometry may make tail safety not applicable. Target 5 additionally requires an exact per-call ID, selected kernel-cache artifact/configuration, autotuner selection or explicit no-autotuner record, assigned GPU/stream, and post-return seal; missing any applicable field leaves target 5 open. A missing F field invalidates the corresponding fault row, but does not change the separately evaluated seven-target witness.

Target	I	L	E	F	Blind predicate
Frozen identity	code, model, data, protocol; role map	—	—	N/A	exact bytes and complete roles
Prefix immutability	document/base start digests	roles; setup and final guards	canonical content digests	N/A	start = final
Private ownership	request/base/private roles	setup, registered-transition, final bindings	normalized storage ranges	N/A	non-vacuous all-pairs disjointness
Tail-safe append	page geometry; document/private roles	first append event	ordered append and copied-tail receipt	N/A	one private copy before write
Host-side launch/config.	callable, source, protocol, selected artifact/config.	post-return GPU/stream assignment	seal; ordered call, query, mask, position, append, scale, autotuner outcome, eager route	N/A	exact per-call binding; no fallback or route contamination
Cross-arm equivalence	factor, query, and schedule identities	final bindings	token, logit, logical-KV, GDN records	N/A	canonical semantic equality
Cross- N consistency	query and fan-out identities	—	the same canonical semantic records	N/A	smallest- N equals each larger N
Fault sensitivity (supporting)	frozen target/gate map	matched-control rebuild and cleanup	injected outcome and failed predicates	required	target, injection, expected gate, outcome

Each tensor-view row has the exact fields `storage_id`, `shape`, `stride`, `storage_offset`, `dtype`, `device`, `storage_nbytes`, `tensor_nbytes`, and `content_sha256`; absolute addresses are forbidden. Within each phase, the producer assigns `storage-0000`, `storage-0001`, and so on in first-observation order to backing storages. Reusing an ID with a different device or storage size is invalid. Phase capture IDs are unique, whereas request and persistent lifecycle-guard IDs must remain stable across setup, the registered transition, and final generation.

Let element size be e , storage offset o , shape s_i , and stride t_i . The independent replay initializes $m = M = o$ and, for every dimension, computes

$$d_i = (s_i - 1)t_i, \quad m += \min(d_i, 0), \quad M += \max(d_i, 0). \quad (6)$$

The conservative half-open byte interval is

$$I = [em, e(M + 1)), \quad (7)$$

which must be non-negative and contained in `storage_nbytes`; also `tensor_nbytes` = $e \prod_i s_i$. Two rows overlap iff their normalized storage IDs match and $I_a^{\text{start}} < I_b^{\text{end}}$ and $I_b^{\text{start}} < I_a^{\text{end}}$. Exact aliasing additionally requires equality of interval, geometry, dtype/device, storage and tensor sizes, and content digest. For correctly live-bound rows at one capture, if represented views sharing physical backing receive the same normalized ID and their shape, stride, offset, and element size exactly describe affine-strided views, I contains every touched byte. The overlap test then cannot miss a shared byte between represented rows, although interleaving may cause false positives. Omitted rows, between-capture changes, and wrapper or storage-ID violations remain outside the claim.

At setup, every recurrent-state coordinate must either exactly alias the persistent state (borrowed-base arm) or be disjoint from it (materialized arm). At the primary registered 32-token transition boundary, a completed request’s 60 mutable rows must be disjoint from the base and every peer; incomplete borrowed requests must remain exact aliases. At final, all request-local GDN states must be disjoint from the persistent GDN base and every peer; document KV remains intentionally shared. Independent binding-token receipts additionally require unchanged tokens for incomplete requests and changed binding/storage tokens for completed requests. The package applies this schema to all 96 timelines and 288 phase artifacts; nine adversarial tests cover exact aliasing, partial overlap, adjacent intervals, negative strides, conflicting ID reuse, pointer injection, capture/guard reuse, and cross-coordinate request–base and request–peer aliasing.

E EXECUTED FAULT ROWS AND INVARIANT MAP

The row-level campaign is kept out of the main paper because it supports diagnosis and audit coverage rather than a separate headline result. Separately, the CPU governance suite passes 17 tests and rejects 28 explicit serialized-artifact tamper cases plus its registered helper/launcher faults. That result concerns schema and orchestration fail-closure only; it is not GPU semantic or performance evidence.

Table 7: Separate primary all-gates-on reference. Every separately rebuilt matched-clean cell passed; every mutant reached its target gate fixed before execution and was restored before disposal.

ID	Injected live fault	Expected/observed gate	Primary outcome
M1	reservation alias	KV.RESERVATION.DISJOINT	target gate reached
M2	sequence swap	KV.SEQUENCE.ID	target gate reached
M3	omit tail COW	KV.TAIL.COW	target gate reached
M4	GDN–base alias	gdn.completed.vs.base.disjoint	target gate reached
M5	GDN peer alias	gdn.completed.vs.peers.disjoint	target gate reached
M6	position off-by-one	POSITION.CANONICAL.VALUES	target gate reached
M7	materialized mask	MASK.CONTRACT	target gate reached
M8	wrong callable	KERNEL.CALLABLE.ID	target gate reached
M9	dense KV view	KV.PAGED.VIEW	target gate reached

Table 8: Same-system outcomes for nine designed faults. The primary campaign enables all gates; a separate target-gate-suppression campaign suppresses only the named gate. “Same” is exact and “N/O” an earlier stop. Rows are not a detection rate; M8’s injected sentinel is not a production detector.

Fault / target contract	All gates on	With target gate suppressed	Token / FP32 logits
M1 Reservation ownership	target gate	completes	Same / Same
M2 Sequence binding	target gate	completes	Same / Same
M3 Tail COW	target gate	other audit: active-block ownership	N/O / N/O
M4 GDN base isolation	target gate	other audit: peer isolation	N/O / N/O
M5 GDN peer isolation	target gate	completes; logits change ($L_\infty = 0.75$, rel. $L_2 = 0.0314$)	Same / Changed
M6 Positions	target gate	completes	Same / Same
M7 Mask contract	target gate	completes	Same / Same
M8 Callable identity	target gate	injected payload sentinel	N/O / N/O
M9 Paged-KV view	target gate	paired-view production assertion	N/O / N/O

Designer–executor-separated fault detail. The five fixed per-fault outcomes summarized in the main text are shown below; they remain constructed controls rather than a defect-population sample.

Table 9: Designer–executor-separated constructed fault outcomes. A separate designer saw only a frozen earlier manuscript PDF and fixed all five definitions, loci, payloads, clean requirements, and expected first predicates before executor preparation or candidate output. The cases are separated from executor preparation and outcome inspection, not from the predicate vocabulary. All matched clean cases pass; every mutant completes its registered horizon and is rejected first by the frozen primary predicate. Semantic comparisons are secondary: “Same” means canonical exact equality with the matched clean case, KV/GDN denote terminal logical state, and “N/C” means full-logit call cardinality is not comparable. These five fixed outcomes are not a population detection rate.

Frozen fault	Isolated mutation	First failed predicate	Tokens / logits / KV / GDN
HF01 delayed tail detach	append write precedes the otherwise present tail copy	tail-copy-before-first-write	Same / Same / Same / Same
HF02 inactive-lane scribble	one bit in the unused lane of a shared physical K page	physical-prefix immutability	Same / Same / Same / Same
HF03 duplicate commit	one correct feedback call commits twice; first output discarded	ordered-call cardinality	Same / N/C / Changed / Changed
HF04 effective-scale drift	first registered attention scale is exactly doubled	attention effective scale	Changed / Changed / Changed / Changed
HF05 stale GDN binding token	storage correctly rebinds; lifecycle token remains stale	completed binding-token advance	Same / Same / Same / Same

Table 10: Representative failure families and intended checks. Rows naming M1–M9 refer to the nine executed primary-factorial live mutants; the other rows summarize schema, relational, numerical, or replay checks. None is a complete oracle.

Check family	Intended sensitivity and observed boundary
Frozen identity	wrong model/data/query bank, changed protocol, stale shard (schema adversaries)
Prefix immutability	M3 omits partial-tail COW; physical document digest and ownership gates fail
Private ownership	M1 reservation alias, M4 request–base GDN alias, M5 request–peer GDN alias
Sequence/view binding	M2 sequence swap and M9 dense-KV replacement
Host-side selected launch/configuration	M6 position offset, M7 materialized mask, M8 wrong callable
Factorial exactness	behavior changes when either KV or GDN setup ownership changes
Bounded numerical oracle	numerical error in a selected fused full-attention result
Raw replay	internally inconsistent trajectory, storage, allocator, or aggregate fields

The primary all-gates-on execution is a first-gate localization study. Every registered gate fires before the injected path produces a post-injection token, logit, logical-KV digest, or GDN digest; its output cells are therefore N/O rather than misses. Lifecycle/ownership/storage receipts supply the first gate for M1 and M3–M5; request-binding/call receipts supply it for M2 and M6–M9. The full audit rejects all nine at their predeclared gates under passing matched controls. The separate target-gate-suppression matrix then suppresses one named gate per mutant while leaving the remaining observers active. It provides the per-fault comparison in Table 8, but does not estimate a population rate, unseen-fault coverage, or false-positive rate.

F RELATION TO PRIOR WORK AND AUDIT OBLIGATIONS

Table 11: Primary-source positioning. “Evidence focus” describes what the cited work evaluates; it is not a claim that the work lacks all other checks.

Work	Primary mechanism	Evidence focus / relation to this work
PagedAttention	Paged KV allocation and sharing (Kwon et al., 2023)	Serving throughput and memory utilization; supplies the block-sharing substrate.
SGLang / Prompt Cache	Prefix reuse in structured or modular prompts (Zheng et al., 2024; Gim et al., 2024)	Runtime performance and output behavior; establish reuse and positional handling as prior art.
CoMem	One persistent residual per token at split j ; bounded selection; suffix replay (Liu et al., 2026a;b)	Supplies the split-replay foundation; we do not claim that primitive as new.
Q-COMEM (ours)	Joint low-bit split residual and required lower KV/convolution/recurrent state	Adds complete hybrid-state packing, a capacity–latency–quality evaluation, and an auditable request-local ownership boundary; not a model-weight compressor or production scheduler.
Marconi	Admission and eviction for hybrid-model prefixes (Pan et al., 2025)	Hybrid cache hit rate and time to first token; closest state-heterogeneous caching system, but a different policy-level contribution.
SCBench	Lifecycle-oriented KV-cache benchmark (Li et al., 2025)	Breadth over tasks and cache phases; exposes breadth absent from our bounded declared configurations.
MT-DLComp / CRADLE	Metamorphic compiler testing and cross-backend differential checking (Xiao et al., 2022; Pham et al., 2019)	Direct testing precedents; our audit adds ownership witnesses and a layerwise dense oracle, but not an independent end-to-end backend.
DNNMem	Component-wise GPU-memory accounting (Gao et al., 2020)	Supports denominator separation; it is an estimator, not a request-ownership audit.
FORKAUDIT	Typed, phase-indexed ownership/call/accounting trace with explicit coverage semantics	Validates the ownership component through a mutation-tested KV-by-GDN factorial, bounded host-side call receipts, a preproducer expected-slot census, numerical sweeps, and one bounded Falcon-H1 comparison; it is not a security attestation or general end-to-end correctness oracle.

Table 12: Closest established precedent for each audit obligation. The final column states the narrower integration contributed by FORKAUDIT; it does not claim that any individual mechanism is new.

Audit obligation	Closest established precedent	Integrated obligation in this work
Frozen identity	Controlled variants in metamorphic and differential testing (Chen et al., 2018; Xiao et al., 2022; Pham et al., 2019)	Bind model, data, query bank, callable contract, and protocol bytes to every ownership receipt.
Document immutability and tail COW	PagedAttention block sharing and copy-on-write for shared prefixes (Kwon et al., 2023)	Replay physical document pages, padding, active tables, and the partial-tail detach event.
Recurrent-state ownership	Hybrid recurrent states and their in-place update constraint (Yang et al., 2025; Pan et al., 2025)	Witness request-base and request-peer ranges before and after the frozen schedule’s registered transition boundary.
Host-side selected launch/configuration	Cross-backend differential checking exposes implementation-dependent behavior (Pham et al., 2019)	Hold the backend fixed; record callable, positions, mask, scale, and append history; and bind the selected Triton kernel-cache artifact/configuration per attention call while keeping GDN at an explicit eager-route boundary.
Cross-arm and cross- N relations	Metamorphic relations over semantics-preserving transformations (Chen et al., 2018; Xiao et al., 2022)	Treat ownership policy and resident fan-out as transformations that must preserve registered request observables.
Bounded numerical oracle	Independent backend comparison (Pham et al., 2019)	Reconstruct selected attention inputs for dense IEEE-FP32 math and selected post-normalization GDN inputs for a candidate-import-free NumPy recurrence.
Memory denominators	Component-wise GPU-memory accounting (Gao et al., 2020)	Keep page geometry, allocator current/peak, and unmeasured process/service capacity as separate claim domains.

G SUPPORTING CONTROL EXPERIMENTS

Mac common-mode motivation. The archival Apple-M4-Pro control is kept unpooled from the H20 experiments and serves only to show why split-family agreement does not replace a dense or independently implemented oracle.

Table 13: Unpooled Mac motivation context (Apple M4 Pro, five fresh MLX processes; median over 30 query executions at split depth 7). Vanilla dense is the external reference. CoMem BF16/Q8/Q4 variants use the same selected documents. Store size is the six-document residual corpus; online time includes query Write, dequantization, and suffix Read. This small teacher-forced demonstration set is not primary ForkAudit evidence and supports no cross-hardware or production-speed claim.

Method	Store (MiB)	Online (ms)	Speedup	Top-1 vs. dense
Vanilla dense prompt	–	194.70	1.00×	100.00%
CoMem BF16 split	1.998	168.88	1.15×	71.89%
CoMem Q8 split	1.061	169.17	1.15×	71.89%
CoMem Q4 split	0.562	169.25	1.15×	71.89%

Q8/Q4 retain 100% top-1 agreement with the BF16 split path, yet all split rows inherit only 71.89% agreement with vanilla dense. The speed column is a small teacher-forced mechanism benchmark, not production RAG throughput.

Historical alias regression and repair. The pre-fix, repaired, and materialized lanes isolate one previously observed recurrent-state alias. This is a fixed mechanism study, not a natural-bug frequency or detector-recall estimate.

Table 14: Retrospective reproduction and repair of one historical alias regression. Ranks 0–2 reuse three archived input coordinates; ranks 3–7 are additional frozen inputs, not additional defects or statistical replicates. Semantic columns compare each path with a fresh materialized control. “Base inv.” is within-path persistent-base immutability, retained as a conventional state-invariant catch.

Coordinates	Path	Ownership result	Token	FP32 logits	Request GDN	Logical KV	Base inv.
Archived (3)	Pre-fix	binding gate 3/3	exact 3/3	exact 3/3	exact 3/3	exact 3/3	fail 3/3
Archived (3)	Repaired	storage-clean 3/3	exact 3/3	exact 3/3	exact 3/3	exact 3/3	pass 3/3
Additional (5)	Pre-fix	binding gate 5/5	exact 5/5	exact 5/5	exact 5/5	exact 5/5	fail 5/5
Additional (5)	Repaired	storage-clean 5/5	exact 5/5	exact 5/5	exact 5/5	exact 5/5	pass 5/5

Primary allocator endpoints. The four ownership cells use one frozen post-priming allocator denominator and are kept separate from deployment Store.

Table 15: Primary $8 \times H20-3e$ allocator results at $N = 32$ (median across ranks; rank values coincide). Full-copy is the factorial control; the paged-prefix row is the closest same-stack prefix-sharing baseline. The remaining rows isolate recurrent-base ownership. All cells retain the persistent document source. Values are GiB above the frozen post-priming baseline; generation increment is the additional peak above allocation present at generation start.

Empirical arm	KV / GDN setup	Final	Peak	Gen. incr.
Full-copy factorial control	Full-copy KV / Materialized GDN base	4.901	4.920	0.019
GDN ownership ablation	Full-copy KV / Borrowed GDN base	4.890	4.907	1.951
Paged-prefix baseline	Shared-doc KV / Materialized GDN base	2.229	2.843	0.019
Audited hybrid fork	Shared-doc KV / Borrowed GDN base	2.229	2.843	1.950

Published systems context. Table 16 reports each cited system only in its native protocol and keeps all non-commensurate metrics outside our matched panels. This unpooled table is literature context, not ForkAudit evidence.

Table 16: Unpooled published-systems context in each paper’s native protocol. These values show the scale and evaluation breadth of related systems, *not* a cross-paper leaderboard or ForkAudit evidence: models, hardware, batching, baselines, and metrics differ. Consequently, no row is merged with our same-stack H20 Table 18, and the absence of a common score must not be read as a negative result.

System	Native method and study setting	Reported efficiency in its native protocol	Quality and comparability boundary
PagedAttention / vLLM (Kwon et al., 2023)	Paged KV allocation plus serving engine <i>Setting:</i> LLaMA-7B on A10G and LLaMA-13B on A100-40GB; vLLM serving engine	2–4× serving throughput at similar latency vs. FasterTransformer/Orca	No common LongBench F1; exact cache-management method
Prompt Cache (Gim et al., 2024)	Modular attention-state reuse <i>Setting:</i> Llama2/MPT/Falcon; RTX 4090, A40, and A100; HuggingFace-style prototype	1.5–3.1× TTFT (CPU-resident modules); 3.7–11.7× (GPU-resident) on its 12-dataset LongBench suite	LongBench accuracy evaluated, but with different models and hardware
SGLang / RadixAttention (Zheng et al., 2024)	Prefix reuse plus structured-program runtime <i>Setting:</i> Llama-7B/Mixtral-8x7B and structured programs; A10G/A100; SGLang	Up to 6.4× throughput vs. prior inference systems	Application-suite result; no common LongBench F1
ChunkAttention (Ye et al., 2024)	Prefix-aware KV chunks and attention kernel <i>Setting:</i> attention microbenchmarks; mainly A100-80GB, plus A100-40GB/RTX 4090; CUDA 11.8	3.2–4.8× attention-kernel speedup for 1K–4K shared prompts	Kernel result; no end-to-end common task score
Hydragen (Juravsky et al., 2024)	Exact batched shared-prefix attention <i>Setting:</i> CodeLlama-13B/34B shared-prefix batches; A100 (with H100/L40S microbenchmarks); custom HuggingFace runtime	Up to 32× CodeLlama-13B end-to-end throughput in large shared-prefix batches	Exact attention construction; different model and batching regime
Palu (Chang et al., 2025)	Low-rank latent KV with reconstruction/fusion <i>Setting:</i> Llama/Mistral/LongChat quality studies; Llama-2-7B latency on one RTX 4090; HuggingFace plus custom kernels	50% KV compression and up to 1.89× RoPE-attention speedup; up to 2.91× with quantization	Llama/Mistral/LongChat native study; our bounded Qwen3.5 operator transfer is reported separately
Preble (Srivatsa et al., 2025)	Distributed prompt-sharing scheduling <i>Setting:</i> Mistral-7B/Llama-3-70B request traces; two–eight GPUs across A6000/H100 clusters; distributed serving	1.5–14.5× average-latency and 2–10× p99-latency improvement	Distributed serving traces; no common task score
Marconi (Pan et al., 2025)	Admission/eviction for hybrid-model prefixes <i>Setting:</i> Attention–Mamba2 7B geometry; LMSys/ShareGPT/SWE-bench traces; CloudLab policy simulation	Up to 34.4× token-hit rate and up to 71.1% or 617 ms lower TTFT	Hybrid models and trace-driven policy study; no common LongBench F1
HYPIC (Liu et al., 2026c)	Position-independent caching for hybrid full/linear-attention models <i>Setting:</i> Qwen3.5-35B-A3B main panels at TP=2 on 8 H20-3e; SGLang 0.5.14	3.25× lower TTFT on average and 1.66× higher QPS over Prefix Cache	Four models/five workloads; our same-slice TP=1 timing and retained-state controls are reported separately

Same-checkpoint transfer detail. Table 17 is the matched complement to the published context table: it uses our frozen Qwen3.5 checkpoint and H20-3e, binds the official Hydragen and Palu source revisions, and independently replays all raw operator sidecars. Because it does not execute all layers or LongBench, it is kept separate from Table 18.

Table 17: Bounded same-checkpoint H20 transfers. Hydragen uses frozen Qwen3.5 layer-3 post-RoPE attention; Palu uses eight 256-token PG19 calibration rows and one disjoint held-out row. Errors are relative ℓ_2 against the matched uncompressed operator; state ratios are logical selected-operator ratios, not process memory or end-to-end results.

Port	Setting	Matched numerical result	Logical state result	Timing / boundary
Hydragen port	$N = 8$	Hydragen-dense rel. $\ell_2 = 0.00255$	8.50 vs. 64.48 MiB (7.59× smaller)	0.287 vs. 0.0965 ms; no speedup
Hydragen port	$N = 32$	Hydragen-dense rel. $\ell_2 = 0.00262$	10.00 vs. 257.94 MiB (25.80× smaller)	0.290 vs. 0.259 ms; no speedup
Palu port	rank 64/head	held-out K/V 0.487/0.549 (plain 0.711/0.839)	512 vs. 2,048 B/token (4.00× smaller)	projection only; no fused kernel
Palu port	rank 128/head	held-out K/V 0.317/0.364 (plain 0.543/0.668)	1,024 vs. 2,048 B/token (2.00× smaller)	projection only; no fused kernel
Palu port	rank 192/head	held-out K/V 0.149/0.194 (plain 0.341/0.459)	1,536 vs. 2,048 B/token (1.33× smaller)	projection only; no fused kernel

Hydragen avoids prefix replication but has no speedup here. Palu improves over plain SVD, but residual error forbids an end-to-end quality claim. Neither is a jointly optimized combined system.

G.1 EXTENDED DEPLOYMENT AND SERVING CONTEXT

Table 2 isolates the matched Transformers block used for the main Q-COMEM claim. The expanded table below repeats that block only to place separately measured SGLang/HYPIC context beside it. The blocks use different runtimes and harnesses, so timing and F1 are interpreted within block and are not pooled. These serving controls are not FORKAUDIT validation evidence and do not establish robust quality preservation.

Table 18: Illustrative same-checkpoint H20 retained Store and mean-F1 point estimates on an eight-item Qasper/2WikiMQA slice. All rows use Qwen3.5-35B-A3B, one H20-3e per item, a 4,096-token cap, greedy decoding, and at most 32 generated tokens. The Transformers (three repeats/item) and HYPIC (Liu et al., 2026c) blocks are unpooled; HYPIC timing/F1 and Store use separate 24- and 16-cell TP=1 cohorts. Timing is compared only within blocks; no robust-quality, capacity, or ForkAudit-ownership claim is made.

Method	Persistent state	Store (MiB)	TTFT (s)	TPOT (ms)	tok/s	F1
<i>Transformers 5.14.1 in-process deployment cohort</i>						
No-cache dense recompute	—	—	0.649	648.75	1.36	39.42
Full-prefix KV cache	Q16	140.34	0.163	54.11	13.38	39.14
CoMem state	Q8	15.89	0.673	55.44	7.15	39.14
CoMem mixed state	Q4/Q4/Q8	10.01	0.674	54.97	7.19	39.01
CoMem per-layer mixed	mixed	9.74	0.673	55.07	7.19	39.11
CoMem state	Q4	8.41	0.674	55.02	7.19	36.09
<i>HYPIC commit 98147c0 timing/F1; receipt-instrumented Store; SGLang 0.5.14, TP=1</i>						
Official HYPIC code: full recompute	—	—	0.217	49.21	14.77	39.14
Official HYPIC code: prefix cache	Radix prefix	139.53	0.072	50.15	19.07	39.14
HYPIC (transition + seam 8)	segment transition	324.09	0.101	52.13	17.67	39.63

Bold marks the two illustrative Store-F1 points; these are not a headline validation result. Relative to full-prefix Q16, Q8/per-layer mixed reduce Store 88.68%/93.06%; unrounded mean-F1 deltas are 0.000/−0.022 (39.115 vs. 39.137 for mixed). “—” means no retained entry, and Q16 denotes BF16. Full-prefix Q16 is CoMem’s matched semantic comparator; no-cache recompute uses a different document/query execution boundary. Q4/Q4/Q8 gives residual/attention/linear bits. The per-layer policy (Q4 residual, Q16 projections, layer bits [8, 8, 4, 4, 8, 8]) was selected on a separate, disjoint four-item calibration set and frozen before this eight-item validation. Store is median retained-document tensor payload and excludes metadata, pools, process deltas, and capacity. Published HYPIC panels use TP=2; this adaptation uses TP=1.

Matched serving-framework and HYPIC context. Table 19 gives separately timed vLLM and SGLang cache-off/on controls plus the official-code HYPIC Full Recompute, Prefix Cache, and full HYPIC arms under the shared Qwen3.5 streaming boundary. It remains a timing/quality-only panel and authorizes only within-framework interpretation; the separate 16-cell retained-state receipt cohort appears in Table 18. vLLM and SGLang each record 8/8 prefix hits and preserve 8/8 cache-off predictions. Within its TP=1 block, HYPIC gives 39.63, 0.101 s, and 17.67 tokens/s; published Qwen3.5 panels use TP=2, so this is an adaptation. Its 24 timing cells bind a 4,404-entry terminal artifact ledger. Across the 16 retained-state cells, Prefix Cache retains a median 139.53 MiB and HYPIC retains 324.09 MiB; Prefix Cache and HYPIC retain 8.78× and 20.39× as much Store as

CoMem Q8, respectively, under the separate Store estimand. This remains single-stream; it does not measure continuous batching.

Table 19: Unpooled timing/quality-only serving context on Qwen3.5 and the same eight-item Qasper/2WikiMQA slice. Each row uses one H20-3e per item, a 4,096-token cap, greedy decoding, and at most 32 generated tokens. Times are OpenAI-compatible streaming client wall-clock; F1 is the eight-item mean. Hit is interpreted only against the disabled/full-recompute reference inside each framework block. No cross-framework speedup, retained-state, scheduler-throughput, capacity, or broad-quality claim is made.

Framework	Prefix reuse	F1	TTFT (s)	TPOT (ms)	tok/s	Hit
vLLM 0.26	off	39.63	1.445	56.31	4.70	–
vLLM 0.26	align on	39.63	0.179	53.99	14.00	8/8
SGLang 0.5.17	off	39.14	0.283	53.38	12.73	–
SGLang 0.5.17	Radix on	39.14	0.146	53.15	15.83	8/8
HYPIC code	full recompute	39.14	0.217	49.21	14.77	–
HYPIC code	prefix cache	39.14	0.072	50.15	19.07	8/8
HYPIC	transition+seam 8	39.63	0.101	52.13	17.67	8/8

vLLM uses experimental Mamba align; SGLang 0.5.17 uses extra_buffer Radix caching. The HYPIC rows use the authors’ commit 98147c0 on SGLang 0.5.14, TP=1, full transition_rope_recompute, and an eight-token seam. Its full and prefix rows are controls from that same codebase; published TP=2 results and the separate retained-state receipt cohort are not inserted here.

Hybrid-prefix policy context. The separate preregistered Marconi trace extends the eight frozen workloads with a fixed synthetic follow-up turn and evaluates only prefix-policy token reuse under the artifact’s native Attention–Mamba2 geometry. It neither reruns LongBench quality nor supplies Qwen3.5 serving speed, so it is not pooled with the preceding panel or the primary audit.

Table 20: Unpooled preregistered policy-simulator context on a 128-event synthetic multi-turn extension of the eight frozen workloads. Values are token-hit percentages under the official Marconi artifact’s native Attention–Mamba2 geometry. The wrapper calls each policy’s own eviction routine after insertion to close actual tree size to the declared budget. These are not Qwen3.5 quality, latency, or throughput measurements and do not validate the primary ownership audit.

Policy	5 GB	10 GB	20 GB
vLLM+	2.07	6.18	46.83
SGLang+	90.78	93.86	93.86
Marconi	93.86	93.86	93.86

Table 21: Normative record classes. “Mandatory” means that absence makes the corresponding target open rather than silently not applicable.

Record class	Mandatory payload	Independent replay
Identity	frozen code, model, data, protocol, and object-to-role map	exact bytes and role coverage
Lifecycle	setup, registered-transition, and final inventories and bindings	normalized ranges, aliases, and disjointness
Execution	ordered call/append receipts, semantic records, named allocator snapshots	frozen schedule, digests, and phase equations
Target/fault	predeclared target, gate, matched control, injection, and outcome	failed-predicate set and classified outcome

Aligned cancellation and reclamation lifecycle. The separate lifecycle-transfer cohort uses eight ranks, eight distinct PG-19 training books, an aligned 4,096-token prefix, and $N = 4$. It is not pooled with the primary factorial.

Table 22: Lifecycle-transfer obligations on the existing Qwen3.5/vLLM Q16 adapter. Counts are deterministic rank outcomes, not stochastic replicates.

Obligation	Result
Frozen static/code/model/weight/raw bindings	pass; 14/14 weights, 8/8 shards
Full-vocabulary logits and generated tokens vs. uninterrupted control	exact, 8/8
Final logical KV and final GDN state	exact, 8/8
Aligned prefix, immutable document, and zero partial-tail copy	pass, 8/8
Cancelled-slot zero-scrub and exact reservation reassignment	pass, 8/8
Stale epoch-0 rejection and matched epoch-1 acceptance	8/8; 0 escapes, 0 wrong gates

Scheduler-managed request-step interleaving. The separate formal extension keeps four requests resident, cancels slot 3 after its second round, zero-scrubs and increments its lease epoch, reassigns the exact reservation, and completes the replacement interleaved with the three survivors. It is not pooled with the primary factorial or the aligned lifecycle-transfer cohort.

Table 23: Formal scheduler-managed request-step interleaving on the frozen Qwen3.5/ForkAudit vLLM-Q16 stack. Clean counts are deterministic rank outcomes; fault counts are preregistered expected-gate outcomes, not a detection rate.

Case	Frozen geometry or intervention	Outcome
Clean geometry 1	page 64; 3,985-token prefix; 17-token document tail; 18-event cancel/replacement lifecycle	all semantic/storage gates and uninterrupted-control comparisons pass, 8/8
Clean geometry 2	page 128; 4,033-token prefix; 65-token document tail; same lifecycle	all semantic/storage gates and uninterrupted-control comparisons pass, 8/8
FH1	dispatch cancelled epoch-0 lease	STALE_SLOT_LEASE, 16/16
FH2	dispatch lease under another request	LEASE_REQUEST_MISMATCH, 16/16
FH3	reclaim before zero scrub	RECLAIM_NOT_ZERO, 16/16

Candidate-import-free GDN transition oracle. The oracle uses one frozen PG-19 window and four query-time GDN transitions at global model-layer indices 0, 10, 20, and 38. A single preregistration frozen before the successful run binds the post-native-q/k-normalization boundary, scale $1/\sqrt{128}$, tolerances, and seeded faults. Its NumPy reference starts at that captured boundary; no candidate module is imported by the reference.

Table 24: PDF-visible oracle selection and replay boundary. Every listed row was locked before its successful candidate execution; neither oracle validates upstream activation construction or end-to-end model behavior.

Oracle	Frozen rows / selection rule	Captured boundary	Decision rule	Replay availability
Attention	rank r uses layer $4r + 3$ and round r , for $r \in \{0, \dots, 7\}$; request 0, $N = 1$, shared-document KV + borrowed GDN; frozen rank/layer/round map	post-RoPE Q/K/V, positions, mask, scale, and ordered append shadows	dense CPU-FP32 rel. $L_2 \leq 0.005$	8 raw rows and numerical sidecars; primary-package one-command replay
GDN	global model layers 0, 10, 20, 38; four query tokens from one frozen PG-19 window	post-native-q/k-normalization recurrence; scale $1/\sqrt{128}$	output rel. $L_2 \leq 0.005$; state rel. $L_2 \leq 10^{-4}$; seeded fault must reject	4 clean + 4 fault decisions from 40 sidecars; separate one-command replay

Table 25: Bounded GDN recurrence oracle. Relative errors compare native outputs and final states with the candidate-import-free NumPy FP32 recurrence.

Global layer	Output rel. L_2	State rel. L_2	Clean	Seeded fault	Rejected
0	0.001584	1.23×10^{-7}	yes	omit decay	yes
10	0.001652	1.36×10^{-7}	yes	complement β	yes
20	0.001572	1.39×10^{-7}	yes	pre-decay memory	yes
38	0.001638	1.30×10^{-7}	yes	roll value heads	yes

Table 26: Preregistered expanded captured-boundary numerical sweep. Clean rows use candidate-import-free CPU/NumPy FP32 replay; every clean row has one seeded wrong-operator control. Attention begins after RoPE and GDN begins after native q/k normalization, so the table does not validate upstream activation construction, independent capture, or end-to-end model behavior.

Operator	Inputs	Layers	Clean rows	Positions / transitions	Max output rel. L_2	Controls rejected
Attention	2	10/10	20/20	160 positions	0.001897	20/20
GDN	2	12/30	24/24	192 transitions	0.002073 (2.08×10^{-7} state)	24/24

Full target-gate-suppression matrix. The preregistered target-gate-suppression execution contains all M1–M9 faults and a fresh matched clean path for each. Across eight distinct H20s, all 18 cases complete their declared horizon or an allowed classified stop, all 24 FP32 sidecars verify, and no case is operationally invalid. Five mutant paths complete semantics: token equality catches 0/5, exact logits catch M5 only ($L_\infty = 0.75$, relative $L_2 = 0.0314$), and M1/M2/M6/M7 leave both exact. M3/M4 reach other audit gates; M9 reaches the paired-view production assertion; M8 reaches the exact injected payload sentinel, which is not a production detector. The frozen aggregator initially read the inherited execution identifier from a detached-manifest field that the manifest schema omits. A disclosed postexecution wrapper derives that identifier from the canonical receipt embedded in all eight manifest-bound primary-package shards, changes only the generated comparison-row `run_id` from null to the preregistered value, and reproduces the summary byte-for-byte without modifying candidate outputs or frozen sources.

H EXPANDED LIMITATIONS AND CLAIM BOUNDARY

The primary factorial covers one model, one H20 hardware family, one page size, 16-bit KV, and $N \leq 32$. Its ownership conclusion is tied to one 4,095-token partial-tail geometry, a registered 32-token continuation block, and the subsequent frozen sequential schedule; it is not a claim about an arbitrary first token or transition length. The lifecycle cohort adds one aligned 4,096-token geometry and a fixed cancel/scrub/reclaim sequence. The scheduler-interleave extension adds two geometries and deterministic request-step interleaving, but those requests execute sequentially on one CUDA stream.

The separate two-stream cohort has two host workers and distinct CUDA streams with positive full-model-call interval overlap. Its intervals include queuing, resource waits, and host-enqueue gaps; there is no per-kernel overlap trace. Cancellation occurs only after synchronization, so the cohort does not test arbitrary or in-flight reclamation. The original out-of-process GDN observer removes same-process reconstruction and live judgment labels. The later preproducer census independently fixes the expected slot set, but correct slot-ID-to-live-tensor binding remains trusted; both sides also trust PyTorch tensor/storage and CUDA-IPC semantics, and mutation is paused for each snapshot. The separate selected-cell hook binds an actual builder return through the production serializer and direct clone lineage, but still runs with the producer in one honest process. These cohorts do not supply OS/driver allocation ground truth, independent receiver-side model execution, KV/kernel attestation, or protection against a transient write followed by restoration.

No cohort evaluates native continuous/ragged batching, eviction, multi-document serving, general scheduler cancellation, process-level NVML memory, OOM capacity, or a broad downstream benchmark. Full-copy KV and materialized GDN bases are controlled ownership factors rather than optimized production baselines. The GDN path is the Transformers Torch implementation, not an

optimized recurrent kernel. The target-5 rerun closes only the declared honest-process host-side boundary: attention calls bind the selected Triton kernel-cache artifact/configuration and autotuner record, whereas GDN calls bind eager source routes. It does not attest device-side completion, compiled GDN or underlying ATen/CUDA operators, driver/device binaries, a malicious runtime, or another model/runtime/hardware stack. The live-binding rerun strengthens ownership provenance only in one selected cell and six rows per phase; other factorial cells and coordinates retain the trace-relative binding boundary.

The 60-item Qasper/2WikiMQA quality cohort is the broadest direct answer-quality evidence in this paper, but it is a single archival validation run on one model/H20 stack. Its raw shards and outcome calculations replay, whereas the original mutable code directory and absence of pre-run full source/model-weight ledgers prevent a source-frozen regeneration claim. Its paired interval does not establish equivalence or noninferiority. It records neither TTFT nor Recall. The separate systems execution uses the indices-6–9 subset of this cohort and supplies TTFT/TPOT/throughput point estimates, not an independent replication, population uncertainty, or a speedup over full-prefix reuse; it also has no Recall measurement. The two panels are not pooled, and blank table cells preserve these gaps explicitly.

The auxiliary serving panel uses one request stream per engine and compares each cache-enabled framework only with its own disabled reference. It does not measure continuous batching, eviction, admission control, maximum concurrency, or best-tuned performance. CUDA graphs were disabled in both pinned runs; the reported values describe that fixed configuration only.

The HYPIC timing/quality comparison uses the authors' exact tracked commit on SGLang 0.5.14 in one fixed CUDA 12.9 / PyTorch 2.11 configuration, adapting Qwen3.5 execution from the paper's TP=2 to TP=1 so every arm uses one H20-3e per workload. Cache-hit authority comes from OpenAI response usage. Each of its 24 timing/quality cells has one formal measurement after a discarded, prefix-disjoint warmup. A separate receipt-instrumented run uses the same model/slice and fixed TP=1 serving configuration to produce 8 Prefix Cache and 8 HYPIC retained-state cells; each completed cell binds raw output plus pre-measurement and terminal ownership receipts, and all 16 pass an external hash-bound replay. The instrumentation adds target-entry byte-range receipts and a lifecycle-idempotency overlay; Store therefore does not come from the unmodified timing run. Neither cohort establishes scheduler throughput, QPS, capacity, the paper's TP=2 latency, or accuracy robustness beyond this eight-item slice.

The same-checkpoint related-work transfer is likewise operator-bounded: Hydragen uses one captured layer-3 shared-prefix attention family, and Palu uses one layer with eight calibration rows and one held-out row. Neither executes all model layers, a fused Palu kernel, LongBench, or a jointly optimized Hydragen/Palu-plus-CoMem serving path.

The page equations are deterministic identities, not learned laws; the primary factorial has one observation per rank-configuration cell and no stochastic replication. Most runtime tensors are represented by hashes and equality receipts rather than archived in full. The attention and GDN oracle rows are the exception and ship numerical sidecars, but cover only selected captured-input operator boundaries rather than an end-to-end model. The expanded numerical sweep covers two short inputs, all ten attention layers, and 12 of 30 GDN layers; it does not establish length/input robustness or full recurrent-layer coverage. Both the original and expanded GDN runs start after native q/k normalization and therefore do not validate that upstream boundary. Pointer-free normalized intervals plus raw-byte hashes make the primary captured storage evidence replayable. The preproducer census supplies the expected semantic slot set, and a second producer/observer execution reduces single-run capture risk, but both still trust correct slot-ID-to-live-tensor binding, IPC semantics, and paused capture; neither protects against a malicious producer. The selected-cell live binding directly checks the actual builder, serializer, and clone lineage, but shares the same honest runtime and does not generalize to unselected coordinates. The separate Falcon track compares against an official full-model reference only on its frozen Transformers/H20 naive path and supports no performance, memory, quality, or cross-runtime claim. The nine mutants are targeted positive controls; the target-suppression matrix adds a same-system comparison for those same designed faults, not blind or naturally occurring defects. Four complete with unchanged measured tokens and exact logits, one changes only full logits, and four stop earlier; N/O cells are not misses, while the injected M8 sentinel is not a production detector. This is not a general detection rate, a blind-fault result, or superiority over every conventional test suite. The scheduler extension's three frozen faults are expected-gate

outcomes, not a detection rate or evidence of fault-set completeness. Broader systems claims would additionally require a native production serving scheduler, per-kernel overlap evidence, in-flight cancellation tests, optimized controls, alignment/length/input sweeps, and process-level capacity accounting.

The five designer–executor-separated controls were fixed before executor preparation, but remain deliberately constructed, fixed-stack mutations. Passing all five does not estimate unseen-fault recall, false-positive rate, or robustness to arbitrary runtime, model, schedule, or adversarial capture defects.

The retrospective case study concerns one organically encountered integration defect in our borrowed-state path. Only ranks 0–2 reproduce its archived input coordinates; ranks 3–7 are additional frozen inputs, and none is a statistically independent defect. The study does not estimate natural-defect prevalence or recall. Output and terminal-state differentials miss the alias, while a persistent-base content invariant catches it; FORKAUDIT contributes authenticated transition-time owner/layer/family localization rather than exclusive detection over all conventional testing.

I ARTIFACT AND CLAIM MAP

Artifact footprint. The source-complete primary package, derived summaries, raw ledgers, and expanded numerical sidecars are content-bound and locally replayable. The table below maps each manuscript claim to its machine-readable source and validation boundary; supporting measurements not used by a manuscript claim remain in the internal experiment registry.

Table 27: Artifact and claim map. Relative paths are rooted at the accompanying research package. Completed rows authorize only their stated bounded claims. Metadata-redacted derivatives preserve numerical fields and state digests, while most full tensors still require the pinned environment for independent numerical recomputation.

Claim/evidence	Machine-readable source
Primary factorial replay package	evidence/round_04_rr2_package/
Replay and package manifest	evidence/round_04_rr2_package/replay/run_replay.sh and evidence/round_04_rr2_package/MANIFEST.json
Registered, derived, and raw evidence	upstream/forkeaudit-summary.json, derived/derived_summary_v2.json, and upstream/raw/ under the primary replay package
Executed source and storage tests	executed_source/gpu/ and replay/test_storage_witness.py under the primary replay package
Identity, lifecycle, scheduler, and GDN extensions	evidence/rr2_terminal_closure_a1_20260821/, evidence/lifecycle_transfer_reviewer_package/, evidence/round23_a2_scheduler_interleave/, and evidence/gdn_transition_oracle_preregistered_20260819d/
Mac and H20 deployment controls	evidence/mac_m4_motivation/, evidence/h20_deployment_benchmark/, and their generators under scripts/
Q-CoMem 60-item Store-F1 validation	evidence/qcomem_mixed_validation_60item_20260812d/; exact raw-outcome replay: replay/run_replay.sh; numerical summary and claim boundary: validation_report.json and claim_boundary.json; registry entry E-QCOMEM-60-VALIDATION-20260812D-A
HYPIC retained-state receipts	evidence/experiment_registry.json, entry E-R34-HYPIC-STORE-EXTERNAL-ACCEPTANCE; the anonymous package assigns a neutral artifact alias
Target-gate-suppression preregistration and raw ledger	evidence/r28_full_detector_matrix/formal_run_20260824a/preregistration/detector-matrix-v2-plan.json evidence/r28_full_detector_matrix/formal_run_20260824a/receipts/raw-artifacts.sha256

Continued on next page

Table 27 continued.

Claim/evidence	Machine-readable source
Target-gate-suppression summary and replay receipt	evidence/r28_full_detector_matrix/formal_run_20260824a/detector-matrix-v2-summary.postexec-corrected.json evidence/r28_full_detector_matrix/formal_run_20260824a/postexecution-correction-replay-receipt.json; disclosed field-source amendment: evidence/r28_full_detector_matrix/postexecution-rr2-run-binding-correction.json
vLLM, SGLang, HYPIC timing/quality, and Marconi controls	evidence/related_work_same_protocol/; exact package members and one-command replays reside under this directory
Hydragen/Palu same-checkpoint transfers	evidence/round24_related_work_transfer/ and its executed_source/ directory
Published context and table generator	literature/reported_system_context.json and scripts/generate_related_work_context_table.py
Out-of-process GDN observer	evidence/r33_independent_capture/formal_h20/result/preregistration/preregistration.json evidence/r33_independent_capture/formal_h20/result/raw/out-of-process-gdn-capture.json evidence/r33_independent_capture/formal_h20/result/replay/out-of-process-gdn-replay.json evidence/r33_independent_capture/formal_h20/result/receipts/terminal-files.sha256
Independent preproducer slot census	evidence/experiment_registry.json, entry E-R39-INDEPENDENT-SLOT-CENSUS-A; execution package, formal archive, preexecution receipt, aggregate, controls, and terminal ledger are path- and SHA-bound there
Dual-producer repeat	evidence/experiment_registry.json, entry E-R39-DUAL-PRODUCER-REPEAT-A; execution package, formal archive/sidecar, summary digest, and exact replay boundary are bound there
Falcon-H1 bounded transfer	evidence/experiment_registry.json, entry E-R39-FALCON-H1-TRANSFER-V2-A; formal archive/sidecar, count clarification, validation, aggregate, and terminal ledger are path- and SHA-bound there
Primary compiled dispatch	evidence/experiment_registry.json, entry E-R40-PRIMARY-COMPILED-DISPATCH-V11-A; the immutable source package retains its pre-run HOLD metadata, while the separately hash-bound post-run aggregate and mirror authorize the bounded target-5 result
Selected-cell live binding	evidence/experiment_registry.json, entry E-R40-INDEPENDENT-LIVE-BINDING-V29-A, and evidence/r40_independent_live_binding_v29_postrun_audit_mirror/; the immutable package retains pre-run metadata, while the terminal result and hash-identical post-run audit authorize only the bounded selected-cell claim
Designer-executor-separated faults	evidence/r33_fresh_faults/author_freeze/designer_input/round32_input.pdf evidence/r33_fresh_faults/author_freeze/PROTOCOL.md evidence/r33_fresh_faults/executor_attempt_b/AMENDMENT.md evidence/r33_fresh_faults/formal_h20/RESULT_VERIFICATION.json evidence/r33_fresh_faults/formal_h20/r33-fresh-faults-20260825b/rank-run-*/rank-*/mutant-case.json evidence/r33_fresh_faults/formal_h20/r33-fresh-faults-20260825b/summary.json evidence/r33_fresh_faults/formal_h20/r33-fresh-faults-20260825b/terminal-files.sha256

Continued on next page

Table 27 continued.

Claim/evidence	Machine-readable source
Historical alias regression and repair	evidence/r35_historical_alias_regression/design_decision.md evidence/r35_historical_alias_regression/preregistration.json evidence/r35_historical_alias_regression/preexecution-freeze-receipt.json evidence/r35_historical_alias_regression/preexecution-resource-amendment.json evidence/r35_historical_alias_regression/formal_h20/RESULT_VERIFICATION.json evidence/r35_historical_alias_regression/formal_h20/r35-historical-alias-20260826a/aggregate.json evidence/r35_historical_alias_regression/formal_h20/r35-historical-alias-20260826a/replay/
Bounded two-stream lifecycle	evidence/r29_true_concurrency/formal_run_20260825b/raw/formal-result.json evidence/r29_true_concurrency/formal_run_20260825b/replay/independent-replay.json evidence/r29_true_concurrency/formal_run_20260825b/receipts/raw-and-replay.sha256
Expanded captured-boundary numerical sweep	evidence/r30_expanded_oracle_sweep/; 272 local arrays under raw/sidecars/; canonical result oracle-result.json; local replay source r30_expanded_oracle_reference.py; validation validation_report.json; terminal binding terminal-products.sha256
Local replay and storage accounting	evidence/experiment_registry.json, entry E-R40-LOCAL-REPLAY-STORAGE-COST-V1-A; protocol, Attempt-B amendment, aggregate, independent audit, and terminal ledger reside under evidence/r40_ci_cost_accounting_v1/
Manuscript claim index	evidence/claim_evidence_map.tsv