

FORKAUDIT: TRUSTED-CAPTURE VALIDATION OF HYBRID-LLM OWNERSHIP TRACES

Anonymous authors

Paper under double-blind review

ABSTRACT

Forking many continuations from one long prefix creates an ownership question that output equality alone cannot resolve: agreeing branches may still alias mutable recurrent state or share a defect. We present FORKAUDIT, a trusted-capture validator for ownership traces, instantiated on one fixed Qwen3.5-35B-A3B/H20 stack. Assuming an honest, mandatory-event-complete producer, detached replay checks immutable document state, phase-conditioned request/base and request/peer storage relations, tail copy-on-write, the registered recurrent-state transition, call contracts, and artifact-byte bindings. This is trace-relative validation, not an independent live monitor.

The primary RR2 run uses batch-one sequential round-major calls on one CUDA stream per rank. Across eight books, 96 KV-by-GDN ownership configurations and 288 adjacent-fan-out comparisons match on tokens, canonical CPU-FP32 full logits, logical KV, and final GDN state. A preregistered follow-up numerical sweep checks 20 attention rows spanning all ten full-attention layers and 24 GDN rows spanning 12 of 30 recurrent layers over two inputs: all 44 clean rows pass their captured-boundary FP32 oracles, and all 44 seeded wrong-operator controls are rejected. Nine designed faults reach their preregistered first gates. Of nine target-suppressed cases, five complete semantics: four preserve both tokens and canonical full logits, while M5 preserves tokens but changes logits. These are per-fault observations, not a detection rate. At $N = 32$, shared KV reduces the median final allocated delta from 4.901 to 2.229 GiB. Full live capture costs a median $4.321 \times$ in one frozen request-step cohort, positioning this implementation as an offline debugging or CI validator rather than a serving fast path. The results establish fixed-stack feasibility at declared observation points; capture honesty, compiled dispatch, unseen-fault coverage, and end-to-end correctness remain outside the conclusion.

1 INTRODUCTION

Best-of- N decoding, agents, tree search, and document question answering create many continuations from a common prefix. Modern serving systems exploit this structure through paged KV caches, prefix trees, and scheduling policies (Kwon et al., 2023; Zheng et al., 2024; Gim et al., 2024; Srivatsa et al., 2025). For a Transformer, “sharing the prefix” can often be described as sharing read-only KV blocks and giving each continuation private append blocks. Hybrid models complicate this picture: full-attention layers coexist with recurrent or linear-attention layers whose compact state is updated in place unless the runtime installs an explicit request-local view (Yang et al., 2025; Lenz et al., 2025; Pan et al., 2025).

This creates two evaluation hazards. First, output equality is only a relational observation: branches can agree while sharing mutable state, and compared implementations can share a defect. An ownership audit must additionally witness prefix immutability, phase-conditioned request/base and request/peer storage relations, tail copy-on-write, rebinding at the registered 32-token RR2 transition, and the call contract that produced the result. Second, “memory saving” may refer to logical payload, allocated cache pages, framework allocator peaks, process-level device memory, or admitted service capacity. Those quantities are not interchangeable.

We propose FORKAUDIT, a trusted-capture trace-validation protocol for this setting. It partitions document and request state, records page reservations and phase-aware recurrent rebinding, fingerprints the Python attention callable, and combines a 2×2 ownership factorial with relational and selected numerical checks. Its replayable claims assume an honest capture producer, as formalized in Section 3; it is neither an independent live monitor nor a cache policy, kernel, or scheduler.

The reusable contribution is the integration of a typed phase-indexed trace schema, fail-closed mandatory-coverage semantics, and ownership predicates that span token-indexed KV and mutable recurrent state. Paging, prefix reuse, copy-on-write, hashing, and metamorphic relations are established ingredients; our fixed-stack adapter is a case study showing that the integrated protocol can be instantiated, not a claim that those ingredients are new.

The primary contribution is a replayable ownership-assurance boundary, evaluated by one fixed primary case study and separately labeled support:

- **Typed trace obligations.** Mandatory records cover immutable document KV, private append pages, GDN bases and registered transition-boundary rebinding, tail COW, and call contracts. Trace coverage and replay verdict are reported separately and remain conditional on the honest-producer boundary.
- **Evidence beyond outputs.** Lifecycle, storage, and call predicates complement semantic relations. Nine designed mutants localize predeclared first violations; a separate target-gate-suppression matrix compares the same faults with token equality, exact FP32 logits, redundant audit gates, and production assertions. Five target-suppressed cases complete semantics; four preserve both tokens and exact logits while M5 changes logits. This is per-fault evidence, not held-out coverage or a detection rate.
- **Explicitly bounded numerical evidence.** Primary RR2 covers 96 ownership configurations, 288 cross- N comparisons, selected attention checks, and four allocator endpoints. A preregistered follow-up expands the captured-boundary numerical check to 20 attention and 24 GDN rows over two inputs, while remaining neither upstream nor end-to-end validation.

The primary RR2 inference remains case-specific: one Qwen3.5 model, one H20 family, 16-bit KV, one partial-tail geometry, a registered 32-token RR2 continuation block, batch-one model calls, and sequential round-major steps on one CUDA stream per rank. The source-distinct observer shares live objects, phase labels, and PyTorch’s storage API. The two-stream cohort establishes overlapping full-model-call intervals but neither individual-kernel simultaneity nor native continuous batching. Mac, deployment, HYPIC, operator-transfer, replay-cost, and live-overhead cohorts remain unpooled context or support and do not broaden the primary claim.

2 BACKGROUND AND RELATED WORK

Paged and prefix-aware inference. PagedAttention, Prompt Cache, RadixAttention, ChunkAttention, Hydragen, and Preble establish block sharing, prefix reuse, shared-prefix computation, and scheduling as prior art (Kwon et al., 2023; Gim et al., 2024; Zheng et al., 2024; Ye et al., 2024; Juravsky et al., 2024; Srivatsa et al., 2025). We ask instead what execution evidence demonstrates isolation when several state types coexist.

Hybrid state and cache management. Gated Delta Networks maintain mutable compact state unlike token-indexed KV (Yang et al., 2025); Jamba, Marconi, and HYPIC address hybrid architectures or reusable hybrid prefixes (Lenz et al., 2025; Pan et al., 2025; Liu et al., 2026c). Palu, SqueezeAttention, and HeadKV instead compress or budget KV (Chang et al., 2025; Wang et al., 2025; Fu et al., 2025). FORKAUDIT targets this kind of ownership boundary on our fixed stack rather than replacing those policies; Appendix Table 13 separates published-protocol numbers from our common-model controls.

Persistent intermediate residual reuse. CoMem persists a depth- j residual per document token for later selection and suffix execution (Liu et al., 2026a;b), motivating our persistent document object. Its bounded deployment slice supplies context; FORKAUDIT claims neither a new selector nor a new quality-latency frontier.




Figure 1: FORKAUDIT in one view. A common prefix forks into requests with distinct KV and recurrent-state lifetimes. Four ownership cells feed private append and recurrent-rebinding checks. The right-hand audit rail maps ownership/call receipts, selected FP32 attention checks, and positive-control faults into registered verdicts. The callouts summarize 96 configurations, eight attention rows, nine faults, and the $N = 32$ endpoint. Compiled dispatch remains open; this bounded schematic is not evidence or an independent-recapture claim.

Evaluation and metamorphic testing. SCBench emphasizes cache lifecycles (Li et al., 2025); metamorphic testing, MT-DLComp, and CRADLE motivate preserved relations and independent checks (Chen et al., 2018; Xiao et al., 2022; Pham et al., 2019); DNNMem separates memory denominators (Gao et al., 2020). We combine these ideas for heterogeneous state; Appendix Table 10 gives the component-wise boundary.

3 THE FORKAUDIT CONTRACT AND CONDITIONAL ASSURANCE BOUNDARY

Consider a document prefix D and resident requests $r \in \{1, \dots, N\}$. Let $\mathcal{A}$ be the full-attention layers and $\mathcal{G}$ the recurrent linear-attention layers. For $\ell \in \mathcal{A}$, the state is a read-only set of document pages K_ℓ^D plus a request-private reservation K_ℓ^r . For $\ell \in \mathcal{G}$, it is a persistent prefix base G_ℓ^D plus the request-local recurrent state G_ℓ^r . A logical fork is

$$\mathcal{F}_r(D) = (\{K_\ell^D, K_\ell^r\}_{\ell \in \mathcal{A}}, \{G_\ell^D, G_\ell^r\}_{\ell \in \mathcal{G}}). \quad (1)$$

Threat model and assurance boundary. The runtime may violate registered ownership, lifecycle, storage-view, or Python-call contracts. We trust the capture producer to record every mandatory event faithfully and bind each byte object to its archived artifact. Detached RR2 SHA-256 ledgers and schema replay expose missing, inconsistent, or modified records, but are not a second live observer. A separate source-distinct observer re-reads the same live candidate-created PyTorch objects without importing candidate capture rows or verdicts; it removes one direct code-sharing path but remains same-process, receives the same phase/completion labels, and trusts the same tensor/storage API. It is therefore not independent producer recapture or external ground truth. Replay can reject violations represented in the trace, not coherent producer omission/fabrication, an unattested host or compiled kernel/autotuning choice, transient mutation restored between snapshots, or unrecorded semantics. Accordingly, we report mandatory-trace *coverage* separately from the replay *verdict*; neither term denotes independent evidence completeness.

The checklist has seven audit targets. A target may be fully witnessed, partially witnessed, or left open; passing the case study must not silently upgrade a partial witness into a proof.

1. **Frozen identity.** Model, data, query bank, layer geometry, kernel descriptor, and protocol hashes are fixed before execution.
2. **Prefix immutability.** K_ℓ^D and G_ℓ^D have identical digests before and after all continuations.

3. **Private ownership.** KV reservations of distinct shared-document requests are disjoint in the common arena; independent materialized arenas have disjoint storage. At setup, a GDN request either owns a materialized base or holds an explicitly read-only persistent-base alias. At RR2’s registered 32-token transition boundary, all 60 mutable tensors must occupy ranges disjoint from the base and every peer.
4. **Tail-safe append.** If the last document page is partial, its valid tokens are copied once to a private page before a request appends.
5. **Dispatch provenance.** Both controls invoke the same recorded attention Python callable with the same query, mask, position, and append contracts; fallback and full-KV concatenation are forbidden. This does not fingerprint a compiled Triton binary or autotuning choice.
6. **Cross-arm equivalence.** For the same query and schedule, the materialized and shared ownership arms must have exactly equal generated tokens, full-vocabulary logits, final logical KV, and request recurrent state.
7. **Cross- N prefix consistency.** For each query observed at two or more fan-outs, its result at every larger fan-out must match its smallest- N result.

Targets 2–5 are pre-output ownership and trace obligations, whereas targets 6–7 are relational semantic checks. Matching tokens or logits is therefore insufficient. For the nine positive controls, lifecycle/ownership/storage receipts are first for M1 and M3–M5, and request-binding/call receipts for M2 and M6–M9. All stop before post-injection semantics: bounded first-gate localization, not unseen-fault coverage. A separate execution suppresses only the named target gate and records the next classified outcome under the predeclared precedence; it does not pool the two executions into a rate.

Conditional trace-validity statement. Let Σ be the frozen case specification, E an execution, and $\tau = C_\Sigma(E)$ the producer trace. For target i , let $M_i(\Sigma)$ be its mandatory record slots and Φ_i its registered replay predicates. Coverage is *complete* only when every slot in $M_i(\Sigma)$ has exactly one schema-valid, byte-bound record. We define

$$\text{Pass}_i(\tau) \iff [\text{Coverage}_i(\tau) = \text{complete}] \wedge \text{Bind}_\Sigma(\tau) \wedge \bigwedge_{\phi \in \Phi_i} \phi(\tau). \quad (2)$$

If capture is honest and mandatory-event-complete, and manifest binding, digests, and replay are correct, $\text{Pass}_i(\tau)$ implies that every registered predicate held at its captured observation point. A missing or modified mandatory record, or a represented predicate violation, cannot produce a pass. This statement is trace-relative: it provides no conclusion about coherent omission or fabrication, transient writes restored between observations, unrecorded compiled dispatch, common-mode semantics, or unseen-fault error rates.

Minimal adoption workflow and decision rule (normative). In our fixed Qwen/H20 case study, FORKAUDIT is instantiated in four steps. (1) *Freeze roles*: bind code, model, data, protocol, and runtime objects before outputs. (2) *Capture*: emit setup, registered-transition, and final lifecycle events with canonical digests, normalized ranges, ordered append/dispatch receipts, outputs, and allocator snapshots. (3) *Replay*: recompute all applicable predicates. Coverage is *complete* when every declared mandatory witness exists, *partial* for sampled, proxy, or lower-scope evidence, and *open* when a required witness is absent; the replay verdict is separately pass, fail, or not evaluated. (4) *Test sensitivity only when claimed*: pre-freeze the target-to-gate map and retain all outcomes. In the fixed Qwen run, terminal identity has complete trace coverage and passes replay; per-call compiled dispatch/autotuning coverage is partial. Appendix D gives the exhaustive target-to-record dependency matrix and the pointer-free storage schema used by this audit.

Appendix Table 17 gives the four normative record classes and their replay obligations.

Here and below, “exact” for a registered tensor-valued observable means equality of a canonical serialization: contiguous CPU-FP32 bytes with shape and dtype bound into the record, or the correspondingly specified logical-KV or GDN byte digest. It is not a claim that an unarchived device tensor was compared bit-for-bit before canonicalization.

Targets 6 and 7 are relational self-consistency oracles. They can reveal a layout-specific error or an N -dependent change in recorded observables, respectively. The primary factorial additionally captures the same post-RoPE query, logical K/V, position values, mask, and effective scale for one rank-specific full-attention row selected by the pre-execution static plan and compares the fused result with dense IEEE-FP32 attention. A separate recurrent oracle captures four transitions originally selected at layers 0, 10, 20, and 38 after the native BF16 q/k normalizer, then evaluates the transition with a candidate-import-free NumPy FP32 token recurrence. A fresh single preregistration, frozen before the successful run, binds this post-native-normalization boundary, the implicit query scale $1/\sqrt{128}$, the four rows, tolerances, and seeded faults. These are separately implemented numerical checks for selected operator boundaries, not for upstream activation construction, q/k normalization itself, every GDN layer, or the complete end-to-end model.

Figure 2 shows the four primary-factorial cells and the lifecycle they share. The KV axis contrasts one document-plus-private arena per request with a shared document and disjoint private tail/append pages. The GDN axis contrasts per-request setup materialization with a borrowed immutable base followed by the same functional rebind at the registered RR2 transition. Memory and witness cells are rebuilt separately so storage guards do not retain tensors during allocator measurement.




Figure 2: ForkAudit’s ownership factorial. The two left axes choose one KV and one GDN setup policy for the request group: append tails are always private, and a borrowed GDN base becomes private at the registered 32-token RR2 transition. The dashed lifecycle spine marks setup, the registered 32-token RR2 transition, and final capture. Left-to-right, the bracketed bottom row records document KV, private append, persistent GDN, request-private GDN, storage binding, call/semantic, and fault-to-gate evidence. Offline replay validates that bundle; within trusted capture, lifecycle/storage trace coverage is complete, dispatch coverage is partial, and numerical/fault support is selected. This mapped schematic is not experimental evidence.

4 TRUSTED-CAPTURE TRACE-VALIDATION PROTOCOL

4.1 FROZEN CASE-STUDY GEOMETRY

We evaluate primary feasibility on one fixed Qwen3.5-35B-A3B configuration (Qwen Team, 2026a;b): ten vLLM-0.26 Triton full-attention layers with BF16 (artifact “Q16”) KV and 128-token pages, plus thirty Transformers-torch GDN layers. RR2 uses eight H20-3e ranks, one book/rank, with batch-one model calls executed sequentially on one CUDA stream per rank. The ledger binds the Python callable, not the compiled binary; each request’s GDN digest covers 60 tensors. Claims are conditioned on this exact stack, input geometry, and schedule.

Eight ranks use distinct PG-19 books (Rae et al., 2020), each with one 4,095-token document and 32-query-token registered transition. Four cells evaluate $N \in \{1, 8, 32\}$ and eight subsequent argmax outputs. Here “state-appended” counts tokens consumed into KV/GDN state: the first model call consumes the 32-token query, the next seven consume generated tokens 1–7, and generated token 8 is output but is not fed into another model call, giving $32 + 7 = 39$. These 39 state-appended input tokens exercise COW on the 127-token final document page. Appendix Table 5 separates every supporting cohort.

4.2 AUDIT SEQUENCE

In primary RR2, each cell binds identity, rebuilds allocator and witness executions, records setup-registered-32-token-transition-final ownership, and compares tokens, logits, logical KV, and GDN state. Replay adds fixed-axis projections, one dense-FP32 sidecar per rank, and fresh matched-clean/live-fault pairs. Aggregation requires every byte binding; Appendix Table 6 states target strength.

4.3 CONTROLS, SCHEDULES, AND ORACLES

Primary RR2 requests follow sequential round-major steps on one CUDA stream per rank; this remains the scope of its factorial and memory claims. Full logits use shape/dtype-bound CPU-FP32 bytes; logical KV, 60 GDN tensors, bases, calls, storage, and allocator records use typed digests. Detached replay cannot reconstruct unarchived tensors or recapture the producer. A preregistered NumPy GDN oracle checks four layers after native q/k normalization, with fixed scale, tolerances, and four wrong recurrences.

Two unpooled supporting protocols probe other boundaries. The source-distinct observer reads the same candidate-created GDN tensors at setup, registered transition, and final generation for two fresh $N = 2$ policy cells using the fixed 4,095-token document and a 32-token registered transition, while constructing its own descriptors, opaque identities, relations, and lifecycle verdicts. On one 4,033-token document with two 16-token query requests at $N = 2$, the bounded two-stream protocol barrier-launches two request-local full-model calls from two host workers onto distinct CUDA streams and records CUDA-event-bracketed call intervals; cancellation, zero-scrub, reclamation, and same-slot replacement occur only after synchronization.

A separate paired protocol measures live audit cost on one H20. Each pair rebuilds one persistent 4,033-token document and two resident requests outside the timer, then executes byte-identical 16-token queries through the same `MultiForkHitLedger` and recorded Python attention callable. The baseline retains mandatory functional adapter checks but disables optional observers, ownership receipts, hashing, and audit persistence; instrumented execution enables the complete capture/receipt path. One warmup pair is discarded and five measured pairs alternate arm order and frozen request slot. Each arm synchronizes CUDA and resets peak-allocation statistics before timing; the timer ends after synchronized model completion and, for the instrumented arm, receipt persistence. The common full-vocabulary FP32 copy and semantic hash occur after timer stop.

4.4 MEMORY DENOMINATORS

We separate logical state, pages, synchronized allocator bytes, and unmeasured service capacity. Deployment “Store” is instead the union of tensor byte ranges owned by a retained document entry; it excludes metadata, pools, NVML deltas, and admission capacity.

4.5 GOVERNANCE AND LIVE MUTATION CHECKS

Pre-output ledgers bind code, model, data, queries, and protocol. Nine fresh matched controls pair with live mutations spanning allocation/binding, COW, GDN aliasing, positions, masks, callable identity, and paged-view type; each mutant first stops at its predeclared gate, establishing localization, not a rate; RR2 records no post-injection token, logit, logical-KV, or GDN digest. In a separate preregistered 18-case matrix, only that gate is suppressed. The precedence then records another FORKAUDIT gate, an allowlisted production assertion, the exact M8 payload sentinel, or completed-

horizon token and FP32 logit equality. Unallowlisted exceptions are invalid; the sentinel is not a production detector.

5 RESULTS

5.1 PRIMARY FIXED-STACK TRACE VALIDATION

Targets 1–4 and 6–7 have complete mandatory-trace coverage and a passing replay verdict for the frozen RR2 case. Dispatch provenance (target 5) passes at the recorded Python-call scope but has partial coverage because no per-call compiled-binary or autotuning-choice binding exists. Selected numerical oracles and designed mutations are bounded support, not additional target coverage. Source-distinct observation, two-stream lifecycle, live-capture cost, local replay, deployment, and related-work cohorts remain unpooled.

5.2 RELATIONAL EQUALITY HOLDS FOR EVERY REGISTERED OBSERVABLE

Across $8 \times 3 \times 4 = 96$ configurations (192 rebuilt allocator/witness cells), tokens, full logits, final logical KV, final GDN, both fixed-axis projections, and all 288 adjacent-fan-out comparisons match exactly. These relations cannot exclude a shared defect. Ownership receipts add that borrowed GDN starts as 60 immutable base aliases/request, materialized setup is disjoint, and both rebind at the registered 32-token transition boundary to 60 request-private tensors disjoint from bases and completed peers. Every phase passes; $N = 1$ directly exercises request–base disjointness.




Figure 3: Primary validation. Left: all selected attention rows beat the pre-fixed tolerance. Center: four ownership cells match on registered observables. Right: a separate RR2 all-gates-on reference reaches every target gate with passing controls. Numerical evidence, relational equality, and positive controls are distinct assurance levels; the last is not a rate.

Expanded captured-boundary numerical sweep. The primary dense-FP32 attention oracle obtains relative- L_2 $0.001484\text{--}0.001743 < 0.005$ on eight rows. A preregistered R30 follow-up uses two distinct PG-19 inputs and checks 20 attention rows spanning all ten full-attention layers (160 query positions) plus 24 GDN rows spanning 12 of 30 recurrent layers (192 token transitions). All 44 clean rows pass: maximum attention relative- L_2 is 0.001897; maximum GDN output/state relative- L_2 is $0.002073/2.08 \times 10^{-7}$. All 20 attention and 24 GDN seeded wrong-operator controls are rejected. The CPU/NumPy FP32 replay imports no candidate code, but begins from producer-captured post-RoPE attention and post-native-q/k-normalization GDN inputs. It therefore expands operator coverage without becoming upstream, end-to-end, or independent-capture validation; Appendix Table 23 fixes the boundary.

Table 1: Same-system outcomes for nine designed faults. RR2 enables all gates; R28 suppresses only the named gate in a separate matched case. “Same” is exact and “N/O” an earlier stop. Rows are not a detection rate; M8’s injected sentinel is not a production detector.

Fault / target contract	All gates on	With target gate suppressed	Token / FP32 logits
M1 Reservation ownership	target gate	completes	Same / Same
M2 Sequence binding	target gate	completes	Same / Same
M3 Tail COW	target gate	other audit: active-block ownership	N/O / N/O
M4 GDN base isolation	target gate	other audit: peer isolation	N/O / N/O
M5 GDN peer isolation	target gate	completes; logits change ($L_\infty = 0.75$, rel. $L_2 = 0.0314$)	Same / Changed
M6 Positions	target gate	completes	Same / Same
M7 Mask contract	target gate	completes	Same / Same
M8 Callable identity	target gate	injected payload sentinel	N/O / N/O
M9 Paged-KV view	target gate	paired-view production assertion	N/O / N/O

Table 1 adds a prospective same-system comparison. All nine controls pass and no case is invalid. Five suppressed mutants complete: tokens catch none, exact logits catch M5, and M1/M2/M6/M7 retain both. M3/M4 reach redundant audit gates, M9 a production assertion, and M8 only its injected sentinel. The receipts thus add observation for four designed violations on this horizon, not a rate, blind-fault result, or blanket superiority. Appendix Table 8 gives the gates.

Same-stack cancellation/reclamation checks match all 16 clean rank–geometry cells and reach the expected gate in all 48 preregistered schedule-fault trials (Tables 19–20); those cohorts do not test concurrent kernels, continuous batching, or completeness.

Source-distinct GDN observation. Across two fresh $N = 2$ GDN-policy cells and three phases/cell, a separately implemented same-process observer matches all 1,080 candidate row descriptors and all 96,660 ownership-relation coordinates, with zero mismatches. Its own 6/6 phase and 2/2 lifecycle verdicts pass. The serialized package contains 2,160 observer before/after rows plus 1,080 candidate rows; 1,080 is therefore the comparison denominator, not the total serialized-row count. Both paths inspect the same candidate-created tensors, consume the same phase/completion labels, and trust PyTorch’s storage API, so this is source-distinct PyTorch-level corroboration rather than independent or external recapture.

Bounded two-stream lifecycle. Two host threads on distinct CUDA streams produce positive full-model-call interval overlap both before cancellation (260.802 ms) and after quiescent zero-scrub, reclamation, and same-slot replacement (95.439 ms). All four lifecycle-indexed full-vocabulary samples match the serialized reference across 993,280 FP32 scalar pairs with zero maximum error; all four tokens match, although the samples contain only three unique logit digests. Twenty terminal logical-KV layer-digest pairs and two GDN slot aggregates covering 120 tensors also match, and ownership/lifecycle receipts pass. Cancellation occurs only after CUDA synchronization. The interval includes stream queuing, resource waits, and host-enqueue gaps; it does not prove simultaneous execution of individual kernels, native continuous batching, in-flight cancellation safety, or performance.

Paired live-audit cost. All five measured optional-audit-disabled/full-capture pairs, plus the discarded warmup, have exact full-vocabulary logits, identical token 353, and passing capture and ownership receipts. The median paired wall-time increment is 512.160372 ms (range 498.554564–542.293280 ms); the median paired ratio is $4.3210106\times$. Every measured instrumented arm writes 2,950,400 capture bytes and a 112,711-byte receipt, totaling 3,063,111 bytes (2.921 MiB). The paired incremental-peak delta alternates $-1.377, +1.377, -1.377, +1.377, -1.377$ MiB with frozen slot/order; its median -1.377 MiB is retained but is not evidence that auditing saves memory.

This is five-pair instrumentation cost for one frozen input and one 16-token request model call. It excludes document prefill and the common final-logit copy and does not estimate production latency, throughput, QPS, capacity, native serving overhead, or statistical significance. The measured $4.321\times$ path is an offline debugging/CI mode in this implementation, not the serving fast path; Appendix Table 12 reports all five measured pairs.

Table 2: Unpooled Mac motivation context (Apple M4 Pro, five fresh MLX processes; median over 30 query executions at split depth 7). Vanilla dense is the external reference. CoMem BF16/Q8/Q4 variants use the same selected documents. Store size is the six-document residual corpus; online time includes query Write, dequantization, and suffix Read. This small teacher-forced demonstration set is not primary ForkAudit evidence and supports no cross-hardware or production-speed claim.

Method	Store (MiB)	Online (ms)	Speedup	Top-1 vs. dense
Vanilla dense prompt	–	194.70	1.00×	100.00%
CoMem BF16 split	1.998	168.88	1.15×	71.89%
CoMem Q8 split	1.061	169.17	1.15×	71.89%
CoMem Q4 split	0.562	169.25	1.15×	71.89%

Q8/Q4 retain 100% top-1 agreement with the BF16 split path, yet all split rows inherit only 71.89% agreement with vanilla dense. The speed column is a small teacher-forced mechanism benchmark, not production RAG throughput.

Ownership factors move different memory components. At $N = 32$, switching full-copy to shared-document KV reduces the final allocated delta by 2.672 GiB with materialized GDN setup and 2.661 GiB with borrowed setup. Borrowing changes allocation timing, not the shared-KV end state: both GDN policies finish at 2.229 GiB. Appendix Table 18 gives the four endpoints and keeps the earlier $12.33\times$ pool versus $1.0398\times$ model-inclusive capacity cohort separate.

5.3 UNPOOLED MAC MOTIVATION CONTEXT EXPOSES COMMON-MODE AGREEMENT

An archival Apple-M4-Pro control illustrates common-mode agreement. Across five fresh processes on six selected document/query packs, CoMem BF16/Q8/Q4 (Liu et al., 2026a;b) share all answer-position top-1 decisions, yet agree with vanilla dense on only 71.89%. Table 2 motivates independent oracles; it is not primary ForkAudit or cross-hardware evidence.

5.4 UNPOOLED SAME-CHECKPOINT DEPLOYMENT AND NEAREST-WORK CONTEXT

Table 3 places two unpooled end-to-end blocks beside an independent retained-state cohort. All use the same checkpoint, eight LongBench items, one H20-3e/item, 4,096-token cap, greedy decoding, and at most 32 new tokens. Our Transformers block compares vanilla, full-prefix reuse, and CoMem variants; the official HYPIC (Liu et al., 2026c) block uses SGLang. Harnesses differ, so timing comparisons stay within block.

The rows expose state–quality–latency tradeoffs: full-prefix reuse is fastest in our block, while CoMem Q8 records 39.14 F1 with 15.89 MiB and mixed Q4/Q8 records 39.01 F1 with 10.01 MiB. Within the separate official-code block, Prefix Cache is the fastest row; HYPIC records 39.63 F1 at 0.101 s TTFT and 17.67 tokens/s. This TP=1 adaptation is not its published TP=2 reproduction. These eight items are context, not a complete LongBench or robustness result.

Table 3: Unpooled same-checkpoint H20 context on the eight-item Qasper/2WikiMQA slice. The first block is an in-process Transformers deployment cohort with three timing repeats per item. HYPIC (Liu et al., 2026c) timing/F1 comes from a separate 24-cell TP=1 official-code cohort; its Store values come from an independent receipt-instrumented 8+8-cell cohort. All use Qwen3.5-35B-A3B, one H20-3e per item, a 4,096-token cap, greedy decoding, and at most 32 generated tokens. Timing is interpreted within block; no cross-block speedup, production-capacity, broad-quality, or ForkAudit-ownership claim is made.

Method	Persistent state	Store (MiB)	TTFT (s)	TPOT (ms)	tok/s	F1
<i>Transformers 5.14.1 in-process deployment cohort</i>						
Vanilla dense	—	—	0.649	648.75	1.36	39.42
Full-prefix KV cache	Q16	140.34	0.163	54.11	13.38	39.14
CoMem state	Q8	15.89	0.673	55.44	7.15	39.14
CoMem mixed state	Q4/Q4/Q8	10.01	0.674	54.97	7.19	39.01
CoMem per-layer mixed	mixed	9.74	0.673	55.07	7.19	39.11
CoMem state	Q4	8.41	0.674	55.02	7.19	36.09
<i>HYPIC commit 98147c0 timing/F1; receipt-instrumented Store; SGLang 0.5.14, TP=1</i>						
Official HYPIC code: full recompute	—	—	0.217	49.21	14.77	39.14
Official HYPIC code: prefix cache	Radix prefix	139.53	0.072	50.15	19.07	39.14
HYPIC (transition + seam 8)	segment transition	324.09	0.101	52.13	17.67	39.63

Store is the median retained-document tensor payload: persistent document bytes for CoMem/full-prefix reuse and the union of target-entry-owned physical tensor byte ranges for Prefix Cache/HYPIC. It excludes Python metadata, preallocated pools, NVML/process deltas, and admission capacity. HYPIC timing/F1 uses streaming-client wall-clock from 24 cells; HYPIC Store uses a separate 16-cell externally replayed receipt cohort. Published Qwen3.5 HYPIC panels use TP=2, whereas this fixed same-slice adaptation uses TP=1.

In independent Store cells, Prefix Cache retains 139.53 MiB/document and HYPIC 324.09 MiB ($2.32\times$), versus CoMem Q8’s 15.89 MiB under the stated target-owned tensor-range estimand. This excludes metadata, pools, NVML, and capacity, and does not turn the separate timing blocks into a speedup claim.

Further unpooled context. vLLM/SGLang cache controls preserve 8/8 predictions (Table 15); bounded Hydragen/Palu transfers cover one captured layer family (Table 14), not a leaderboard, all-layer result, or combined system.

6 DISCUSSION AND LIMITATIONS

Conditional on honest capture and registered byte bindings, primary RR2 witnesses request/base and request/peer relations across ten attention and thirty GDN layers for its fixed 32-token registered transition and subsequent schedule. The suppression matrix adds observation for four designed faults beyond measured tokens/exact logits. The expanded follow-up numerical sweep covers all ten attention and 12 of 30 GDN layers over two short inputs, but it still begins from captured intermediates and does not close common-mode paths. The same-process source-distinct observer removes direct reuse of candidate capture rows and verdicts while retaining shared objects, labels, process, and PyTorch API. The two-stream cohort shows overlapping full-model-call intervals around a quiescent lifecycle transition, not simultaneous kernels or arbitrary in-flight cancellation. The measured $4.321\times$ live-audit ratio is the cost of the protocol’s full live-capture hooks, ownership receipts, hashing, and synchronous persistence path in one frozen 16-token request model call; this implementation is intended for offline debugging/CI rather than a production serving fast path. Evidence remains one model and H20 family with bounded inputs and schedules; lower-precision KV, compiled-binary identity, capacity, and end-to-end correctness remain untested. Context cohorts stay unpooled; Appendix I is exhaustive.

7 REPRODUCIBILITY STATEMENT

Manifest-first CPU replay covers the primary RR2 package and named replayable supporting cohorts, not every contextual table. RR2 rechecks 536 artifacts, 96 cells, 288 cross- N comparisons, eight attention rows, and nine fault pairs. R28 rechecks 18 cases, eight rank JSON files, and 24 FP32 sidecars; its disclosed postexecution wrapper changes only the inherited RR2 run-ID source and

preserves every candidate byte. The source-distinct-observer replay recomputes 1,080 descriptor and 96,660 relation comparisons. The bounded two-stream replay recomputes four full-logit comparisons and both interval overlaps while trusting archived lifecycle flags and the shared lifecycle helper.

The R30 expanded numerical sweep binds one preexecution source ledger and input manifest and includes 272 numerical sidecars totaling 140,199,936 bytes. A candidate-import-free CPU/NumPy replay rechecks 20 attention and 24 GDN clean rows plus 44 seeded controls. The validation result and terminal products ledger have SHA-256 values 56473175...9e41 and f125de7c...4e57, respectively. These are captured-boundary checks, not independent producer recapture or full-model replay.

The live-overhead formal result was frozen before replay-v1 exposed a raw-versus-derived request-GDN schema error. A disclosed postexecution, candidate-import-free replay-only amendment preserved the formal result and all raw artifacts, made no GPU rerun, recomputed row schema/order/token relations and `rows_sha256`, and passed replay v2. The formal result, replay v2, and terminal raw/replay ledger have SHA-256 values 3ccf86e2...8bc22, 7ede1b7e...cb4e, and 81808a4b...ca3a, respectively.

A preregistered Apple-M4-Pro practicality run executes one discarded warmup and five complete warm-cache RR2-package replays; 5/5 exit zero. The measured 630-file package occupies 892,284,156 logical bytes (850.95 MiB), and replay wall time has median/min/max 102.706536/101.372219/103.061831 s. This is local warm-cache reviewer cost, not live capture or serving latency. HYPIC separately binds 4,404 timing artifacts and hash-replays all 16 Store cells. The single-request fair-control row ships a verified summary but not its raw trials, so it is explicitly outside the replay-complete set. None of these ledgers provides producer recapture or compiled-dispatch attestation; Appendix J maps exact artifacts and boundaries.

8 CONCLUSION

On one fixed Qwen3.5/H20 stack, primary sequential RR2 shows that FORKAUDIT can validate a typed ownership trace containing tail-COW, registered GDN-transition, call, storage, and semantic receipts. The expanded numerical sweep passes 44 captured-boundary clean rows and rejects 44 seeded controls; four designed suppressed paths in the primary campaign retain canonical tokens/logits while their named trace gates reject them in separate all-gates-on runs. The conclusion is conditional: at declared observation points, the mandatory trace is covered and its registered predicates replay successfully under honest capture. It is not independent live recapture, compiled-dispatch attestation, unseen-fault coverage, or end-to-end correctness.

REFERENCES

- Chi-Chih Chang, Wei-Cheng Lin, Chien-Yu Lin, Chong-Yan Chen, Yu-Fang Hu, Pei-Shuo Wang, Ning-Chi Huang, Luis Ceze, Mohamed Abdelfattah, and Kai-Chiang Wu. Palu: KV-cache compression with low-rank projection. In *International Conference on Learning Representations*, volume 2025, pp. 50222–50249, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/7da6e0e00702c60607a6ae05c802ef85-Abstract-Conference.html.
- Tsong Yueh Chen, Fei-Ching Kuo, Huai Liu, Pak-Lok Poon, Dave Towey, T. H. Tse, and Zhi Quan Zhou. Metamorphic testing: A review of challenges and opportunities. *ACM Computing Surveys*, 51(1):4:1–4:27, 2018. doi: 10.1145/3143561. URL <https://doi.org/10.1145/3143561>.
- Yu Fu, Zefan Cai, Abedelkadir Asi, Wayne Xiong, Yue Dong, and Wen Xiao. Not all heads matter: A head-level KV cache compression method with integrated retrieval and reasoning. In *International Conference on Learning Representations*, volume 2025, pp. 99269–99290, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/f649556471416b35e60ae0de7c1e3619-Abstract-Conference.html.
- Yanjie Gao, Yu Liu, Hongyu Zhang, Zhengxian Li, Yonghao Zhu, Haoxiang Lin, and Mao Yang. Estimating GPU memory consumption of deep learning models. In *Proceedings of the 28th ACM Joint European Software Engineering Conference and Symposium on the Foundations*

- of *Software Engineering*, pp. 1342–1352. Association for Computing Machinery, 2020. doi: 10.1145/3368089.3417050. URL <https://doi.org/10.1145/3368089.3417050>.
- In Gim, Guojun Chen, Seung-seob Lee, Nikhil Sarda, Anurag Khandelwal, and Lin Zhong. Prompt cache: Modular attention reuse for low-latency inference. In *Proceedings of Machine Learning and Systems*, volume 6, pp. 325–338, 2024. URL https://proceedings.mlsys.org/paper_files/paper/2024/hash/a66caal703fe34705a4368c3014c1966-Abstract-Conference.html.
- Jordan Juravsky, Bradley Brown, Ryan Ehrlich, Daniel Y. Fu, Christopher Ré, and Azalia Mirhoseini. Hydragen: High-throughput LLM inference with shared prefixes. *arXiv preprint arXiv:2402.05099*, 2024. URL <https://arxiv.org/abs/2402.05099>.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles*, pp. 611–626. ACM, 2023. doi: 10.1145/3600006.3613165. URL <https://doi.org/10.1145/3600006.3613165>.
- Barak Lenz, Opher Lieber, Alan Arazi, Amir Bergman, Avshalom Manevich, Barak Peleg, Ben Aviram, Chen Almagor, Clara Fridman, Dan Padnos, Daniel Gissin, Daniel Jannai, Dor Muhlgay, Dor Zimberg, Edden Gerber, Elad Dolev, Eran Krakovsky, Erez Sa, Erez Schwartz, Gal Cohen, Gal Shachaf, Haim Rozenblum, Hofit Bata, Ido Blass, Inbal Magar, Itay Dalmedigos, Jhonathan Osin, Julie Fadlon, Maria Rozman, Matan Danos, Michael Gokhman, Mor Zusman, Naama Gidron, Nir Ratner, Noam Gat, Noam Rozen, Oded Fried, Ohad Leshno, Omer Antverg, Omri Abend, Or Dagan, Orit Cohavi, Raz Alon, Ro’i Belson, Roi Cohen, Rom Gilad, Roman Glozman, Shahar Lev, Shai Shalev-Shwartz, Shaked Meirom, Tal Delbari, Tal Ness, Tomer Asida, Tom Ben Gal, Tom Braude, Uriya Pumerantz, Joshua Cohen, Yonatan Belinkov, Yuval Globerson, Yuval Levy, and Yoav Shoham. Jamba: Hybrid transformer–mamba language models. In *International Conference on Learning Representations*, volume 2025, pp. 67959–67984, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/a9ed43fa31dc8b4a7d7a673d713dcb5f-Abstract-Conference.html.
- Yucheng Li, Huiqiang Jiang, Qianhui Wu, Xufang Luo, Surin Ahn, Chengruidong Zhang, Amir Abdi, Dongsheng Li, Jianfeng Gao, Yuqing Yang, and Lili Qiu. SCBench: A KV cache-centric analysis of long-context methods. In *International Conference on Learning Representations*, volume 2025, pp. 66063–66093, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/a540b17fb2295c736d5afd6c507acf66-Abstract-Conference.html.
- Hanzuo Liu, Xuan Qi, Chunyu Liu, Haotian Zhong, Yulong Wang, Key, Rayying, Alex Lamb, and Mingyu Gao. CoMem: Reusing transformer depth across queries with persistent intermediate residuals. In *COLM 2026 Workshop on Efficient Reasoning*, 2026a.
- Hanzuo Liu, Xuan Qi, Chunyu Liu, Haotian Zhong, Yulong Wang, Rayying, Key, Alex Lamb, and Mingyu Gao. Understanding is done early: A depth division of labor in large language models and its use for unbounded-context memory. *arXiv preprint arXiv:2607.28263*, 2026b. doi: 10.48550/arXiv.2607.28263. URL <https://arxiv.org/abs/2607.28263>.
- Yifei Liu, Juntong Wu, Yang Liu, Junhao Hu, Minghao Li, Xiaoxu Chen, and Weihang Chen. HYPIC: Accelerating hybrid-attention LLM serving with position-independent caching. *arXiv preprint arXiv:2607.01299*, 2026c. doi: 10.48550/arXiv.2607.01299. URL <https://arxiv.org/abs/2607.01299>.
- Rui Pan, Zhuang Wang, Zhen Jia, Can Karakus, Luca Zancato, Tri Dao, Yida Wang, and Ravi Netravali. Marconi: Prefix caching for the era of hybrid LLMs. In *Proceedings of Machine Learning and Systems*, volume 7. MLSys, 2025. URL https://proceedings.mlsys.org/paper_files/paper/2025/hash/7c180af017258d239bac6248d1eb26ac-Abstract-Conference.html.
- Hung Viet Pham, Thibaud Lutellier, Weizhen Qi, and Lin Tan. CRADLE: Cross-backend validation to detect and localize bugs in deep learning libraries. In *2019 IEEE/ACM 41st International*

- Conference on Software Engineering (ICSE), pp. 1027–1038. IEEE, 2019. doi: 10.1109/ICSE.2019.00107. URL <https://doi.org/10.1109/ICSE.2019.00107>.
- Qwen Team. Qwen3.5: Towards native multimodal agents. Official Qwen technical blog, February 2026a. URL <https://qwen.ai/blog?id=qwen3.5>.
- Qwen Team. Qwen3.5-35B-A3B-Base model card. Hugging Face model repository, 2026b. URL <https://huggingface.co/Qwen/Qwen3.5-35B-A3B-Base/blob/main/README.md>. Accessed 2026-08-16.
- Jack W. Rae, Anna Potapenko, Siddhant M. Jayakumar, Chloe Hillier, and Timothy P. Lillicrap. Compressive transformers for long-range sequence modelling. In *International Conference on Learning Representations*, 2020. URL https://iclr.cc/virtual_2020/poster_SylKikSYDH.html.
- Vikranth Srivatsa, Zijian He, Reyna Abhyankar, Dongming Li, and Yiyang Zhang. Preble: Efficient distributed prompt scheduling for LLM serving. In *International Conference on Learning Representations*, volume 2025, pp. 37057–37082, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/5bc342f48de8264779952fac378f96dc-Abstract-Conference.html.
- Zihao Wang, Bin Cui, and Shaoduo Gan. SqueezeAttention: 2d management of KV-cache in LLM inference via layer-wise optimal budget. In *International Conference on Learning Representations*, volume 2025, pp. 23593–23606, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/3b0a8df568ec496a717566a7f8158aaa-Abstract-Conference.html.
- Dongwei Xiao, Zhibo Liu, Yuanyuan Yuan, Qi Pang, and Shuai Wang. Metamorphic testing of deep learning compilers. *Proceedings of the ACM on Measurement and Analysis of Computing Systems*, 6(1):1–28, 2022. doi: 10.1145/3508035. URL <https://doi.org/10.1145/3508035>.
- Songlin Yang, Jan Kautz, and Ali Hatamizadeh. Gated delta networks: Improving Mamba2 with delta rule. In *International Conference on Learning Representations*, volume 2025, pp. 29687–29707, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/4904fad153f6434a7bcf04465d4be2cc-Abstract-Conference.html.
- Lu Ye, Ze Tao, Yong Huang, and Yang Li. ChunkAttention: Efficient self-attention with prefix-aware KV cache and two-phase partition. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 11608–11620, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.623. URL <https://aclanthology.org/2024.acl-long.623/>.
- Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying Sheng. SGLang: Efficient execution of structured language model programs. In *Advances in Neural Information Processing Systems*, volume 37, pp. 62557–62583. Curran Associates, Inc., 2024. doi: 10.52202/079017-2000. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/724be4472168f31balc9ac630f15dec8-Abstract-Conference.html.

A ETHICS STATEMENT

This work introduces no human-subject data collection and evaluates only systems behavior on the PG-19 training split. It does not assess model safety, bias, privacy, or downstream deployment quality. Memory efficiency can lower deployment cost but may also increase the scale of model use; those broader effects are outside this controlled systems study.

B COMPARISON LATTICE AND SAMPLING UNITS

This appendix describes the earlier capacity cohort only; its counts are not pooled with the pre-specified-window primary factorial.

Token trajectories	8 pass	16 pass	32 pass	64 pass	128 pass	256 pass
Paired full-vocabulary logits	64 pass	128 pass	256 pass	512 pass	1,024 pass	2,048 pass
Logical KV (10 layers)	80 pass	160 pass	320 pass	640 pass	1,280 pass	2,560 pass
GDN state (60 tensors/request)	8 pass	16 pass	32 pass	64 pass	128 pass	256 pass
Attention callable contract	1,280 pass	2,560 pass	5,120 pass	10,240 pass	20,480 pass	40,960 pass
	$N = 1$	$N = 2$	$N = 4$	$N = 8$	$N = 16$	$N = 32$

Resident fan-out (eight independent PG-19 books/ranks per column)

Totals: 504 request pairs; 4,032 paired logit comparisons; 5,040 KV layer digests; 504 GDN digests; 80,640 attention-call records. All compared artifacts passed.

Figure 4: Deterministic comparison counts by resident fan-out. Eight books are independent workload sources, each with one 32-query bank. Query sets are nested prefixes: $N = 1$ reuses query 0, $N = 2$ queries 0–1, and so on. Consequently, 504 paired request instances reduce to 256 unique queries, and only request indices 0–15 on each rank occur at two or more fan-outs. The cells count derived comparisons, not stochastic replicates.

Within this earlier two-arm cohort, all 504 paired request instances also match on the registered token, full-logit, logical-KV, and GDN records, and the 80,640 archived attention-call receipts use one fixed Python callable without a recorded path violation. These archival controls do not replace the primary four-cell RR2 evidence in Section 5. The nested fan-outs cover 256 unique queries rather than 504 independent examples: 128 queries recur at a larger fan-out and yield 248 non-vacuous cross- N comparisons, all matching. Across the two arms, the ledger contains 4,032 paired full-vocabulary step comparisons and 8,064 arm-specific model steps. Full logits and states are not archived directly; replay instead checks shape/dtype-bound canonical CPU-FP32 logit bytes and logical-KV/GDN digests.

C PROTOCOL GEOMETRY

Table 4: Frozen primary-factorial execution geometry. The eight ranks provide distinct books and query banks, not repeated stochastic measurements.

Component	Frozen value
Model	Qwen3.5-35B-A3B; 10 attention + 30 GDN layers
Hardware	$8 \times$ H20-3e; one book/rank
Software	Python 3.11.13; PyTorch 2.11.0+cu129; Transformers 5.14.1; vLLM 0.26.0+cu129; Triton 3.6.0; CUDA 12.9
KV backend	vLLM fused attention; BF16/Q16; page 128
Document / query	4,095 / 32 tokens; PG-19 train only
Generation	8 output tokens; 39 state-appended input tokens (32 query + 7 generated-token feedback)
Resident counts	1, 8, 32; four KV-by-GDN ownership cells
RR2 schedule	all requests resident; batch-one sequential round-major calls on one CUDA stream/rank
Validation	96 factor configurations; separate memory/witness cells
Fault campaigns	$N = 2$; RR2 has nine live mutants with rebuilt controls; R28 has nine fresh clean plus nine target-gate-suppressed cases on 8 H20s

Table 5: Cohort-to-claim authorization. Cohorts are never pooled; each row supports only the listed inference.

Cohort	Frozen scope	Authorized claim / explicit exclusion
Primary factorial (RR2)	8 books; $N = 1, 8, 32$; four KV $\times$ GDN cells; 8 decode steps	ownership lifecycle, cross-cell equivalence, selected attention oracle, live-mutant sensitivity, and four-cell allocator endpoints; not transfer or scheduler evidence
Earlier capacity	8 books; $N = 1, 2, 4, 8, 16, 32$; two KV arms	cross- N comparison counts and source-retaining capacity denominators only; not a primary-factorial replicate
Fresh-process control	one independent $N = 2$ process; one clean path and M1	input reconstruction, cleanup, and one intended-gate reproduction; not pooled with the factorial
Lifecycle transfer	8 books; aligned 4,096-token prefix; $N = 4$; two cancelled slots	exact outputs/state, zero-scrub, exact reuse, and stale-lease rejection on the same adapter; not scheduler, concurrency, or performance evidence
Scheduler interleave	8 books; $N = 4$; page/prefix/tail geometries 64/3,985/17 and 128/4,033/65; cancel slot 3 after round 2	deterministic request-step interleaving, zero-scrub, epoch change, exact reuse, uninterrupted-control equality, and three frozen scheduler-fault gates; not concurrent kernels, continuous-batching performance, detection rate, or production end-to-end correctness
Source-distinct GDN observer	one H20; fixed 4,095-token document and 32-token registered transition; two fresh $N = 2$ shared-KV cells (shared-base/materialized GDN); three phases/cell	zero mismatch over 1,080 candidate descriptors and 96,660 ownership relations; 6/6 phase and 2/2 lifecycle verdicts pass; same process/objects/labels/PyTorch API, not independent recapture, KV observation, or external ground truth
Bounded two-stream lifecycle	one H20; one 4,033-token document; two 16-token query requests at $N = 2$; two host threads and distinct CUDA streams; pre-cancel and post-reclaim phases; cancellation after synchronization	positive call-interval overlap plus exact four lifecycle outputs, 20 logical-KV layer pairs, two GDN slot aggregates, ownership, scrub, and reuse; not simultaneous kernels, native continuous batching, in-flight cancellation, production scheduling, or performance
Live audit overhead	one frozen Qwen/H20 PG-19 input; 4,033 document tokens, 16 query tokens; one discarded warmup plus five alternating optional-audit-disabled/full-capture pairs; mandatory functional adapter checks retained in both arms	paired request-step wall/incremental-peak/artifact cost with exact semantic and receipt gates; not prefill, final-logit-copy, serving latency, throughput, QPS, capacity, cross-input generality, or statistical significance
Local complete-replay overhead	Apple M4 Pro; already-unpacked 630-file RR2 package; one warmup plus five measured warm-cache complete replays	850.95 MiB logical footprint and 102.707 s median local replay; not H20 live capture, cold download/extraction, service latency, device memory, or integration effort
GDN transition oracle	1 frozen book/window; layers 0, 10, 20, 38; 4 query tokens	clean recurrence outputs/states and four seeded wrong-transition rejections conditional on post-normalization inputs; not all GDN layers or end-to-end semantics
R30 expanded numerical sweep	2 distinct PG-19 inputs; all 10 attention layers and 12 of 30 GDN layers; 20 attention and 24 GDN clean rows plus one seeded control/row	captured-boundary numerical agreement and rejection of all 44 seeded wrong operators; not upstream-activation, independent-capture, full-model, or end-to-end validation

Continued on next page

Table 5 continued.

Cohort	Frozen scope	Authorized claim / explicit exclusion
Target-gate suppression matrix	9 fresh clean + 9 target-suppressed cases; same fixed model/data/ $N = 2$; 8 distinct H20s	per-fault comparison with redundant audit gates, exact allowlisted production assertions, and token/FP32-logit equality; M8’s injected sentinel is not a detector; not a blind-fault study or rate
Mac common-mode control	Apple M4 Pro; six documents/queries; five fresh MLX processes; vanilla dense plus CoMem BF16/Q8/Q4 at depth 7	separates split-family agreement from agreement with vanilla dense; motivation only and separate from the H20 case
H20 deployment context	8 LongBench validation workloads; vanilla dense, full-prefix KV reuse, and CoMem Q8/Q4/mixed variants (six configurations total) with 3 timing repeats	same-stack reusable-state, TTFT, TPOT, tokens/s, and mean F1 comparison; not pooled with RR2 and not scheduler/continuous-batching evidence
Matched serving controls	8 LongBench validation workloads; vLLM 0.26 and SGLang 0.5.17 cache off/on plus 24 fixed-configuration HYPIC timing/quality cells; one H20-3e per workload	within-block cache-hit, quality, and client-wall timing context; not pooled with in-process CoMem timings, not a cross-framework leaderboard, and not continuous-batching evidence
HYPIC retained-state receipts	8 Prefix Cache + 8 HYPIC cells on the same fixed TP=1 model/slice; receipt-instrumented SGLang lifecycle	median target-entry-owned physical tensor payload only; not timing, NVML/process memory, capacity, continuous batching, or ForkAudit ownership evidence
Related-work operator transfer	Hydragen at $N = 8/32$ on one captured layer-3 attention family; Palu on one layer with 8 calibration and 1 held-out PG19 row	same-checkpoint numerical and logical-state trade-offs only; not an all-layer implementation, LongBench/quality result, speedup claim, or jointly optimized system

Table 6: Trace coverage and replay verdict for the seven contract targets in the frozen case. Equality denotes canonical byte-digest identity as recorded. Overall trace coverage remains partial because dispatch lacks a per-call compiled-binary and autotuning-choice binding. A passing verdict is conditional on the trusted producer capture and target boundary in Section 3; it is not independent producer recapture or comprehensive system correctness.

Target	Trace coverage	Replay verdict	Decisive witness	Boundary
Frozen identity	complete	pass	identical preflight/terminal source ledgers, terminal model rehash, and stable data/query/protocol/execution bindings	one validated formal run; not independent producer recapture
Prefix immutability	complete	pass	physical document-KV digests and GDN base guards across all three phases	one frozen configuration/geometry
Private ownership	complete	pass	all-pairs request-base and request-peer range replay at the registered 32-token RR2 transition	pointer-free conservative intervals
Tail-safe append	complete	pass	one-time 127-token private-tail copy before append	one partial-tail geometry
Dispatch provenance	partial	pass at	Python callable, query/position/mask/append/scale receipts; four compiled cache artifacts bound	no per-call compiled-binary or explicit autotuning-choice binding
Cross-arm equivalence	complete	pass	tokens, shape/dtype-bound CPU-FP32 logit digests, logical KV, and GDN	one primary RR2 sequential schedule
Cross- N consistency	complete	pass	288 all-arm adjacent-fan-out comparisons over 72 rank-query identities	relational, not an independent oracle

Selected attention and GDN numerical oracles, the nine targeted live faults, the lifecycle/scheduler extensions, and the source-distinct, bounded two-stream, live-overhead, and local-replay cohorts are supporting checks rather than additional contract targets. Their captured-input, same-adapter, same-process, schedule, and measurement boundaries are reported separately.

D POINTER-FREE STORAGE WITNESS SCHEMA

Target-to-record dependencies. The following matrix is normative. I, L, E, and F denote the Identity, Lifecycle, Execution, and Target/fault classes in Table 17; parenthesized entries name mandatory fields. Missing any marked field makes that target open. A preregistered geometry may make tail safety not applicable; the absence of a compiled-binary fingerprint leaves dispatch partial rather than silently full. A missing F field invalidates the corresponding fault row, but does not change the separately evaluated seven-target witness.

Target	I	L	E	F	Blind predicate
Frozen identity	code, model, data, protocol; role map	—	—	N/A	exact bytes and complete roles
Prefix immutability	document/base start digests	roles; setup and final guards	canonical content digests	N/A	start = final
Private ownership	request/base/private roles	setup, registered-transition, final bindings	normalized storage ranges	N/A	non-vacuous all-pairs disjointness
Tail-safe append	page geometry; document/private roles	first append event	ordered append and copied-tail receipt	N/A	one private copy before write
Dispatch provenance	callable and protocol identity	—	ordered call, query, mask, position, append, scale	N/A	schedule and callable equality; no fallback
Cross-arm equivalence	factor, query, and schedule identities	final bindings	token, logit, logical-KV, GDN records	N/A	canonical semantic equality
Cross- N consistency	query and fan-out identities	—	the same canonical semantic records	N/A	smallest- N equals each larger N
Fault sensitivity (supporting)	frozen target/gate map	matched-control rebuild and cleanup	injected outcome and failed predicates	required	target, injection, expected gate, outcome

Each tensor-view row has the exact fields `storage_id`, `shape`, `stride`, `storage_offset`, `dtype`, `device`, `storage_nbytes`, `tensor_nbytes`, and `content_sha256`; absolute addresses are forbidden. Within each phase, the producer assigns `storage-0000`, `storage-0001`, and so on in first-observation order to backing storages. Reusing an ID with a different device or storage size is invalid. Phase capture IDs are unique, whereas request and persistent lifecycle-guard IDs must remain stable across setup, the registered transition, and final generation.

Let element size be e , storage offset o , shape s_i , and stride t_i . The independent replay initializes $m = M = o$ and, for every dimension, computes

$$d_i = (s_i - 1)t_i, \quad m += \min(d_i, 0), \quad M += \max(d_i, 0). \quad (3)$$

The conservative half-open byte interval is

$$I = [em, e(M + 1)), \quad (4)$$

which must be non-negative and contained in `storage_nbytes`; also `tensor_nbytes` = $e \prod_i s_i$. Two rows overlap iff their normalized storage IDs match and $I_a^{\text{start}} < I_b^{\text{end}}$ and $I_b^{\text{start}} < I_a^{\text{end}}$. Exact aliasing additionally requires equality of interval, geometry, `dtype/device`, storage and tensor sizes, and content digest. This bounding interval can conservatively flag an exotic interleaved view, but it cannot miss a shared byte.

At setup, every recurrent-state coordinate must either exactly alias the persistent state (borrowed-base arm) or be disjoint from it (materialized arm). At the registered 32-token RR2 transition boundary, a completed request’s 60 mutable rows must be disjoint from the base and every peer; incomplete borrowed requests must remain exact aliases. At final, all request-local GDN states must be disjoint from the persistent GDN base and every peer; document KV remains intentionally shared. Independent binding-token receipts additionally require unchanged tokens for incomplete requests and changed binding/storage tokens for completed requests. The package applies this schema to all 96 timelines and 288 phase artifacts; nine adversarial tests cover exact aliasing, partial overlap, adjacent intervals, negative strides, conflicting ID reuse, pointer injection, capture/guard reuse, and cross-coordinate request–base and request–peer aliasing.

E COMPLETE CAPACITY LEDGER

Table 7: Memory scaling under three explicitly named estimands. The controlled pool retains the common source in both arms and is therefore a controlled contrast, not an optimized production full-copy baseline. Post-pack and full recorded-lifecycle PyTorch peaks are allocator counters, not NVML usage or service capacity.

N	Controlled pool (MiB)		Post-pack prod. peak (GiB)		Full lifecycle peak (GiB)	
	Materialize	Shared	Materialize	Shared	Materialize	Shared
1	170	85	64.896	64.813	67.434	67.434
2	260	90	65.048	64.881	67.434	67.434
4	440	100	65.345	65.011	67.434	67.434
8	800	120	65.936	65.270	67.434	67.434
16	1520	160	67.126	65.792	67.434	67.434
32	2960	240	69.498	66.837	69.498	67.434

The source-retaining analytic pool is exactly $80 + 90N$ MiB for Source+Materialize and $80 + 5N$ MiB for Shared-Documents. The post-pack production peak is the maximum of resident setup and generation after the phase reset. The full recorded-lifecycle peak is the maximum over prefill, Q16 pack, resident setup, and generation. Values were identical across all eight ranks; this deterministic run does not support a stochastic confidence interval.

These values are source-retaining controls: the deterministic pools are $80 + 90N$ MiB for materialization and $80 + 5N$ MiB for sharing. Removing the retained materialized source is only an arithmetic counterfactual, $80(N - 1)$ MiB; at $N = 32$ it gives a 2,480 MiB difference and has no allocator measurement. For the executed source-retaining cohort, the post-pack production peaks differ by 2,857,268,224 bytes (3.8289%; ratio 1.0398 $\times$), while the full recorded-lifecycle peaks differ by a 2,216,007,168-byte amount (2.9696%; ratio 1.0306 $\times$). These phase-bounded peaks are PyTorch allocated counters, not NVML memory or admitted service capacity.

F EXECUTED FAULT ROWS AND INVARIANT MAP

The row-level campaign is kept out of the main paper because it supports diagnosis and audit coverage rather than a separate headline result. Separately, the CPU governance suite passes 17 tests and rejects 28 explicit serialized-artifact tamper cases plus its registered helper/launcher faults. That result concerns schema and orchestration fail-closure only; it is not GPU semantic or performance evidence.

Table 8: Separate RR2 all-gates-on reference. Every separately rebuilt matched-clean cell passed; every mutant reached its target gate fixed before execution and was restored before disposal.

ID	Injected live fault	Expected/observed gate	RR2 outcome
M1	reservation alias	KV.RESERVATION.DISJOINT	target gate reached
M2	sequence swap	KV.SEQUENCE.ID	target gate reached
M3	omit tail COW	KV.TAIL.COW	target gate reached
M4	GDN-base alias	gdn.completed.vs.base.disjoint	target gate reached
M5	GDN-peer alias	gdn.completed.vs.peers.disjoint	target gate reached
M6	position off-by-one	POSITION.CANONICAL.VALUES	target gate reached
M7	materialized mask	MASK.CONTRACT	target gate reached
M8	wrong callable	KERNEL.CALLABLE.ID	target gate reached
M9	dense KV view	KV.PAGED.VIEW	target gate reached

Table 9: Representative failures each target is designed to expose. These are the nine executed primary-factorial live mutants; the checks are not complete oracles.

Target	Intended sensitivity and observed boundary
Frozen identity	wrong model/data/query bank, changed protocol, stale shard (schema adversaries)
Prefix immutability	M3 omits partial-tail COW; physical document digest and ownership gates fail
Private ownership	M1 reservation alias, M4 request-base GDN alias, M5 request-peer GDN alias
Sequence/view binding	M2 sequence swap and M9 dense-KV replacement
Dispatch provenance	M6 position offset, M7 materialized mask, M8 wrong callable
Factorial exactness	behavior changes when either KV or GDN setup ownership changes
Bounded numerical oracle	numerical error in a selected fused full-attention result
Raw replay	internally inconsistent trajectory, storage, allocator, or aggregate fields

The RR2 execution is a first-gate localization study. Every registered gate fires before the injected path produces a post-injection token, logit, logical-KV digest, or GDN digest; its output cells are therefore N/O rather than misses. Lifecycle/ownership/storage receipts supply the first gate for M1 and M3–M5; request-binding/call receipts supply it for M2 and M6–M9. The full audit rejects all nine at their predeclared gates under passing matched controls. The separate R28 matrix then suppresses one named gate per mutant while leaving the remaining observers active. It provides the per-fault comparison in Table 1, but does not estimate a population rate, held-out-fault coverage, or false-positive rate.

G COMPONENT-BY-COMPONENT NOVELTY BOUNDARY

Table 10: Primary-source positioning. “Evidence focus” describes what the cited work evaluates; it is not a claim that the work lacks all other checks.

Work	Primary mechanism	Evidence focus / relation to FORKAUDIT
PagedAttention	Paged KV allocation and sharing (Kwon et al., 2023)	Serving throughput and memory utilization; supplies the block-sharing substrate.
SGLang / Prompt Cache	Prefix reuse in structured or modular prompts (Zheng et al., 2024; Gim et al., 2024)	Runtime performance and output behavior; establish reuse and positional handling as prior art.
CoMem	One persistent residual per token at split j ; bounded selection; suffix replay (Liu et al., 2026a;b)	Motivates the reusable document object; the bounded deployment table supplies same-stack context, while the contribution here remains the hybrid-state ownership audit.
Marconi	Admission and eviction for hybrid-model prefixes (Pan et al., 2025)	Hybrid cache hit rate and time to first token; closest state-heterogeneous caching system, but a different policy-level contribution.
SCBench	Lifecycle-oriented KV-cache benchmark (Li et al., 2025)	Breadth over tasks and cache phases; exposes breadth absent from our single reference case.
MT-DLComp / CRADLE	Metamorphic compiler testing and cross-backend differential checking (Xiao et al., 2022; Pham et al., 2019)	Direct testing precedents; our audit adds ownership witnesses and a layerwise dense oracle, but not an independent end-to-end backend.
DNNMem	Component-wise GPU-memory accounting (Gao et al., 2020)	Supports denominator separation; it is an estimator, not a request-ownership audit.
FORKAUDIT	Typed, phase-indexed ownership/call/accounting trace with explicit coverage semantics	Integrates a mutation-tested KV-by-GDN ownership factorial, captured-boundary attention/GDN numerical sweeps, and same-adapter lifecycle/scheduler-contract checks; no new cache policy, production scheduler, independent recapture, or end-to-end oracle.

Table 11: Closest established precedent for each audit obligation. The final column states the narrower integration contributed by FORKAUDIT; it does not claim that any individual mechanism is new.

Audit obligation	Closest established precedent	Integrated obligation in this work
Frozen identity	Controlled variants in metamorphic and differential testing (Chen et al., 2018; Xiao et al., 2022; Pham et al., 2019)	Bind model, data, query bank, callable contract, and protocol bytes to every ownership receipt.
Document immutability and tail COW	PagedAttention block sharing and copy-on-write for shared prefixes (Kwon et al., 2023)	Replay physical document pages, padding, active tables, and the partial-tail detach event.
Recurrent-state ownership	Hybrid recurrent states and their in-place update constraint (Yang et al., 2025; Pan et al., 2025)	Witness request-base and request-peer ranges before and after the frozen schedule’s registered transition boundary.
Dispatch provenance	Cross-backend differential checking exposes implementation-dependent behavior (Pham et al., 2019)	Hold the backend fixed and record callable, position values, mask, scale, and append history per call.
Cross-arm and cross- N relations	Metamorphic relations over semantics-preserving transformations (Chen et al., 2018; Xiao et al., 2022)	Treat ownership policy and resident fan-out as transformations that must preserve registered request observables.
Bounded numerical oracle	Independent backend comparison (Pham et al., 2019)	Reconstruct selected attention inputs for dense IEEE-FP32 math and selected post-normalization GDN inputs for a candidate-import-free NumPy recurrence.
Memory denominators	Component-wise GPU-memory accounting (Gao et al., 2020)	Keep page geometry, allocator current/peak, and unmeasured process/service capacity as separate claim domains.

H SUPPORTING CONTROL EXPERIMENTS

Paired live-audit cost detail. The main text reports the paired summary and its offline-debugging boundary; Table 12 exposes every measured pair.

Table 12: All five measured live-audit pairs. B retains mandatory functional adapter checks but disables optional audit hooks; A enables complete capture plus ownership/receipt persistence. One warmup pair is excluded. Incr. B/A reports the incremental peak allocated above the allocation present at each arm’s timer start; Art. is the instrumented audit artifact. Both Δ columns are A minus B. The peak sign alternates with frozen order/slot and does not support a memory-saving claim.

Pair	Order	B ms	A ms	Δ ms	A/B	Incr. B/A MiB	Δ peak MiB	Art. MiB
0	B→A	153.620	695.913	542.293	4.5301	79.957/78.580	−1.377	2.921
1	A→B	149.398	665.831	516.433	4.4568	78.580/79.957	+1.377	2.921
2	B→A	157.068	665.128	508.060	4.2347	79.957/78.580	−1.377	2.921
3	A→B	153.939	652.494	498.555	4.2386	78.580/79.957	+1.377	2.921
4	B→A	154.218	666.379	512.160	4.3210	79.957/78.580	−1.377	2.921
Median	–	153.939	665.831	512.160	4.3210	–	−1.377	2.921

Published systems context. Table 13 reports each cited system only in its native protocol and keeps all non-commensurate metrics outside our matched panels. This unpooled table is literature context, not ForkAudit evidence.

Table 13: Unpooled published-systems context in each paper’s native protocol. These values show the scale and evaluation breadth of related systems, *not* a cross-paper leaderboard or ForkAudit evidence: models, hardware, batching, baselines, and metrics differ. Consequently, no row is merged with our same-stack H20 Table 3, and the absence of a common score must not be read as a negative result.

System	Native method and study setting	Reported efficiency in its native protocol	Quality and comparability boundary
PagedAttention / vLLM (Kwon et al., 2023)	Paged KV allocation plus serving engine <i>Setting</i> : LLaMA-7B on A10G and LLaMA-13B on A100-40GB; vLLM serving engine	2–4× serving throughput at similar latency vs. FasterTransformer/Orca	No common LongBench F1; exact cache-management method
Prompt Cache (Gim et al., 2024)	Modular attention-state reuse <i>Setting</i> : Llama2/MPT/Falcon; RTX 4090, A40, and A100; HuggingFace-style prototype	1.5–3.1× TTFT (CPU-resident modules); 3.7–11.7× (GPU-resident) on its 12-dataset LongBench suite	LongBench accuracy evaluated, but with different models and hardware
SGLang / RadixAttention (Zheng et al., 2024)	Prefix reuse plus structured-program runtime <i>Setting</i> : Llama-7B/Mixtral-8x7B and structured programs; A10G/A100; SGLang	Up to 6.4× throughput vs. prior inference systems	Application-suite result; no common LongBench F1
ChunkAttention (Ye et al., 2024)	Prefix-aware KV chunks and attention kernel <i>Setting</i> : attention microbenchmarks; mainly A100-80GB, plus A100-40GB/RTX 4090; CUDA 11.8	3.2–4.8× attention-kernel speedup for 1K–4K shared prompts	Kernel result; no end-to-end common task score
Hydragen (Juravsky et al., 2024)	Exact batched shared-prefix attention <i>Setting</i> : CodeLlama-13B/34B shared-prefix batches; A100 (with H100/L40S microbenchmarks); custom HuggingFace runtime	Up to 32× CodeLlama-13B end-to-end throughput in large shared-prefix batches	Exact attention construction; different model and batching regime
Palu (Chang et al., 2025)	Low-rank latent KV with reconstruction/fusion <i>Setting</i> : Llama/Mistral/LongChat quality studies; Llama-2-7B latency on one RTX 4090; HuggingFace plus custom kernels	50% KV compression and up to 1.89× RoPE-attention speedup; up to 2.91× with quantization	Llama/Mistral/LongChat native study; our bounded Qwen3.5 operator transfer is reported separately
Preble (Srivatsa et al., 2025)	Distributed prompt-sharing scheduling <i>Setting</i> : Mistral-7B/Llama-3-70B request traces; two–eight GPUs across A6000/H100 clusters; distributed serving	1.5–14.5× average-latency and 2–10× p99-latency improvement	Distributed serving traces; no common task score
Marconi (Pan et al., 2025)	Admission/eviction for hybrid-model prefixes <i>Setting</i> : Attention–Mamba2 7B geometry; LMSys/ShareGPT/SWE-bench traces; CloudLab policy simulation	Up to 34.4× token-hit rate and up to 71.1% or 617 ms lower TTFT	Hybrid models and trace-driven policy study; no common LongBench F1
HYPIC (Liu et al., 2026c)	Position-independent caching for hybrid full/linear-attention models <i>Setting</i> : Qwen3.5-35B-A3B main panels at TP=2 on 8 H20-3e; SGLang 0.5.14	3.25× lower TTFT on average and 1.66× higher QPS over Prefix Cache	Four models/five workloads; our same-slice TP=1 timing and retained-state controls are reported separately

Same-checkpoint transfer detail. Table 14 is the matched complement to the published context table: it uses our frozen Qwen3.5 checkpoint and H20-3e, binds the official Hydragen and Palu source revisions, and independently replays all raw operator sidecars. Because it does not execute all layers or LongBench, it is kept separate from Table 3.

Table 14: Bounded same-checkpoint H2O transfers. Hydragen uses frozen Qwen3.5 layer-3 post-RoPE attention; Palu uses eight 256-token PG19 calibration rows and one disjoint held-out row. Errors are relative ℓ_2 against the matched uncompressed operator; state ratios are logical selected-operator ratios, not process memory or end-to-end results.

Port	Setting	Matched numerical result	Logical state result	Timing / boundary
Hydragen port	$N = 8$	Hydragen-dense rel. $\ell_2 = 0.00255$	8.50 vs. 64.48 MiB (7.59× smaller)	0.287 vs. 0.0965 ms; no speedup
Hydragen port	$N = 32$	Hydragen-dense rel. $\ell_2 = 0.00262$	10.00 vs. 257.94 MiB (25.80× smaller)	0.290 vs. 0.259 ms; no speedup
Palu port	rank 64/head	held-out K/V 0.487/0.549 (plain 0.711/0.839)	512 vs. 2,048 B/token (4.00× smaller)	projection only; no fused kernel
Palu port	rank 128/head	held-out K/V 0.317/0.364 (plain 0.543/0.668)	1,024 vs. 2,048 B/token (2.00× smaller)	projection only; no fused kernel
Palu port	rank 192/head	held-out K/V 0.149/0.194 (plain 0.341/0.459)	1,536 vs. 2,048 B/token (1.33× smaller)	projection only; no fused kernel

Hydragen avoids prefix replication but has no speedup here. Palu improves over plain SVD, but residual error forbids an end-to-end quality claim. Neither is a jointly optimized combined system.

Matched serving-framework and HYPIC context. Table 15 gives separately timed vLLM and SGLang cache-off/on controls plus the official-code HYPIC Full Recompute, Prefix Cache, and full HYPIC arms under the shared Qwen3.5 streaming boundary. It remains a timing/quality-only panel and authorizes only within-framework interpretation; the separate 16-cell retained-state receipt cohort appears in Table 3. vLLM and SGLang each record 8/8 prefix hits and preserve 8/8 cache-off predictions. Within its TP=1 block, HYPIC gives 39.63, 0.101 s, and 17.67 tokens/s; Published Qwen3.5 panels use TP=2, so this is an adaptation. Its 24 timing cells bind a 4,404-entry terminal artifact ledger. Across the 16 retained-state cells, Prefix Cache retains a median 139.53 MiB and HYPIC retains 324.09 MiB; CoMem Q8 is 8.78× and 20.39× smaller, respectively, under the separate Store estimand. This remains single-stream; it does not measure continuous batching.

Table 15: Unpooled timing/quality-only serving context on Qwen3.5 and the same eight-item Qasper/2WikiMQA slice. Each row uses one H20-3e per item, a 4,096-token cap, greedy decoding, and at most 32 generated tokens. Times are OpenAI-compatible streaming client wall-clock; F1 is the eight-item mean. Hit is interpreted only against the disabled/full-recompute reference inside each framework block. No cross-framework speedup, retained-state, scheduler-throughput, capacity, or broad-quality claim is made.

Framework	Prefix reuse	F1	TTFT (s)	TPOT (ms)	tok/s	Hit
vLLM 0.26	off	39.63	1.445	56.31	4.70	–
vLLM 0.26	align on	39.63	0.179	53.99	14.00	8/8
SGLang 0.5.17	off	39.14	0.283	53.38	12.73	–
SGLang 0.5.17	Radix on	39.14	0.146	53.15	15.83	8/8
HYPIC code	full recompute	39.14	0.217	49.21	14.77	–
HYPIC code	prefix cache	39.14	0.072	50.15	19.07	8/8
HYPIC	transition+seam 8	39.63	0.101	52.13	17.67	8/8

vLLM uses experimental Mamba align; SGLang 0.5.17 uses extra_buffer Radix caching. The HYPIC rows use the authors’ commit 98147c0 on SGLang 0.5.14, TP=1, full transition.rope.recompute, and an eight-token seam. Its full and prefix rows are controls from that same codebase; published TP=2 results and the separate retained-state receipt cohort are not inserted here.

Hybrid-prefix policy context. The separate preregistered Marconi trace extends the eight frozen workloads with a fixed synthetic follow-up turn and evaluates only prefix-policy token reuse under the artifact’s native Attention–Mamba2 geometry. It neither reruns LongBench quality nor supplies Qwen3.5 serving speed, so it is not pooled with the preceding panel or the primary audit.

Table 16: Unpooled preregistered policy-simulator context on a 128-event synthetic multi-turn extension of the eight frozen workloads. Values are token-hit percentages under the official Marconi artifact’s native Attention–Mamba2 geometry. The wrapper calls each policy’s own eviction routine after insertion to close actual tree size to the declared budget. These are not Qwen3.5 quality, latency, or throughput measurements and do not validate the primary ownership audit.

Policy	5 GB	10 GB	20 GB
vLLM+	2.07	6.18	46.83
SGLang+	90.78	93.86	93.86
Marconi	93.86	93.86	93.86

Table 17: Normative record classes. “Mandatory” means that absence makes the corresponding target open rather than silently not applicable.

Record class	Mandatory payload	Independent replay
Identity	frozen code, model, data, protocol, and object-to-role map	exact bytes and role coverage
Lifecycle	setup, registered-transition, and final inventories and bindings	normalized ranges, aliases, and disjointness
Execution	ordered call/append receipts, semantic records, named allocator snapshots	frozen schedule, digests, and phase equations
Target/fault	predeclared target, gate, matched control, injection, and outcome	failed-predicate set and classified outcome

Primary allocator endpoints. Table 18 keeps the four ownership cells and their allocator denominators together under the primary 8-rank cohort.

Table 18: Primary $8 \times H20-3e$ allocator results at $N = 32$ (median across ranks; rank values coincide). Vanilla full-copy is the external control; the paged-prefix row is the closest same-stack prefix-sharing baseline. The remaining rows isolate recurrent-base ownership. All cells retain the persistent document source. Values are GiB above the frozen post-priming baseline; generation increment is the additional peak above allocation present at generation start.

Empirical arm	KV / GDN setup	Final	Peak	Gen. incr.
Vanilla full-copy	Full-copy KV / Materialized GDN base	4.901	4.920	0.019
GDN ownership ablation	Full-copy KV / Borrowed GDN base	4.890	4.907	1.951
Paged-prefix baseline	Shared-doc KV / Materialized GDN base	2.229	2.843	0.019
Audited hybrid fork	Shared-doc KV / Borrowed GDN base	2.229	2.843	1.950

Aligned cancellation and reclamation lifecycle. The separate lifecycle-transfer cohort uses eight ranks, eight distinct PG-19 training books, an aligned 4,096-token prefix, and $N = 4$. It is not pooled with the primary factorial.

Table 19: Lifecycle-transfer obligations on the existing Qwen3.5/vLLM Q16 adapter. Counts are deterministic rank outcomes, not stochastic replicates.

Obligation	Result
Frozen static/code/model/weight/raw bindings	pass; 14/14 weights, 8/8 shards
Full-vocabulary logits and generated tokens vs. uninterrupted control	exact, 8/8
Final logical KV and final GDN state	exact, 8/8
Aligned prefix, immutable document, and zero partial-tail copy	pass, 8/8
Cancelled-slot zero-scrub and exact reservation reassignment	pass, 8/8
Stale epoch-0 rejection and matched epoch-1 acceptance	8/8; 0 escapes, 0 wrong gates

Scheduler-managed request-step interleaving. The separate formal extension keeps four requests resident, cancels slot 3 after its second round, zero-scrubs and increments its lease epoch, reassigns the exact reservation, and completes the replacement interleaved with the three survivors. It is not pooled with the primary factorial or the aligned lifecycle-transfer cohort.

Table 20: Formal scheduler-managed request-step interleaving on the frozen Qwen3.5/ForkAudit vLLM-Q16 stack. Clean counts are deterministic rank outcomes; fault counts are preregistered expected-gate outcomes, not a detection rate.

Case	Frozen geometry or intervention	Outcome
Clean geometry 1	page 64; 3,985-token prefix; 17-token document tail; 18-event cancel/replacement lifecycle	all semantic/storage gates and uninterrupted-control comparisons pass, 8/8
Clean geometry 2	page 128; 4,033-token prefix; 65-token document tail; same lifecycle	all semantic/storage gates and uninterrupted-control comparisons pass, 8/8
FH1	dispatch cancelled epoch-0 lease	STALE_SLOT_LEASE, 16/16
FH2	dispatch lease under another request	LEASE_REQUEST_MISMATCH, 16/16
FH3	reclaim before zero scrub	RECLAIM_NOT_ZERO, 16/16

Candidate-import-free GDN transition oracle. The oracle uses one frozen PG-19 window and four query-time layer transitions at layers 0, 10, 20, and 38. A single preregistration frozen before the successful run binds the post-native-q/k-normalization boundary, scale $1/\sqrt{128}$, tolerances, and seeded faults. Its NumPy reference starts at that captured boundary; no candidate module is imported by the reference.

Table 21: PDF-visible oracle selection and replay boundary. Every listed row was locked before its successful candidate execution; neither oracle validates upstream activation construction or end-to-end model behavior.

Oracle	Frozen rows / selection rule	Captured boundary	Decision rule	Replay availability
Attention	rank r uses layer $4r + 3$ and round r , for $r \in \{0, \dots, 7\}$; request 0, $N = 1$, shared-document KV + borrowed GDN; frozen rank/layer/round map	post-RoPE Q/K/V, positions, mask, scale, and ordered append shadows	dense CPU-FP32 rel. $L_2 \leq 0.005$	8 raw rows and numerical sidecars; RR2 one-command replay
GDN	layers 0, 10, 20, 38; four query tokens from one frozen PG-19 window	post-native-q/k-normalization recurrence; scale $1/\sqrt{128}$	output rel. $L_2 \leq 0.005$; state rel. $L_2 \leq 10^{-4}$; seeded fault must reject	4 clean + 4 fault decisions from 40 sidecars; separate one-command replay

Table 22: Bounded GDN recurrence oracle. Relative errors compare native outputs and final states with the candidate-import-free NumPy FP32 recurrence.

Layer	Output rel. L_2	State rel. L_2	Clean	Seeded fault	Rejected
0	0.001584	1.23×10^{-7}	yes	omit decay	yes
10	0.001652	1.36×10^{-7}	yes	complement β	yes
20	0.001572	1.39×10^{-7}	yes	pre-decay memory	yes
38	0.001638	1.30×10^{-7}	yes	roll value heads	yes

Table 23: Preregistered R30 captured-boundary numerical sweep. Clean rows use candidate-import-free CPU/NumPy FP32 replay; every clean row has one seeded wrong-operator control. Attention begins after RoPE and GDN begins after native q/k normalization, so the table does not validate upstream activation construction, independent capture, or end-to-end model behavior.

Operator	Inputs	Layers	Clean rows	Positions / transitions	Max output rel. L_2	Controls rejected
Attention	2	10/10	20/20	160 positions	0.001897	20/20
GDN	2	12/30	24/24	192 transitions	0.002073 (2.08×10^{-7} state)	24/24

Full target-gate-suppression matrix. The preregistered R28 execution contains all M1–M9 faults and a fresh matched clean path for each. Across eight distinct H20s, all 18 cases complete their declared horizon or an allowed classified stop, all 24 FP32 sidecars verify, and no case is operationally invalid. Five mutant paths complete semantics: token equality catches 0/5, exact logits catch M5 only ($L_\infty = 0.75$, relative $L_2 = 0.0314$), and M1/M2/M6/M7 leave both exact. M3/M4 reach other audit gates; M9 reaches the paired-view production assertion; M8 reaches the exact injected payload sentinel, which is not a production detector. The frozen aggregator initially read the inherited run ID from a detached-manifest field that the manifest schema omits. A disclosed postexecution wrapper derives that ID from the canonical receipt embedded in all eight manifest-bound RR2 shards, changes only the generated comparison-row `run_id` from null to the preregistered value, and reproduces the summary byte-for-byte without modifying candidate outputs or frozen sources.

Single-request fair control. The supporting run evaluates eight workloads with four materialize and four shared trials per workload under the same recorded attention callable. All full-logit trajectories match. Median per-workload shared/materialize ratios are 1.0090 for the cached-document request TTFT, 1.0133 for the complete continuation first model step, and 1.0006 for TPOT. Incremental peak allocated memory decreases, but reserved-memory peak does not. Because timings are short, fixed-scope, and not a deployment scheduler measurement, we draw no speed claim. The anonymous supplement includes this verified summary as context, not its raw trials; it is therefore not a primary replayable result of this paper package.

I EXPANDED LIMITATIONS AND CLAIM BOUNDARY

The primary factorial covers one model, one H20 hardware family, one page size, 16-bit KV, and $N \leq 32$. Its ownership conclusion is tied to one 4,095-token partial-tail geometry, a registered 32-token RR2 continuation block, and the subsequent frozen sequential schedule; it is not a claim about an arbitrary first token or transition length. The lifecycle cohort adds one aligned 4,096-token geometry and a fixed cancel/scrub/reclaim sequence. The scheduler-interleave extension adds two geometries and deterministic request-step interleaving, but those requests execute sequentially on one CUDA stream.

The separate two-stream cohort has two host workers and distinct CUDA streams with positive full-model-call interval overlap. Its intervals include queuing, resource waits, and host-enqueue gaps; there is no per-kernel overlap trace. Cancellation occurs only after synchronization, so the cohort does not test arbitrary or in-flight reclamation. The source-distinct observer shares the process, candidate-created tensor objects, phase/completion labels, and PyTorch storage API; snapshot equality cannot exclude a transient write followed by restoration.

No cohort evaluates native continuous/ragged batching, eviction, multi-document serving, general scheduler cancellation, process-level NVML memory, OOM capacity, or a broad downstream benchmark. Full-copy KV and materialized GDN bases are controlled ownership factors rather than optimized production baselines. The GDN path is the Transformers Torch implementation, not an optimized recurrent kernel. The compiled-cache ledger does not bind each call to a selected binary or expose an explicit autotuning artifact.

The auxiliary serving panel uses one request stream per engine and compares each cache-enabled framework only with its own disabled reference. It does not measure continuous batching, eviction, admission control, maximum concurrency, or best-tuned performance. CUDA graphs were disabled in both pinned runs; the reported values describe that fixed configuration only.

The live-overhead cohort measures one 16-token request model call on one frozen input, with one discarded warmup and five measured pairs. Its timer includes backend setup/teardown and, for the audit arm, observer copies, ownership checks, hashing, and synchronous artifact persistence; it excludes prefill and the final full-logit copy. The $4.321\times$ median paired ratio is therefore neither GPU-kernel overhead alone nor a serving-latency, throughput, or capacity estimate. Five pairs support no significance claim. The incremental-peak delta alternates sign with frozen slot/order and authorizes no memory-saving conclusion.

The local replay measurement is warm-cache on one Apple M4 Pro after one unreported warmup. It does not measure cold download, extraction, page-cache population, H2O capture, engineering integration, or production operation.

The HYPIC timing/quality comparison uses the authors' exact tracked commit on SGLang 0.5.14 in one fixed CUDA 12.9 / PyTorch 2.11 configuration, adapting Qwen3.5 execution from the paper's TP=2 to TP=1 so every arm uses one H2O-3e per workload. Cache-hit authority comes from OpenAI response usage. Each of its 24 timing/quality cells has one formal measurement after a discarded, prefix-disjoint warmup. A separate receipt-instrumented run uses the same model/slice and fixed TP=1 serving configuration to produce 8 Prefix Cache and 8 HYPIC retained-state cells; each completed cell binds raw output plus pre-measurement and terminal ownership receipts, and all 16 pass an external hash-bound replay. The instrumentation adds target-entry byte-range receipts and a lifecycle-idempotency overlay; Store therefore does not come from the unmodified timing run. Neither cohort establishes scheduler throughput, QPS, capacity, the paper's TP=2 latency, or accuracy robustness beyond this eight-item slice.

The same-checkpoint related-work transfer is likewise operator-bounded: Hydragen uses one captured layer-3 shared-prefix attention family, and Palu uses one layer with eight calibration rows and one held-out row. Neither executes all model layers, a fused Palu kernel, LongBench, or a jointly optimized Hydragen/Palu-plus-CoMem serving path.

The page equations are deterministic identities, not learned laws; there is one observation per rank-configuration cell and no stochastic replication. Most runtime tensors are represented by hashes and equality receipts rather than archived in full. The attention and GDN oracle rows are the exception and ship numerical sidecars, but cover only selected captured-input operator boundaries rather than an end-to-end model. The R30 expansion covers two short inputs, all ten attention layers, and 12 of 30 GDN layers; it does not establish length/input robustness or full recurrent-layer coverage. Both the original and expanded GDN runs start after native q/k normalization and therefore do not validate that upstream boundary. Pointer-free normalized intervals plus raw-byte hashes make the captured storage evidence replayable, not a proof against a malicious producer before capture. The nine mutants are targeted positive controls; the target-suppression matrix adds a same-system comparison for those same designed faults, not blind or naturally occurring defects. Four complete with unchanged measured tokens and exact logits, one changes only full logits, and four stop earlier; N/O cells are not misses, while the injected M8 sentinel is not a production detector. This is not a general detection rate, a blind-fault result, or superiority over every conventional test suite. The scheduler extension's three frozen faults are expected-gate outcomes, not a detection rate or evidence of fault-set completeness. Broader systems claims would additionally require a native production serving scheduler, per-kernel overlap evidence, in-flight cancellation tests, optimized controls, alignment/length/input sweeps, and process-level capacity accounting.

J ARTIFACT AND CLAIM MAP

Artifact footprint. The original source-complete RR2 release contains 628 files totaling 892,144,066 uncompressed bytes (850.81 MiB); its verified raw ledger accounts for 536 artifacts and 888,785,811 bytes (847.61 MiB). The later package snapshot used by the local practicality measurement contains 630 regular files and 892,284,156 logical bytes (850.95 MiB). After one discarded warmup, five complete warm-cache CPU replays all exit zero with 102.706536 s median wall time. These are distinct package snapshots and local archival/replay costs, not online latency, traffic, device memory, or integration cost. Live capture cost is reported only by the separate paired H20 cohort in Table 12.

Table 24: Artifact and claim map. Relative paths are rooted at the accompanying research package. Completed rows authorize only their stated bounded claims. Metadata-redacted derivatives preserve numerical fields and state digests, while most full tensors still require the pinned environment for independent numerical recomputation.

Claim/evidence	Machine-readable source
Primary RR2 replay package	evidence/round_04_rr2_package/
Replay and package manifest	evidence/round_04_rr2_package/replay/run_replay.sh and evidence/round_04_rr2_package/MANIFEST.json
Registered, derived, and raw evidence	upstream/forkaudit-summary.json, derived/derived_summary_v2.json, and upstream/raw/ under the RR2 package
Executed source and storage tests	executed_source/gpu/ and replay/test_storage_witness.py under the RR2 package
Identity, lifecycle, scheduler, and GDN extensions	evidence/rr2_terminal_closure_a1_20260821/, evidence/lifecycle_transfer_reviewer_package/, evidence/round23_a2_scheduler_interleave/, and evidence/gdn_transition_oracle_preregistered_20260819d/
Mac and H20 deployment controls	evidence/mac_m4_motivation/, evidence/h20_deployment_benchmark/, and their generators under scripts/
HYPIC retained-state receipts	evidence/related_work_same_protocol/hypic-retained-state-r34-trial1892234/
R28 preregistration and raw ledger	evidence/r28_full_detector_matrix/formal_run_20260824a/preregistration/detector-matrix-v2-plan.json
R28 canonical summary and replay receipt	evidence/r28_full_detector_matrix/formal_run_20260824a/receipts/raw-artifacts.sha256 evidence/r28_full_detector_matrix/formal_run_20260824a/detector-matrix-v2-summary.postexec-corrected.json evidence/r28_full_detector_matrix/formal_run_20260824a/postexecution-correction-replay-receipt.json; disclosed field-source amendment: evidence/r28_full_detector_matrix/postexecution-rr2-run-binding-correction.json
vLLM, SGLang, HYPIC timing/quality, and Marconi controls	evidence/related_work_same_protocol/; exact package members and one-command replays reside under this directory
Hydragen/Palu same-checkpoint transfers	evidence/round24_related_work_transfer/ and its executed_source/ directory
Published context and table generator	literature/reported_system_context.json and scripts/generate_related_work_context_table.py

Continued on next page

Table 24 continued.

Claim/evidence	Machine-readable source
R29 source-distinct GDN observer	evidence/r29_independent_observer/formal_run_20260825a/raw/independent-gdn-observer-result.json evidence/r29_independent_observer/formal_run_20260825a/replay/independent-gdn-observer-replay.json evidence/r29_independent_observer/formal_run_20260825a/receipts/terminal-files.sha256
R29 bounded two-stream lifecycle	evidence/r29_true_concurrency/formal_run_20260825b/raw/formal-result.json evidence/r29_true_concurrency/formal_run_20260825b/replay/independent-replay.json evidence/r29_true_concurrency/formal_run_20260825b/receipts/raw-and-replay.sha256
R29 live-audit overhead	evidence/r29_live_overhead/formal_run_20260825b/raw/formal-result.json evidence/r29_live_overhead/formal_run_20260825b/replay/independent-replay-v2.json evidence/r29_live_overhead/formal_run_20260825b/receipts/raw-and-replay-v2.sha256; disclosed replay-only repair: evidence/r29_live_overhead/postexecution-replay-only-amendment-v2.json
R29 local complete-replay overhead	evidence/r29_overhead/local_replay_result.json, evidence/r29_overhead/validation_report.json, and evidence/r29_overhead/terminal-files.sha256
R30 expanded captured-boundary numerical sweep	evidence/r30_expanded_oracle_sweep/; 272 local arrays under raw/sidecars/; canonical result oracle-result.json; local replay source r30_expanded_oracle_reference.py; validation validation_report.json; terminal binding terminal-products.sha256
Manuscript claim index	evidence/claim_evidence_map.tsv